

Season 5 — The Bible Instrument

Research document. Twenty-two episodes. A different study, a different repository, the same posture. Four seasons built an instrument that kept failing honestly. This one takes the method and points it at the question everybody actually came for, and the first thing it does is fence off three quarters of that question and refuse to touch it.

*Two anchors per concept, as before. **Character** for the audience, **clinical** for you. And a third thing runs under this season that did not exist in the first four: **this study has a control, and the control is the whole reason any number in it means anything.***

The shape of the season

Three layers, kept strictly apart, and the separation is the design.

LAYER	THE QUESTION	WHERE IT LANDS
1a	do these texts interlock like independent reports, or collide like legends	run 1, and a control run
1b	when somebody alleges a contradiction, is it one, and what does the rescue cost	671 allegations, nine gates
2	is the historical case for the resurrection honestly built	not whether it happened

What this season will not do, said once and meant

It will not argue that the resurrection happened. Layer 2's own rulebook fixes that ceiling before any scoring, and episode 21 reads it out. The best available outcome there is "the case is honestly built", which is a fact about arguments and not about history.

It will not defend inerrancy. The study costs inerrancy in its second week and does not try to buy it back.

Inerrancy is the doctrine that the biblical texts contain no errors of any kind, including in matters of history and detail, not merely in what they teach. It is a strong position and it is not the only Christian one. Plenty of readers hold that these documents are reliable witnesses that also get details wrong, the way any two honest witnesses do.

That distinction runs under the whole season. **A finding that one of two texts is wrong about a fact is fatal to inerrancy and survivable for reliability**, and conflating the two is what makes this argument unwinnable in public.

And every headline in every file carries the same label: one coder, not blinded, who recognises every text on sight. No amount of process moves it past that, and the season says so roughly once an episode, because a study that states its ceiling twice and then forgets it is worse than one that never stated it.

The three things to watch for

One. Season 4 ended on an objection it could not answer: four texts about one week, written by the victim's own partisans, compared to what? **This season has the answer Season 4 lacked, and it is sitting in a folder called apocrypha.**

Two. A falsifier fires in this season **on the person running it**, by name, in the headline. Watch what he does next, because the structural argument he makes is either the best reasoning in six seasons or an escape hatch, and the episode makes you decide.

Three. The refrain has a new form here. **A check that cannot fail prints reassurance**, and this season finds two fourth-century versions of it, plus one father who states the opposite and gets it right in 340 AD.

Collecting Season 4

Nine rounds, two passes, and everything that worked belonged to the half of the instrument disclosed as compromised before it was built. The season closed on a question it had deferred all the way through:

Four documents about one week, written by the victim's own partisans. Can they be compared to anything at all?

That is not a rhetorical objection. It is the reason the dissent result could never be cleanly interpreted, because "the Gospels contain more objection language than Greek myth" has an obvious explanation that has nothing to do with Girard: **the Gospels were written by people on the victim's side.**

The answer Season 4 did not have

You compare them to **other documents written by partisans about the same events and the same person.**

They exist. They are called the apocryphal gospels, and there are enough of them to build a corpus: the Gospel of Peter, the Protevangelium, the Infancy Gospel of Thomas, the Acts of Pilate. Written by Christians, about Jesus, in the same tradition, with the same commitments.

That is a control, and it is the thing every measurement in Seasons 2, 3 and 4 was missing. The four-bucket instrument compared myth to tragedy, which are different genres by different people about different things. This compares a document to its nearest possible neighbour.

Take a moment on how much that changes, because it is easy to nod at and it is the reason this season exists.

In Season 4 the comparison was Gospels against Greek myth. When the Gospels came out different, there were a dozen available explanations and no way to choose: different language, different century, different genre, different translators, different purpose, and yes, possibly the thing being tested. **A difference between two groups that differ in a dozen ways tells you almost nothing**, because you cannot attribute it.

Here the two groups differ in far fewer ways. Same religion, same subject, same narrative form, overlapping episodes, both written by partisans. **When they come out different, the list of candidate explanations is short enough to argue about**, and a short list is the entire difference between a measurement and an observation.

That is why a cruder instrument here can be worth more than a refined one there, and episode 22 makes that comparison explicitly once both are on the table.

The device: the three layers, and why they are fenced

LAYER	THE QUESTION IT ASKS	WHAT IT CAN CONCLUDE	WHAT IT CANNOT
1a	do the texts interlock or collide	this looks like independent reporting, or it does not	nothing about whether the events happened
1b	is an alleged contradiction real, and what does the rescue cost	how much the accusers' list actually costs	nothing about inerrancy surviving
2	is the case for the resurrection honestly built	the argument is sound or unsound	nothing about whether a miracle occurred

Read the right-hand column, because the fences are the design. Every one of those exclusions is a place where a study like this normally leaks. You test transmission and then quietly talk as though you have tested the events. You cost an accusation and then talk as though you have defended a doctrine.

The ceiling, stated now rather than in a footnote

One coder. Not blinded. Who recognises every text on sight.

The person coding these items has read the Gospels, knows the famous arguments, and can tell Mark from the Gospel of Peter in one sentence. **Every classification runs through a mind that already knows which corpus it is holding**, and no procedure invented in this season fixes that.

So every headline carries the label **strong structured pilot, not a confirmatory result**. Not as modesty. As a fact about what the numbers can bear.

The distinction between a pilot and a confirmatory study is worth having properly, because it is not a matter of degree. **A pilot asks whether the measurement can be made at all**, and its output is a procedure plus an estimate of how hard the procedure was. A confirmatory study asks what the answer is, and its output is a number people are entitled to act on.

The two look identical on the page. Same tables, same percentages, same charts. **What separates them is entirely upstream**: whether the coding was blind, whether there was more than one coder, whether the sample was fixed before anybody looked.

So calling this a pilot is not a hedge to be dropped once the numbers look good. **It is a statement that the numbers were never eligible to be acted on**, whichever way they came out, and it was true before any of them existed.

The clinical anchor

Triaging into three streams and refusing to let them mix.

Chest pain goes to one pathway, the social problem to another, the medication review to a third. It is the same patient and they are genuinely connected, and the reason you keep the streams apart is that **a confident answer on one gets spent on the others** if you let them run together.

The commonest failure is not misdiagnosis. It is a resolved chest pain being treated as a resolved admission.

The character anchor

Rashomon.

Four people describe one killing in a wood and the accounts do not fit. Kurosawa's film is famous for being about truth being unavailable, which is the reading everybody takes away and is not quite what the film shows.

Watch what the accounts still share. There was a man, a woman, a bandit, a grove, a death. The accounts collide on motive, on who struck, on who is to blame, and they interlock on the furniture of the event.

That distinction is this whole season. Not "do the witnesses agree", which real witnesses never do, but **which layer they disagree on**, and whether the disagreements are the kind that independent people produce or the kind that invention produces.

What would count against this

1. **The apocrypha being a bad control.** They are later, they had the canonical gospels in front of them, and some are a different genre. **Episode 5 argues this properly**, and the answer is not that the objection is wrong.
2. **The layers not actually being independent.** If the Layer 2 arguments turn out to need Layer 1 accuracy, the fences do not exist and the framing is wrong. **A falsifier is registered for exactly that**, and episode 21 reports that it signalled.
3. **One coder being fatal rather than limiting.** It might be. The honest position is that no number here survives a hostile reading of that fact, and the season keeps saying so.

Handing off

Episode 2 takes an apologetic argument that has been made for two hundred years and asks the question nobody asks it: what would it look like if it were false?

Questions

1. Name another control for the Gospels, better or worse than the apocrypha, and say what it would cost.
2. The three layers are fenced. **Which fence do you expect to leak first?**
3. If one non-blinded coder is a ceiling, what would the next study need? Be specific about who, and about what they would not be told.

Collecting episode 1

Three fenced layers, a control that Seasons 2 to 4 never had, and a ceiling stated in advance: one coder, not blinded.

The argument, which is old and genuinely clever

It is called **undesigned coincidences**, and it goes back to William Paley in

1790. The shape is this.

One gospel says something odd and unexplained. Another gospel, elsewhere, in passing, for no apparent reason, supplies the fact that makes it make sense. Neither text flags the connection.

The worked example, and it is a good one. John says Jesus asks *Philip* where to buy bread for the crowd, which is a strange person to ask. Luke, in a completely different passage, mentions that the feeding happened near **Bethsaida**. John, elsewhere again, mentions that **Philip was from Bethsaida**.

Three texts, no cross-reference, and the odd detail resolves. **The claim is that forgery does not produce this shape**, because a forger supplies the explanation where the puzzle is.

Work that last clause, because it is the argument's actual engine and it is usually stated too fast.

If you are inventing a story, you write the explanation next to the puzzle. That is what writing is. You do not leave an unmotivated detail in chapter six and place its explanation in a different book by a different author, because you would have no reason to and no way to coordinate it.

What produces that shape is several people describing one world they all know about, none of them writing for the others. Each mentions what is relevant to their own telling, and the details fit because the world was consistent, not because anybody made them fit.

So the claim is not "this is too clever to be invented." It is narrower and better: **invention has a characteristic shape, and this is a different shape.** Which is a testable claim the moment you have some invented texts to compare against, and episode 5 goes and gets some.

Why it has never been an argument you could lose

Here is the problem, and it is the reason this study exists.

Nobody has ever counted the ones that do not work.

The literature is a collection of the successes. Paley found some, Blunt found more, McGrew found more again, and each is presented as a case. **There is no denominator and there is no control.**

WHAT THE LITERATURE HAS	WHAT IT DOES NOT HAVE
a list of coincidences that resolve	how many gaps were checked
each one argued in detail	how many gaps failed to resolve
the strongest cases	whether legends do this too

A list of hits with no misses and no comparison is not evidence. It is a scrapbook, and you can build one for any text long enough to have odd details in it.

That last sentence is testable and you should test it, because it decides how much of this argument survives.

Take any long novel with more than one book in the series. Find an unexplained detail. Then go looking, anywhere in the other books, for something that would explain it. **You will find something**, because long texts contain enormous numbers of details and any two details can be connected by a reader willing to connect them.

The question the apologetic literature never asks is how often that happens **by chance in a text nobody claims is reporting anything**. Without that rate, a handful of coincidences in four gospels is a number with nothing to compare it to.

The undesigned-coincidence argument may well be right. Nothing here says it is wrong. It says that as currently presented it is **unfalsifiable**, because no possible finding could count against it, and a version that could come back negative has never been built.

Which is this series' own sentence, arriving in the apologetic literature rather than in a research repository.

A check that cannot fail prints reassurance.

Two hundred years of collected coincidences is reassurance, printed steadily, by a procedure that has never once returned a negative because **nobody built the thing that could return one**. That is the same defect Season 2 found in a lexicon and Season 3 found in a null test, and it is sitting in an argument that predates both by two centuries.

What turning it into an instrument requires

Three things, and each one is a whole episode later.

1. **A denominator.** Every cross-checkable gap, counted, whether or not it resolves.
2. **A control.** The same procedure run on texts everybody agrees are legendary.
3. **A rule that stops the coder supplying the answer.** A fit counts only if a *text* provides the reconciling fact.

The third one is the sharpest, and it is where this study departs from the tradition it is testing. In the apologetic literature the connection is made by the apologist. Here, if the coder has to supply it, it is not a fit.

The device: the same gap, scored two ways

	JOHN ASKS PHILIP	SCORED AS
the coder knows Philip is from Bethsaida and says so	supplied by the coder	not a fit
John 1:44 states Philip is from Bethsaida, on the page	supplied by a text	FIT

The finding is identical. The provenance is not, and provenance is the only thing that separates evidence from a good reader.

The clinical anchor

Corroborating a history from two relatives interviewed separately.

The daughter mentions a fall three weeks ago. The son, asked about something else, mentions the new blood pressure tablet started a month ago. **Neither of them connected those and you did**, and the connection is real.

Now notice the discipline that makes it evidence rather than a story: **you write down what each person actually said, separately, before you join them**. A note that records only your synthesis has destroyed the thing that made it convincing.

The character anchor

A continuity error that turns out not to be one.

A character has a scar in episode 9 and no scar in episode 4, and the internet calls it a blunder. Then somebody points out that episode 6, in a throwaway line nobody cared about, has him mention a fight.

Nobody wrote that as a payoff. It resolves because the writers were tracking one world. And the reason this is the right anchor is that **fandoms also produce hundreds of these that do not resolve**, and nobody makes a video about those, which is precisely the missing denominator.

What would count against this

1. **Undesigned coincidences being explicable by literary dependence**. If Luke read John, the coincidence is a borrowed detail. **The study handles this by tagging independence**, and episode 6 reports that all five canonical fits are between independent texts.
2. **The whole argument proving less than it seems**. Even at its best it shows the texts behave like independent reports of something. **It does not show what**. A shared legend, independently recorded, produces the same shape.
3. **A denominator being impossible in practice**. It has not been built, and episode 17 treats that as the honest ceiling on the whole layer.

Handing off

Episode 3 builds the classification, four categories, and the interesting one is the category that exists to stop the coder helping.

Questions

1. Build an undesigned coincidence for a novel you know well. **How hard was it?** That difficulty is the argument's real strength or its real weakness.
2. Why is "a shared legend, independently recorded" not excluded by any result in this layer?

3. **Before episode 3:** how would you count a gap that neither resolves nor collides? It is the most common case and it decides the denominator.

Collecting episode 2

Undesigned coincidences are a two-hundred-year-old argument with no denominator, no control, and a habit of letting the person making the argument supply the connection. Turning it into an instrument means fixing all three.

Every cross-checkable gap gets exactly one of four labels

The unit is not a text and it is not a verse. It is a **gap**: a place where two texts describe the same episode and one of them says something the other does not, or says it differently.

You walk the **pericope**, which is the one piece of jargon worth adopting here. A pericope is a **self-contained episode in a gospel**: the burial, the feeding of the five thousand, the empty tomb. Not a chapter, which is a later editorial division imposed on the text, and not a verse. **It is the unit the story itself has.** Comparing gospels means lining up the same pericope in each and reading across, and the word exists because the alternative is saying "the bit where" for the rest of your life.

You walk it, you find every such place, and you label it.

FIT. One text supplies a fact that makes the other text's oddity make sense, and the supplying is done by the text rather than by you.

COLLISION. The two texts assert things that cannot both be true. Split into **hard**, where no reading reconciles them, and **soft**, where a reading exists but it is strained.

AGREEMENT. The two texts say the same thing. This sounds like the good case and it is the least informative label on the list, because agreement is what copying produces as readily as truth.

NULL. The gap neither resolves nor collides. One text mentions a detail, the other is silent, and nothing follows either way.

Why NULL is the category that makes this a study

Almost every gap is a NULL, and the temptation is to stop recording them, because recording a nothing feels like clerical work with no payoff.

NULL is the denominator. It is the count of chances the texts had to interlock and did not, and the count of chances they had to collide and did not. Without it you have a numerator floating in space, which is the exact condition the apologetic literature is in and the reason episode 2 called it a scrapbook.

The interlock ratio is fits divided by fits plus collisions. A text that produces five fits looks impressive until you learn it also produced five collisions, and unremarkable until you learn it had four hundred chances to do either.

So the boring label is the one carrying the argument, and a coder who gets bored and stops logging NULLs has quietly rebuilt the scrapbook while feeling efficient.

And notice the specific way that failure arrives, because it is not a decision anybody makes. Nobody thinks "I shall stop recording nulls." What happens is that the interesting items take longer to code, the nulls take seconds, and after two hours the threshold for

what counts as a gap worth logging has drifted upward without any moment where it moved.

The denominator does not get abandoned. It gets eroded, and the erosion is invisible in the output, because a coding sheet with fewer nulls looks exactly like a corpus with fewer gaps.

The rule that does the real work

A difference is a FIT **only if a text supplies the reconciling fact.**

That sounds like a technicality and it is the entire methodological content of the study. In the apologetic literature the connection is drawn by the person writing the book, and the reader cannot tell whether the connection was in the texts or in the author's knowledge of them. Here that difference is made procedural: **name the verse that supplies it, or it is not a fit.**

And the rule is symmetric, which matters more than the rule itself. It blocks the defender inventing a harmony, and it blocks the sceptic dissolving a collision by saying it feels fine. Both moves are the same move: supplying from outside what the text does not supply. The study calls this the **Ehrman guard**, and every later layer inherits it.

Bart Ehrman is a New Testament textual scholar, probably the best-known sceptical voice on how the gospels were transmitted, and the guard carries his name because of an objection of his about harmonising. Faced with two accounts that will not fit, a defender can always construct a third story in which both are true. Ehrman's point is that doing this is **writing a new gospel**: the reconciling narrative appears in none of the texts, so you have solved the problem by inventing a document nobody has.

Naming the guard after a critic is deliberate. It is a rule taken from the other side and then **applied in both directions**, which is the only version of it worth having, and episode 15 is what happens when somebody measures how often he broke it himself.

Never relax it in one direction only. A guard that binds your opponent and not you is not a guard. It is a rhetorical device with a procedural costume on, and it will return whatever you already believed with a number attached.

The device: one gap, four possible labels

WHAT THE TEXTS DO	LABEL	WHAT IT CONTRIBUTES
John's odd detail explained by Luke's aside	FIT	numerator
Mark says the women told nobody, others say they told	hard COLLISION	numerator, other side
a reading reconciles them but strains	soft COLLISION	numerator, discounted
both say Jesus was buried	AGREEMENT	almost nothing
one names an angel, the other does not	NULL	the denominator

The clinical anchor

The difference between a sign and a symptom, and who supplied it.

A symptom is what the patient reports. A sign is what you elicit. Both go in the notes and they are recorded separately, and the reason is exactly this study's rule: **an examination finding you produced is a different kind of evidence from a complaint the patient volunteered**, and a note that blends them cannot be audited later.

The failure mode is familiar. You suspect a diagnosis, you ask a leading question, the patient agrees, and it is written down as a symptom. Nothing false was recorded and the provenance has been destroyed.

The character anchor

A fan theory that quotes an episode, against one that does not.

Two people explain the same odd moment in a show. The first says "this makes sense because of what she said in season two, episode four," and gives the line. The second says "this makes sense because he was obviously lying."

Both explanations might be correct. Only one of them can be checked by somebody who does not already agree, and the difference between them is not cleverness. The first person is pointing at the text and the second is pointing at their own reading of it.

What would count against this

1. **NULL being unbounded.** How many gaps are there really? The answer depends on how finely you slice a pericope, and **nobody has fixed the slicing rule independently**. This is the denominator problem arriving early and it is not solved in this season.
2. **Soft collisions being a dumping ground.** A category for "collides but a reading exists" is exactly where a motivated coder puts the inconvenient items. The study reports hard and soft separately for this reason, which limits the damage without removing it.
3. **The Ehrman guard being unfalsifiable in practice.** Deciding whether a text "supplies" a fact is itself a judgement. **The guard reduces the judgement, it does not eliminate it**, and episode 15 is what happens when somebody measures how much judgement is left.

Handing off

Episode 4 tries to blind the coder, which cannot be done, and the interesting part is what gets built anyway and what the tooling can and cannot hide.

Questions

1. Take a gap you know from the Gospels. **Label it, then say who supplied the label.**
2. AGREEMENT is described here as almost uninformative. Construct the case where it is not.
3. **Before episode 4:** you cannot unknow the texts. Design the best blinding you can anyway, and name what it fails to hide.

Collecting episode 3

Four labels, a rule that a fit must be supplied by a text, and a guard that binds the defender and the sceptic identically. All of it still runs through one mind that knows every text on sight.

The problem, stated without flinching

The coder is going to read a passage and know immediately whether it is Mark or the Gospel of Peter. Not by an act of scholarship, by recognition. The canonical texts are among the most familiar documents in the language.

And the hypothesis under test is that the canonical texts behave better. So the person applying the labels knows, for every item, which way the label needs to fall for the prediction to pass.

That is the textbook setup for bias, and no amount of good intention touches it, because the bias does not operate through intention. It operates through **how long you look before you accept an answer**, which is the mechanism Season 4 episode 10 named and which is invisible from the inside.

What gets built anyway

Two mechanisms, both enforced by tooling rather than by resolve.

Order blinding. A script emits every item in shuffled order, so the coder does not work through Mark and then work through Peter. The identities of the sources are written to a separate file, `key.json`, and that file is sealed until scoring.

The sealed key deserves a sentence on why it is a file rather than a promise.

A coder who keeps the mapping in his head is not doing anything dishonest, and he is also carrying the tally as he goes. He knows the canon is at four fits and the apocrypha at none, and the next borderline item arrives into a mind that knows what it would do to the running total. **The bias is not in wanting a result. It is in knowing the score while you play.**

Putting the mapping in a file the coder does not open removes the scoreboard. It does not remove recognition, which is the next section, and it removes something real, which is why it is worth the ten lines of code.

Label blinding. The coder does not see which witness is which gospel while coding. The items arrive as anonymous witness codes.

The value of this is real and it is narrow. It removes the **run effect**, where your standards drift across a block of items from one source and you calibrate to that source without noticing. It does nothing about recognition.

What the tooling cannot hide, and the study says so

The texts carry their own content, so the run **cannot be truly blinded to canon against apocrypha the way a naive coder could be.**

Read that as the honest sentence it is. A witness code hides the label. It does not hide that this passage has a talking cross walking out of the tomb, which occurs in exactly one document in the corpus and which the coder will recognise in four words.

So the blinding is partial by construction, and the study's response is not to claim otherwise. It is to state the ceiling on every headline and to separate the findings by how much they depend on judgement, which is episode 7 and is the most useful thing in the layer.

That separation is the real answer to unblindable coding, and it generalises past this study, so it is worth stating as a rule.

When you cannot remove a bias, **sort your findings by how much of them the bias could touch**. Some results depend on a judgement call and some are facts on the page. A study that reports both at one confidence has hidden its weakest number inside its strongest, and a study that sorts them lets a hostile reader keep the half that survives.

The alternative, which is what usually happens, is a limitations paragraph at the end saying the coding was not blinded. **That is true, unusable, and costs the author nothing**, because it does not tell a reader which specific number to discount.

The device: what each mechanism actually buys

MECHANISM	REMOVES	DOES NOT REMOVE
shuffled order	drift across a block, calibrating to one source	recognition of any individual item
sealed key	scoring while watching the tally move	knowing which way the prediction needs it to fall
witness codes	the label	the content
nothing available		that the coder knows the hypothesis

The bottom-left cell is empty and the study leaves it empty. That is what an honest limitations section looks like: not a paragraph of hedging, a row with nothing in it.

The clinical anchor

Reporting imaging blind to the clinical details, and why nobody actually does it.

A radiologist given no history reports what is on the film. A radiologist told "query pneumonia, day three post-op" finds the consolidation faster and misses the thing nobody asked about. The literature on this is not subtle: **context improves sensitivity for the suspected diagnosis and reduces it for everything else**.

And you cannot run a service on blind reporting, because the context genuinely helps. So the compromise is the same one here: **you do not pretend the reporter is blind. You record what they were told**.

The character anchor

A blind audition with the chairs turned around.

It works, and it demonstrably changes who gets picked. And it does not survive somebody walking out and singing a song they released last year, because the mechanism

hides the face and not the voice.

Nobody concludes from that failure that the chairs are pointless. They conclude that the chairs handle one specific bias and that a famous contestant is a different problem needing a different answer, and that the show should stop claiming the audition was blind when it was not.

What would count against this

1. **Partial blinding being worse than none.** A defensible position: it produces the *appearance* of rigour while the main bias runs free, and a reader may discount less than they should. **The counter is that the ceiling is printed on every headline**, which is a real answer only if readers read headlines.
2. **The run effect being the smaller problem.** Almost certainly true. The mechanisms address the bias that could be addressed, which is not the same as addressing the bias that matters.
3. **A second coder being available and not used.** The project decided in July that one coder is the design and the runs are pilots. **That is a resource decision recorded as such**, not a methodological argument, and it should not be dressed as one.

Handing off

Episode 5 builds the control, which is the part Seasons 2 to 4 never had, and the objection to it is strong enough that the episode spends most of its time there.

Questions

1. Design a blinding scheme that would hide the *content*. Is the result still the same study?
2. Is partial blinding worse than none? Argue it properly, then decide.
3. **The empty cell.** Name one thing that would fill it, and say what it costs.

Collecting episode 4

Order blinding and a sealed key, both enforced by tooling, and an empty cell where the main bias sits. Which brings us to the thing that makes the numbers mean anything at all.

The question a control answers

Suppose the canonical gospels interlock at some rate. Suppose they collide at some rate. **Those numbers are meaningless in isolation** and it takes a moment to see why, because they feel like facts.

A rate of one hard collision per thousand words sounds low. Low compared to what? Compared to a transcript? Compared to four independent police statements? Compared to a novel? **You have no idea whether it is low, and neither does anybody else**, until you run the same procedure on texts whose character is not in dispute.

And notice this cuts at the sceptic exactly as hard as at the defender, which is what makes it a control rather than a debating point.

The sceptic's version of the same error is "the gospels contradict each other, therefore they are unreliable." **Compared to what?** Four independent accounts of a real event contradict each other constantly. Police statements from the same incident disagree about the number of assailants, the colour of a car, the order of events. **Contradiction is what independent testimony looks like**, and a set of accounts with no contradictions in it is evidence of collusion rather than truth.

So both sides are quoting a raw count and calling it a verdict. **The defender says the collisions are few and the sceptic says they exist**, and neither has a baseline, and the argument cannot terminate because there is nothing in it that could.

The control, and it is a good one

Six apocryphal texts. The Gospel of Peter, the Protevangelium of James, the Infancy Gospel of Thomas, the Acts of Pilate, and two more.

These are the right control for a specific reason. **They are not a random different genre**. They are Christian documents, about Jesus, written by people with the same commitments, in the same tradition, narrating overlapping events. The Gospel of Peter has a passion and an empty tomb. The infancy gospels have a childhood.

So the comparison is not "sacred text against secular text" or "ancient against modern". It is **partisan documents about Jesus against other partisan documents about Jesus**, which is exactly the objection Season 4 could not answer, made into a measurement.

The objection, which is serious and does not go away

The apocrypha are **later**, and they **had the canonical gospels in front of them**.

That is not a quibble. If the Gospel of Peter is working from Mark, then its collisions with Mark are one text reworking another rather than two independent traditions diverging.

The control is not a clean sample of "what legend looks like". It is a sample of what legend looks like **when it is deriving from the thing you are comparing it to.**

The study records this rather than answering it, and it records a specific instance: the two infancy gospels are tagged as independent by the frozen rulebook and are **probably literarily related.**

What the study does with an objection it cannot answer

It states which way the objection cuts, and notices it cuts both ways.

If the infancy texts are related, their collisions are one author reworking another, not independent invention. That weakens "independent legends collide". **And two related texts that still contradict each other on whether a teacher dies is arguably worse for their reliability rather than better.**

Recorded, not resolved. That phrase appears in the results file and it is the correct posture. An objection you cannot answer gets written down at full strength, with its consequences in both directions, and left standing.

The device: what makes a control good, applied here

PROPERTY A CONTROL NEEDS	DO THE APOCRYPHA HAVE IT
same subject matter	yes , Jesus, overlapping episodes
same commitments	yes , written by believers
same rough genre	mostly , narrative gospel form
independent of the test corpus	no , and this is the wound
character not in dispute	yes , nobody defends the talking cross

Four of five, and the missing one is the important one. That is a good control, honestly described, which is a rarer thing than a perfect one.

The clinical anchor

A negative control on an assay, and what a missing one costs.

You run a test and it comes back positive. Without a negative control you do not know whether the reagent produces that signal on plain buffer. **The result is uninterpretable, not weak.** There is no partial credit.

And the specific failure here is the one that matches the apocrypha problem: **a negative control contaminated with the analyte.** It is still a control, it still catches gross problems, and it will systematically understate your effect because the thing you are controlling with has some of the thing in it.

The character anchor

A mockbuster.

Transformers exists because *Transformers* exists. Watching them side by side tells you a great deal about what a competent film does, because the cheap one is copying the same brief with none of the craft.

And it tells you nothing about how two independent studios would have handled robots, because there is only one brief and one of them is chasing it. That is the shape of the objection and it is why the control is good for embellishment and weak for independence.

What would count against this

1. **The control being contaminated, as above.** The study's own strongest critique of itself in this layer, and it is not resolved.
2. **Six texts being too few.** They are, and the corpora on both sides are small. Episode 8 shows how small, when four units out of twenty-two turn out to move a verdict on their own.
3. **The apocrypha being selected.** They were not chosen at random from the apocryphal literature. They were chosen because they narrate overlapping episodes, which is necessary for the comparison and **is still a selection made by the person running the study.**

Handing off

Episode 6 runs it, and the result splits cleanly into a half that is trustworthy and a half that the person who produced it tells you not to trust.

Questions

1. Name a control for the Gospels that is **not** derived from them. What do you lose?
2. A contaminated negative control understates the effect. **Which direction does that push this study's conclusion, and does that help or hurt the hypothesis?**
3. **Before episode 6:** predict which half of the result will be the weak one.

Collecting episode 5

A control that is four-fifths right, with the missing fifth being the one that matters, recorded and not resolved. Thirty-nine coded items across five pericopes.

What came out

The prediction had five parts and four of them held.

Direction holds. Within each matched comparison the canon interlocks more than the apocrypha. On the empty tomb it is 0.25 against 0.00. On the feeding of the five thousand against the infancy miracles it is 0.67 against 0.00.

Hard collisions hold. The apocrypha collide harder, at 0.95 per thousand words against the canon's 0.22.

Notice what that second number admits. **The canon is not clean.** Mark's ending, where the women say nothing to anyone, is logged as a hard collision against the other accounts, by the person hoping the canon would win. It is in the numerator where it hurts.

That single logging decision is worth more than the paragraph around it, and here is the test for whether it was honest. **The rescue was available.** Mark's ending is textually disputed, the manuscript tradition is genuinely messy there, and a coder who wanted to could have excluded the item on those grounds and written a defensible sentence explaining why.

He did not, and the item sits in the canon's own collision count. That is what it looks like when a rule written in advance binds somebody who would rather it did not, and it is the only reason any of the other numbers in this table are worth reading.

Embellishment holds, and it is the strongest thing in the run. One embellishment in the canon, which is Matthew's raised saints coming out of the tombs. Eight in the apocrypha. Adjusted for length that is **0.22 per thousand against 2.54, about eleven times.**

Independence holds. All five canonical fits are between texts that do not depend on each other, mostly John against the **Synoptics**.

That word needs unpacking, because the whole prediction rests on it. Matthew, Mark and Luke are called the Synoptic gospels, from a Greek word meaning **seen together**, because they can be laid out in three parallel columns and read across. They share wording, order and whole episodes, and the standard explanation is that **Matthew and Luke both used Mark.**

John is the odd one out. Different structure, different episodes, largely different material, and no accepted account has him copying the other three.

So a fit between two Synoptics is nearly worthless as evidence, because copying explains it. **A fit between John and a Synoptic cannot be explained that way.** All five canonical fits are of the second kind, so the interlock advantage is not an artifact of copying, which was the obvious deflating explanation and it is ruled out.

The prediction that failed

Magnitude fails on the canon's own side. The floor was set at 0.60 before scoring. The canon comes in at **0.50**, and only two of its four pericopes clear the floor.

And the reason is worth sitting with, because it is not a technical failure. **The passion narratives are collision-dense enough to drag the canon's own interlock rate below the floor its defender set for it.** The empty tomb scores 0.25. That is the single most theologically loaded pericope in the corpus and it is the worst performer on the canonical side.

There are two readings of that and the episode should give you both rather than pick.

The uncomfortable one. The instrument does worst exactly where the tradition needs it most, and a defender would have to explain why the central claim is supported by the shakiest pericope in the corpus.

The interesting one. Collision density is what you would expect from several independent people describing a chaotic morning nobody understood at the time. The accounts of a car crash disagree more than the accounts of a dinner party, and the disagreement is not evidence that the crash was invented.

Both readings are available from the same number, and nothing in this run chooses between them, which is why the episode reports the number and stops.

The prediction that broke is the one flagged softest in advance, which is either good calibration or a sign that the soft flag was doing the work of a hedge. **Episode 7 is the author deciding which.**

The device: the two halves of one result

FINDING	DEPENDS ON THE CODER'S JUDGEMENT?	TRUST
the apocrypha embellish 11× more	no , these are facts on the page	high
the apocrypha hard-collide more	mostly no	high
the canon interlocks more	yes, entirely	low
the canon's ratio is 0.50	yes	low

A schoolmaster dies in one infancy gospel and merely blushes in the other. A walking cross comes out of the tomb and speaks, in one document and nowhere else. The boy who withers a playmate kills him in one version and heals him in another.

Those are not interpretations. You do not need to know the hypothesis or trust the coder to check any of them, and that is what makes the embellishment finding the solid part of the run.

The clinical anchor

Inter-observer variability, and which findings survive it.

Two radiologists agree almost perfectly on whether there is a fracture. They agree much less on whether a lung field shows early consolidation. **Same imaging, same training, different amounts of judgement in the call.**

So a study reporting both should not report one confidence for both. The fracture count is nearly a fact and the consolidation count is a measurement of the observers as much as the patients, and a paper that averages them into a single result has hidden its own weakest number inside its strongest.

The character anchor

A fishing story growing between two tellings.

The fish gets bigger. That is the recognisable shape of a story being retold by people invested in it, and everybody knows it on sight, which is exactly why it works as a measure.

Now notice what the measurement actually is. Not whether the first telling was true. Only whether the second telling grew, and by how much, against a baseline of somebody retelling the same event without growth. The canon grows once in four and a half thousand words. The apocrypha grow eight times in three thousand.

What would count against this

1. **Thirty-nine items.** It is a very small run, and episode 8 measures exactly how fragile a run this size is.
2. **The 11× being about genre rather than truthfulness.** The infancy gospels are arguably a devotional genre where embellishment is the point, not a failure. If so, the measure is detecting genre and calling it legend.
3. **The empty tomb scoring worst.** The pericope the whole tradition cares most about is the one where the canon performs worst on its own instrument, and **that is either honest reporting or a sign the instrument mismeasures exactly where it matters.**

Handing off

Episode 7 is the author telling you which half of his own result to disbelieve, and it is the most useful thing in the layer.

Questions

1. Mark's ending is logged as a hard collision against the coder's own interest. **What would you have done, honestly?**
2. Is 11× a fact about legend, or about genre? Say what would separate them.
3. **Before episode 7:** which number in the table above would you attack first if you wanted to destroy this study?

Collecting episode 6

Four predictions held, magnitude failed on the canon's own side at 0.50 against a 0.60 floor, and the run splits into a half that is checkable and a half that is not.

The sentence this episode exists for

From the results file, written by the person who produced the number:

The fit-ratio half is the least trustworthy thing here, and I produced it.

Then the reason, which is the part worth memorising.

The fits logged for the canon are the famous ones. Bethsaida and Philip. Passover and the green grass. John's "we". These are the undesigned coincidences that appear in every book on the subject, that the coder has read about, and that he went looking for because he already knew they were there.

Read that list again and ask a question the apologetic literature never asks. **How many gaps were checked to find those three?**

Nobody knows, because the checking happened across two hundred years, distributed across dozens of readers, and only the successes were written down. Paley published what he found. He did not publish the afternoons that produced nothing, and neither did anybody after him.

So the canonical fit count is not the output of a search. **It is the output of two centuries of searching, filtered for hits**, being re-logged in an afternoon by somebody who inherited the filtered list. And it is being compared against a corpus nobody has ever searched even once.

Why that is fatal to the number and not to the study

Think about where selection can enter the interlock ratio. The ratio is fits divided by fits plus collisions. **Collisions announce themselves.** Two texts say incompatible things and you notice, whether or not you were looking.

Fits do not announce themselves. A fit is a connection between two passages that do not reference each other, and you only see it if you go looking in the right place. A coder who has read the apologetic literature knows exactly where to look in the canon, and has read nothing at all about undesigned coincidences in the Gospel of Peter.

So the numerator of the canon's ratio is precisely where a knowing coder's thumb rests, and there was only my thumb. A blind coder scanning every gap without knowing the literature would very likely log fewer, and the canon ratio would fall.

The asymmetry, stated as a mechanism

	CANON	APOCRYPHA
coder has read a literature listing the fits	yes, extensive	no, none exists
effort spent looking for fits	high, and directed	low, and undirected
collisions found regardless of effort	yes	yes
net effect on the ratio	inflated	not inflated

The bias does not need anybody to log a false fit. **Every fit logged can be genuine and the number still comes out wrong**, because the search was not equally thorough on both sides, and an unequal search produces an unequal count from identical texts.

That is a cleaner statement of the problem than "bias", and it is checkable in principle: give somebody the apocrypha and the apologetic method and see how many undesigned coincidences they find. **Nobody has done that**, and until somebody does, the canon's ratio is an upper bound on itself.

The one-line honest reading

The apocrypha demonstrably embellish and hard-contradict far more than the canon, which is robust. The canon appears to interlock more, which is real in direction and inflated in degree by who did the coding.

Two clauses, two different confidences, in one sentence, and neither one is thrown away. **That sentence is what this whole series has been trying to teach how to write.**

Look at what it refuses to do, because the refusals are the craft.

It does not drop the weak half, which would be cleaner and would misrepresent the run. It does not report both halves at one confidence, which would be easier and would hide the strong finding inside the weak one. It does not hedge the strong half to make the sentence feel balanced, which is the commonest failure and the most respectable-looking.

"Real in direction, inflated in degree" is a specific claim about a specific number, and a hostile reader can act on it. That is the test for whether a caveat is doing work: **can somebody who distrusts you use it?** A caveat they cannot act on is decoration.

The clinical anchor

A radiologist who knows the clinical question.

Told "query fracture", they find the fracture. The fracture is really there. And the pneumothorax at the top of the film, which nobody asked about, gets missed at a measurably higher rate.

Nothing false was reported. The report is accurate about everything it mentions. The search was directed, and a directed search returns a biased sample of what was there, which is a different failure from misreading and is invisible in the report itself.

The character anchor

A highlight reel.

Every clip is real. Nothing is staged, nothing is edited to deceive, and the player genuinely did all of it. And you cannot learn anything about how good they are, because you are watching a sample selected by somebody who was looking for good moments.

The fix is not honesty. The person cutting the reel was honest. The fix is a denominator: how many minutes did they play.

What would count against this

1. **This being excessive self-flagellation.** The fits are real, they are independent, and P4 confirmed they are not a copying artifact. **A reader might reasonably weight them higher than the author does.** That reading is available and it is not obviously wrong.
2. **The bias being measurable and not measured.** The check exists and is describable in one sentence. It was not run. **Describing a check you did not run is cheaper than running it,** and this episode should be read with that in mind.
3. **Direction surviving anyway.** Even a heavily discounted canon ratio beats 0.00, because the apocrypha logged **no fits at all.** Zero is hard to reach by thumb.

Handing off

Episode 8 runs a second, tighter control eight months later, and two more falsifiers fire, one of which tells you how much of this depends on four sentences.

Questions

1. Design the experiment that would measure the search asymmetry. What does it cost to run?
2. The apocrypha logged zero fits. **Is zero evidence, or is it the absence of a literature?**
3. Rewrite the one-line honest reading as a headline a publisher would accept. **What did you have to lose?**

Collecting episode 7

The interlock ratio's numerator is where the coder's thumb rests, admitted by the coder, with the mechanism named as an unequal search rather than a dishonest one.

The second control, run much later

A new run, scored in August against a rulebook frozen on 5 August, with the scoring procedure committed the same morning **before any unit was coded**. It measures **restraint**: not whether texts collide, but whether they decline to embellish where a legend would.

Twenty-two units. And this time three falsifiers were registered, and the third one is a technique worth the whole episode.

H3 fires, and it is the same failure as run 1

The gap between the two sets is 0.225, against a threshold of 0.25 fixed before any text was read.

The direction is right and the size is not. And the pre-registration says exactly what that means, in words written in advance:

direction without magnitude, which is what failed P5 in run 1 and it fails here too.

Two independent runs, eight months apart, on different units, failing in the same specific way. That is worth more than either failure alone, because a repeated failure with a shared shape is a finding about the instrument rather than about a run.

H5 fires, and this is the technique

Leave-one-out. Take your twenty-two units. Drop one. Recompute the headline. Put it back, drop the next, recompute. Twenty-two times.

The question it answers is brutally simple: **does any single item, on its own, change the verdict?**

Four of the twenty-two do.

Dropping the unit for Mark's longer ending takes the gap from 0.225 to 0.259 and H3 stops firing. So does dropping Emmaus. So does dropping either of the two least developed units in the Gospel of Peter, from the other side.

So the verdict of this run rests on **four sentences of text**, and if any one of them had been coded differently, or excluded on a defensible technicality, the round would have reported the opposite.

Notice that two of the four sit on each side. That detail matters and it is easy to skim past.

If all four fragile units were canonical, you would suspect the coder of leaning on the canon. If all four were apocryphal, the same suspicion the other way. **Two and two is what fragility looks like when it comes from sample size rather than from a thumb**, and it is the only reassuring thing in the episode.

What it is not is a defence of the result. A verdict that four sentences can flip is a verdict about twenty-two units, and twenty-two units cannot carry a claim about a literature. The right conclusion is not "the finding is wrong", it is "**this run was never large enough to have a finding**", which the pre-registration could have worked out in advance and did not.

Why running it first is the whole point

The pre-registration required leave-one-out to be run **before the headline ratio was computed**. Not after. Before.

Think about what that ordering buys. Run it afterwards, with the headline in view, and every one of those four units becomes a candidate for re-examination. You would look again at Mark's longer ending, and you would find something arguable, **because there genuinely is something arguable about Mark's longer ending**.

Run it first and the fragility is a property of the design rather than a temptation. You learn that your sample is too small before you learn which way it fell.

The device: what leave-one-out tells you that a p-value cannot

QUESTION	ANSWERED BY
could this have happened by chance	a p-value or a null
is the effect big enough to matter	a pre-set threshold
does the whole result rest on one item	leave-one-out, and nothing else

The third row is invisible to the first two. A result can be significant, clear the magnitude bar, and still be one coding decision away from reversing, and no amount of statistical machinery on the aggregate will tell you.

The clinical anchor

A meta-analysis that hinges on one trial.

Twelve studies pooled, a clean forest plot, a significant summary estimate. Remove the 1997 trial and the effect vanishes.

That is a standard sensitivity analysis and it is standard because it fires often. The pooled number was never a summary of twelve studies. It was one study with eleven pieces of decoration, and the only reason anybody knows is that somebody removed them one at a time.

The character anchor

Jenga, near the end.

The tower is standing. It looks like a tower. And there are four blocks in it such that pulling any one brings it down, while the other forty do nothing.

"It is standing" and "it is stable" are different claims, and the only way to tell them apart is to try removing pieces, which is a thing you have to decide to do while it is still standing.

What would count against this

1. **Twenty-two units being too few for the technique to mean much.** True. With a sample this small, some fragility is arithmetically guaranteed. **The finding is not "this study is unusually fragile", it is "this study is too small", and that was knowable in advance.**
2. **H3's threshold being set too high.** 0.25 was fixed in advance, and a threshold of 0.20 would have passed. **Whether 0.25 was honest is checkable by asking what run 1's failure would have needed,** and it is the same shape.
3. **Two failures being a pattern of one instrument.** Both runs were designed by the same person with the same instincts. **Correlated failures are not independent evidence.**

Handing off

Episode 9 leaves Layer 1a entirely and takes up the question everybody actually argues about, with a docket compiled by people who want the answer to come out badly.

Questions

1. Run leave-one-out mentally on run 1's five pericopes. **Which one carries it?**
2. Why must leave-one-out run before the headline? Name what goes wrong otherwise.
3. **Two runs, same failure shape.** Is that stronger evidence or correlated error? You cannot have both.

Collecting episode 8

Two runs, eight months apart, failing the same way: direction without magnitude. And a leave-one-out check showing the second verdict rests on four sentences.

Layer 1a asked whether the texts interlock. Layer 1b asks the question people actually argue about. **When somebody alleges a contradiction, is it one, and what does the rescue cost?**

The first decision, and it settles the study's honesty

You need a list of alleged contradictions. There are two ways to get one.

You can compile it yourself. You know the texts, you know the hard places, and you would produce a thoughtful list. And every reader who distrusts you would be right to, because you also know which items you can answer.

Or you can take the list from people who want it to hurt.

The docket here is **671 raw allegations**, aggregated from three hostile compilations: the Skeptic's Annotated Bible, Infidels.org, and EvilBible. Not a sample of them. The aggregation, entire.

The docket is hostile-sourced and complete, so the coder cannot skip the hard ones.

They are on the list because somebody who wanted the New Testament to look bad put them there, and removing one requires a stated reason on the record.

What that buys, and what it does not

It buys the one thing selection bias cannot survive: **the person being tested did not choose the test.**

And it costs something real, which is why almost nobody does it.

A hostile list is full of bad entries. Items that misread the verse, items that allege a contradiction where one side is a question, the same complaint four times over. **You have to work all of them**, seriously, and write down the reason each one fails, and that is a great deal of tedious labour producing nothing anybody wants to read.

The temptation is to filter first: drop the obviously bad entries, keep the serious ones, and work those properly. **That is fatal, and the reason is that you would be the one deciding which are obviously bad**, which hands back exactly the control that taking the list was meant to give away.

And it changes what the result is a statement about, in a way worth being precise about. This is a **ceiling on the accusers' case**, never a floor on the truth.

If the strongest available compilation of objections yields two hard problems, you have learned something about that compilation. **You have not learned that there are only two problems in the New Testament**, because a contradiction nobody has noticed is not on anybody's list, and no procedure here would find it.

That asymmetry is permanent and it is worth naming precisely, because it is the kind of thing that gets lost the moment a number is quoted.

A hostile docket bounds the accusers and says nothing about the texts. It is the difference between "your best shot failed" and "there was nothing to hit." Those two sentences are often said as though they were one sentence, and only the first is available here.

The version that would say something about the texts is the denominator study in episode 17, which walks every cross-checkable pair whether or not anybody complained about it. **It has not been built**, and until it is, every number in this layer is a fact about three websites.

The device: who picks the items, and what follows

DOCKET SOURCE	PROTECTS AGAINST	WHAT THE RESULT MEANS
the defender compiles it	nothing	worthless
a neutral party compiles it	most selection	a real estimate, unavailable in practice
hostile compilers	the defender skipping hard items	a ceiling on their case

Compare this to run 1, which is the point of putting the two layers in one season. Run 1's fit ratio failed for exactly the opposite reason: **there, the coder picked the items**. Same person, same project, two months apart, and the difference between a trustworthy number and an untrustworthy one is entirely who wrote the list.

The clinical anchor

An incident reporting system, and what it can never measure.

You can count everything that was filed. You can categorise it, trend it, and present it to the board. And the denominator is not errors, it is **errors somebody chose to write up**, which is a different quantity moving for different reasons.

So a falling incident rate is compatible with improving care and equally compatible with staff losing faith in the form. **The number is real and it is about the reporting, not the harm**, and a service that forgets this congratulates itself on a graph.

The character anchor

Letting the prosecution pick the charges.

A defence that says "try me on whatever you like" is making a much stronger claim than one that says "try me on this". Everybody watching understands the difference immediately, without any legal training.

And it has an exact limit that people forget. Surviving the charges brought says nothing about a crime nobody charged. The acquittal is real and it is bounded by the indictment, which somebody else wrote.

What would count against this

1. **The compilers being bad at their job.** These lists are aggregations of popular objections, not scholarly ones. **The best objections in the academic literature may**

not appear, and the ceiling is a ceiling on the popular case.

2. **671 being inflated by duplication.** It is, heavily, and episode 12 measures exactly how much. That is a finding about the lists rather than a flaw in using them.
3. **Hostile sourcing not preventing hostile gating.** They pick the items. **The coder still decides which ones count**, and episode 15 is a falsifier firing on precisely that.

Handing off

Episode 10 asks what a contradiction actually is, and the answer is twenty-three centuries old and much narrower than anybody uses in practice.

Questions

1. Why is "a ceiling on the accusers' case" weaker than it sounds, and why is it still worth having?
2. Name a way a hostile docket could be *unfair to the accusers*.
3. **Before episode 10:** define a contradiction in one sentence, and then find a counterexample to your own definition.

Collecting episode 9

671 allegations from three hostile compilations, taken entire so the hard ones cannot be skipped, producing a ceiling on the accusers' case rather than a floor on the truth.

Now the definition, because everything downstream is arithmetic on it.

The definition, and it is Aristotle's

Two texts contradict when they assert **P and not-P** about:

- **the same subject,**
- **in the same respect,**
- **at the same time,**

and no licensed reading makes both true.

Those three qualifiers are the whole content, and each one blocks a specific mistake that ordinary argument makes constantly.

Working the three qualifiers

Same subject. "Jesus healed one blind man" and "Jesus healed two blind men" contradict only if they describe the same healing. If they are two occasions, both are true and nothing has happened. So before anything else you have to establish that the two texts are talking about one event, and a surprising number of alleged contradictions die here without any special pleading.

Same respect. This is the one that does the most work and causes the most trouble. "The mustard seed is the smallest of all seeds" and modern botany disagree only if the text is making a botanical claim rather than a proverbial one. **A word used in two senses generates a fake contradiction**, and so does a rhetorical "all" read as a strict quantifier.

Same time. "Circumcision is required" and "circumcision is nothing" are both in the New Testament, and they belong to a stated change of covenant. Before against after is not P and not-P.

And notice that the third qualifier is the one most open to abuse, in a way the first two are not. **Same subject and same respect are questions about the texts. Same time invites you to supply a history.**

"These belong to different periods" is a claim, and it is checkable when the text states the change and unfalsifiable when it does not. The circumcision case is fine because the New Testament argues about the transition explicitly and at length. **An appeal to different periods where no text marks a period is the same move as inventing an event**, and the cost scale in episode 13 prices it accordingly.

Why this is not a defensive trick

Because the same three qualifiers destroy defences as readily as accusations, and the study is about to demonstrate that on the church fathers.

A distinction of respect that nobody can attest is not a distinction, it is a wish. Saying two texts use a word in different senses obliges you to show where the second sense is attested, and if you cannot, you have supplied from outside what the text does not supply, which is the guard from episode 3 pointed at yourself.

The rulebook requires it explicitly: **every appeal to a difference of sense must state the specific distinction and cite where the alternative sense is attested.** Not "these could mean different things", which is always true of any two words.

The requirement to **cite where the alternative sense is attested** is the whole weight of that rule, and it is worth seeing why it is not a formality.

Any word in any language can be argued into a second sense. Given a motive, you can propose that "day" means an age, that "all" means many, that "brother" means cousin. **Every one of those proposals is possible and some of them are correct**, and the difference between the correct ones and the convenient ones is whether the sense turns up somewhere else, in a text nobody was arguing about.

Attestation is what stops a sense-distinction being free. And a project that requires it of its opponents and not of itself has built a machine for agreeing with itself, which is what episode 15 catches this project doing at a measured rate of 67.6%.

The device: the same allegation through the three qualifiers

ALLEGATION	QUALIFIER THAT DECIDES IT	VERDICT
one blind man against two	same subject	dies if two occasions, and there is no evidence either way
faith against works	same respect	two emphases, not P and not-P
circumcision required against nothing	same time	stated change of covenant
Judas hanged against Judas burst open	all three hold	survives, and it is one of two

The bottom row is the shape of a real one. Same man, same death, same moment, and no attested sense in which hanging is falling headlong.

The clinical anchor

Two results that look contradictory until you check the units.

A calcium of 2.2 and a calcium of 8.8 are the same measurement in different units. A glucose taken fasting and one taken after lunch are not in conflict. A potassium from a haemolysed sample is not disagreeing with the one from a clean draw.

Most apparent conflicts in a set of notes are resolved by asking what was measured, in what units, at what time, which is precisely same subject, same respect, same time. The remainder, after those three questions, is the interesting list, and it is always much shorter than it looked.

The character anchor

Two people arguing about whether a film is long.

One saw the theatrical cut and one saw the extended edition. They are not disagreeing about anything. They have been arguing for twenty minutes and both are correct, and the argument only ends when somebody asks which version.

Almost every argument you have ever watched fail to resolve has this shape, and the three qualifiers are a procedure for finding out in the first minute rather than the twentieth.

What would count against this

1. **The definition being too strict.** Under it, a great many real tensions in the texts are not contradictions, and somebody could reasonably say the strict definition protects the texts by construction. **The answer is that it is Aristotle's and not the project's**, and that it is applied to defences too, which episode 18 shows costs the defenders more than it costs the accusers.
2. **"Licensed reading" being where all the judgement hides.** It is. That is what the nine gates in episode 11 exist to constrain and what fires a falsifier in episode 15.
3. **Same-respect being unfalsifiable in the hard cases.** Sometimes nobody can settle whether a word is used in two senses. **Those items should be contested, not dissolved**, and whether they were is the thing episode 15 measures.

Handing off

Episode 11 turns the definition into nine gates, each one a specific way an allegation can fail, applied in a fixed order.

Questions

1. Take three contradictions you have heard alleged and run them through the three qualifiers. **How many survive?**
2. "A distinction of respect nobody can attest is a wish." **Apply that to an argument you find persuasive.**
3. **Before episode 11:** if you had to write gates from this definition, how many would you need, and what would the first one be?

Collecting episode 10

Same subject, same respect, same time, and no licensed reading that makes both true. A definition twenty-three centuries old, applied symmetrically.

An allegation now has to pass nine gates before it counts as a real clash. Each gate is a specific way the definition can fail, and they are applied in order.

The nine

A0, not an assertion. One side is a question, a command, a wish, a curse, a hypothetical, or the reported speech of a character the text does not endorse. Job's friends say a great many things the book of Job exists to contradict.

A1, not the same subject. Different events, persons, places or occasions.

A2, not the same respect. Different sense of a shared word, a rhetorical "all" read as strict, figurative against literal, legal against biological, role against person.

A3, not the same time. Different period or dispensation, a stated change of covenant, before against after.

A4, omission only. One text is silent where the other speaks. **Silence is not denial, and naming one of two is not asserting only one.**

A5, tension not negation. Two emphases in tension without P and not-P. Love against wrath, faith against works, sovereignty against freedom.

A6, text against doctrine. The clash is with a later dogma such as inerrancy or a systematic theology, not between two texts. **This gate is important and counterintuitive:** an item that embarrasses a doctrine but leaves both texts true is not a contradiction between texts, and it is filed as such.

A7, designed juxtaposition. The tension is an intentional literary pairing. Proverbs 26:4 and 26:5 tell you not to answer a fool and then to answer a fool, on consecutive lines, on purpose.

A8, reference failure. The cited verses do not say what the entry claims, or the citation is broken.

A8 is the gate that grades the accusers

Most of the gates are judgements. **A8 is not.** You open the verse and either it says what the entry claims or it does not.

That makes it the one gate whose count is an objective finding about the compilations rather than about the coder, and the study reports it as such: **five entries misdescribe their own citations, four put a question on one side.**

Those are small numbers and they are worth having precisely because they are checkable by anybody in ten minutes, which almost nothing else in this season is.

And there is a second reason to care about A8 beyond its size. **It tells you about the process that produced the list.**

An entry that misdescribes its own citation was written by somebody who did not open the verse. Not somebody who read it and disagreed about its meaning, which would be an ordinary dispute. **Somebody who copied a claim from another list.**

So A8 is a small window onto how these compilations are actually assembled, and what it shows is the same thing episode 12 measures at scale: **a literature propagating by copying rather than by checking.** Nine items is not many. The mechanism it exposes is what produces the other two hundred.

The device: the gates ordered by how much judgement they take

GATE	JUDGEMENT REQUIRED	WHO CAN CHECK IT
A8 reference failure	none	anybody with a Bible
A0 not an assertion	very little	anybody
A7 designed juxtaposition	some	a reader of the genre
A1 same subject	real	contested case by case
A3 same time	real	contested case by case
A2 same respect	maximal	almost nobody
A5 tension not negation	maximal	almost nobody

The two bottom rows are flagged in the rulebook, in advance, as the danger. The pre-registration says so plainly: *Tier A is where my judgement concentrates. A2 and A5 are where almost anything can be called a difference of sense.*

And having named them as the danger, it requires their counts reported on their own line so a reader can discount them, **and sets a falsifier on their share.** Episode 15 is that falsifier firing.

The clinical anchor

Exclusion criteria applied at screening, in a fixed order.

Age, then renal function, then the drug interaction, then consent. The order is fixed so the reason for exclusion is recorded consistently, and so that two screeners exclude the same patient for the same stated reason.

And the criteria that require judgement are the ones you audit. Nobody re-checks the date of birth. Everybody re-checks "unable to comply with follow-up", because that is the criterion where a screener's convenience can hide.

The character anchor

A bouncer with a checklist, and the one line that is not on it.

ID, dress code, capacity, guest list. Four objective gates, and everybody in the queue can see whether they passed each one.

Then there is "not tonight, mate." That is the gate doing the real work, it is unappealable, and the person applying it could not fully explain it if asked. Every system of rules has one, and the honest thing is not to pretend otherwise but to **count how often it gets used.**

What would count against this

1. **Nine gates being too many.** More gates means more ways to dismiss an allegation, and a defender who wanted to dissolve everything would design exactly this. **The counter is that they come from the definition rather than from convenience**, and that A6 in particular costs the defence its doctrine.
2. **The order mattering.** An item that could fail at A1 or A2 gets recorded at whichever comes first, and the ordering was chosen by the coder.
3. **A6 being a dodge.** Filing "this contradicts inerrancy" as not-a-contradiction is defensible and it is also exactly what somebody protecting a doctrine would do. **The study's answer is that it does not protect the doctrine**, and episode 17 says so at the top.

Handing off

Episode 12 does the counting before the judging, and finds that a quarter of the accusers' list is the same complaint written down again.

Questions

1. Which gate would you attack first, and would that attack help the accusers or the defenders?
2. **A8 is checkable by anybody.** Find one. Genuinely, go and check three entries from any list of Bible contradictions.
3. **Before episode 12:** guess what fraction of 210 allegations are duplicates. Write the number down.

Collecting episode 11

Nine gates, ordered, with the two most judgement-heavy flagged in advance and a falsifier set on their share. Before any of them are applied there is arithmetic to do, and it is the most checkable finding in the layer.

The funnel

671 raw allegations arrive from the three compilations. Then:

671 to 210. Most entries allege a conflict between a New Testament text and something else, usually the Old Testament, or a fact about the world, or a doctrine. Only **210** allege a contradiction between two New Testament texts, which is the question this study asked. The rest are a different study and are set aside rather than dismissed.

210 to 198. Exact duplicates removed, where the same entry appears in two compilations word for word.

198 to 153. Near-duplicates removed, where the same underlying proposition-pair is alleged in different words. This is the step requiring judgement, and every merge is logged.

The finding, and it is about the accusers rather than the texts

57 of 210, which is 27%, are the same complaint counted again.

More than a quarter of the objections in the aggregated hostile literature are restatements. Not weak objections, not wrong ones. **The same objection, arriving repeatedly, counted each time as though it were new.**

That is a fact about how these lists are built, and it is checkable by anybody who downloads the same data. It does not require trusting the coder, the gates, or anything else in this season.

Why it matters more than it looks

Because the impressive thing about these lists is their **length**. The rhetorical force of "671 contradictions" is entirely in the number, and nobody quoting it has ever divided it by anything.

The funnel shows what the number is made of. Two thirds are not NT-internal at all. A quarter of what remains is repetition. **And none of that involves deciding whether a single allegation is any good**, which has not started yet.

That last sentence is the one to hold, because it is what makes this episode different from everything after it.

Nothing so far requires you to trust the coder about anything. The category test is close to mechanical. The exact dedupe is string matching. **A hostile reader can download the same three lists and reproduce 671 to 210 to 198 without agreeing with a single opinion in this season.**

Which is why the funnel is reported before the gates, in that order, in the results file. **The arithmetic that survives distrust goes first**, and the judgement that does not goes after, so a reader who stops halfway has still been told something true.

Granularity, and the honest problem inside it

The merge from 198 to 153 requires deciding when two allegations are the same allegation, and that decision has a direction.

Merge aggressively and the list shrinks, which flatters the defence. **Merge conservatively and it stays long**, which flatters the accusers. There is no neutral setting, and the coder benefits from one of them.

The study's response is to log every merge in a separate file, so the decisions are inspectable individually rather than presented as a total. That is the right response and it is not a solution: **a reader who does not open the log is trusting the coder**, and almost no reader opens the log.

It is worth being honest about what a log actually accomplishes, because "everything is documented" is a phrase that gets used as though it settled something.

A log does not make a decision correct. **It makes it attackable**, which is a smaller thing and the only thing available. Somebody who thinks the merges were too aggressive can now point at merge number 31 and argue about that specific pair, rather than making a general complaint that cannot be answered.

The value of transparency is not that it prevents error. It is that it converts a dispute about character into a dispute about item 31, and disputes about item 31 can terminate.

The device: the funnel, and where judgement enters

STEP	FROM	TO	JUDGEMENT
NT-internal only	671	210	almost none , a category test
exact dedupe	210	198	none , string matching
near dedupe	198	153	real, and logged item by item
the nine gates	153	51	maximal, and a falsifier fires here

Read the judgement column downward. The study becomes less trustworthy as you go down it, and it reports the numbers in that order so a reader can stop wherever they lose confidence. **Everything above the last row is close to arithmetic.**

The clinical anchor

The same patient counted three times in an audit.

One admission, three specialty referrals, three entries in the dataset. The audit reports three cases of delayed review, and the board hears a number three times larger than the event.

Nobody falsified anything. Each entry is a real referral with a real delay. The unit of counting was never defined, and a number without a defined unit is not a measurement, however carefully it was collected.

The character anchor

A Steam review bombing.

Ten thousand negative reviews arrive in a day. Some are from people who played the game. Many are copies of the same sentence, posted by people repeating a complaint they read somewhere else.

The complaint underneath may be entirely fair. The count is not evidence about the game, it is evidence about how many people repeated one thing, and the two get conflated because a big number feels like a lot of independent judgements.

What would count against this

1. **Repetition being legitimate.** If three compilers independently noticed the same problem, that is convergence rather than padding, and merging them hides it. **A fair point**, and it argues for reporting both numbers, which the study does.
2. **The near-dedupe being self-serving.** It shortens the accusers' list by a step the defender controls. The log exists precisely because this objection is correct in principle.
3. **671 to 210 discarding the strongest items.** The NT-versus-OT stratum is 190 items and is **unscored**. That is not a small omission and the study says so.

Handing off

Episode 13 is the cost scale, which is what happens to the 51 items that survive the gates, and it is the instrument that eventually costs this project a doctrine.

Questions

1. 27% duplication. **Does that make the list weaker, or does it show three compilers converging?** Both readings are available.
2. Design a dedupe rule that a hostile reader would accept in advance.
3. **Before episode 13:** an allegation survives the gates. What are the possible verdicts, and how many do you need?

Collecting episode 12

671 to 210 to 153, with 27% of the accusers' list being the same complaint recounted. Judgement enters at the near-dedupe and dominates at the gates.

Now the 51 items that survive. They are real clashes between two assertions, and the question is no longer whether there is a problem. It is **what fixing it costs**.

The scale

Five levels, and the point of it is that "is it a contradiction" is the wrong question and "what do you have to believe to make it go away" is the right one.

Cost 0, dissolved. The plain text shows no contradiction. Mentioning one of two is not denying two.

Cost 1, licensed. A rescue exists, is directly cited, and is uncontested. A lexeme genuinely has that range. A custom is genuinely attested.

Cost 2, contested. A real rescue exists but rests on a debated reading or a reconstruction. A live difficulty, not a knockout.

Cost 3, ad hoc. The reconciliation needs an event no text states, or an emendation of the text. Ehrman's phrase for this is **writing a new gospel**, and it is the right phrase.

Cost 4, hard. No non-ad-hoc rescue on any channel. A genuine contradiction.

The Ehrman guard, stated once and applied to everybody

A rescue counts as cost 1 or below only if it cites its licence: a lexeme, a verse, a custom, a source. **If the only harmony available is invented events or emended text, it is cost 3, no matter how plausible it sounds.**

The last clause is the one that bites. Plausibility is not licence. A reconstruction can be entirely reasonable, the sort of thing that probably did happen, and still be cost 3, **because the question is not whether it is true but whether a text says it**.

That separation is unfamiliar enough to be worth a second pass, because most argument runs the two together.

Suppose two accounts of a journey disagree about where somebody stopped, and you propose they stopped twice. **That may well be what happened.** People do stop twice. The proposal is not far-fetched and a historian might accept it without blinking.

It is still cost 3, because no text says they stopped twice, and the proposal was generated by the need to resolve the disagreement rather than by any evidence. The scale is not measuring whether your story is true. **It is measuring where your story came from**, and a story generated by the problem it solves has no independent standing however plausible it sounds.

And the guard runs both ways, which is the part that separates this from apologetics with extra steps. It blocks the defender inventing a harmony **and it blocks the coder**

dissolving an item because it feels fine. Both are supplying from outside what the text does not supply.

Why the scale is better than a verdict

A binary verdict on a contradiction produces an argument that never ends, because both sides can always say the word. The scale converts the argument into a question with an answer: **name the channel your rescue runs on, and cite it.**

That reframing does something specific to the debate. It stops being about whether the Bible has contradictions and starts being about the **price list**, and a price list can be audited item by item by somebody who disagrees with you about everything else.

And it changes what winning looks like, which is the deeper effect.

Under a binary, both sides are trying to establish a fact about the Bible, and neither can, so the argument runs forever on the same examples. Under a scale, **both sides are trying to establish a price for a specific item**, and prices can be compared, aggregated, and revised one at a time.

A defender who accepts the scale has conceded that some items cost something. A sceptic who accepts it has conceded that most cost very little. **Neither concession is available in the binary version**, which is why the binary version is the one that never moves.

The device: one item at four prices

PROPOSED RESCUE FOR THE SAME CLASH	CHANNEL CITED	COST
the two texts describe different occasions, and a text says so	verse	1
the word has a second attested sense, here is the lexicon	lexeme	1
the word could mean something else	nothing	3
there was an intervening event nobody records	nothing	3
no rescue on any channel		4

Rows three and four are the ones that sound fine out loud. They are the ones you hear in sermons and in debates, and they are priced at 3 because a sentence that could be said about any two texts whatsoever has no evidential content.

The clinical anchor

Grading an adverse event rather than declaring one.

Nobody asks whether the patient had a side effect. They grade it: mild, moderate, severe, life-threatening, fatal. Five levels, defined in advance, applied by somebody who did not choose the drug.

The grading is what makes the data poolable. A trial reporting "some adverse events" tells you nothing, and a trial reporting grade 3 and above tells you what the drug costs. And the discipline that makes it work is that the grade is fixed by criteria rather than by how the reporter feels about the drug.

The character anchor

A repair quote.

You do not ask the mechanic whether the car has a problem. You ask what it costs to fix, and the answer sorts itself into a scale everybody understands: a bulb, a service item, a part that needs ordering, an engine rebuild, **write it off**.

And the useful conversation is always about the itemised list, not about whether the car is "fine". Somebody who says the car is fine and somebody who says it is scrap can both look at the same quote and find they are arguing about one line on it.

What would count against this

1. **The scale rewarding ingenuity.** A clever defender can always find *some* channel, so a well-read apologist scores lower costs than a less-read one on identical texts. **The scale measures the state of the defence literature as much as the texts.**
2. **Cost 3 being the interesting category and the fuzziest.** The line between a reconstruction and a licensed reading is exactly where the argument lives.
3. **The guard being applied by one person.** Symmetric in principle, and applied by somebody with a view. **Episode 15 measures how much that mattered**, and the answer is: enough to fire a falsifier.

Hanging off

Episode 14 is the tally, and the headline number is smaller than anybody on either side would predict.

Questions

1. Take a harmonisation you have heard. **Price it.** Which channel does it cite?
2. Why is "plausible" priced the same as "invented"? Argue against that choice.
3. **Before episode 14:** 153 distinct allegations. Guess how many survive the gates, and how many land at cost 3 or 4.

Collecting episode 13

A five-level price list and a guard that binds defender and coder identically. 153 distinct allegations went in.

The tally

51 accepted as genuine clashes between two assertions. 102 fail a gate.

Of those 51, priced on the scale:

Two at cost 3. How Judas died, and who accused Peter.

Zero at cost 4.

Zero at cost 4 touching the core, where the core is the claims the case actually rests on: crucified under Pilate, buried, the tomb found open.

What the two are

Judas. Matthew has him hang himself. Acts has him fall headlong and burst open. Same man, same death, same moment, and the reconciliations available all require an event neither text states, which is the definition of cost 3.

Peter's accusers. The gospels differ on who challenged Peter in the courtyard, and the second challenger is **gender-marked in the Greek**, which closes the usual harmonisation. Mark has a definite feminine article, Matthew has a feminine pronoun, Luke has a masculine one and Peter answers with a masculine vocative.

Why gender-marking closes it is worth spelling out, because the harmonisation it kills is normally a strong one.

The standard move is that more people spoke than any single gospel records, so Mark reports one challenger and Luke another and both are true. **That works whenever the texts merely name different people**, and it is not special pleading: a crowd accuses somebody and different witnesses remember different voices.

It stops working when the texts specify the same slot in the sequence and mark it differently. This is the second denial in all three, and Greek marks the challenger's sex in the grammar rather than leaving it open. To keep the harmonisation you now need a second challenge, of a different sex, in the same slot, that nobody narrates. **Which is an invented event, which is cost 3.**

Read the zero carefully, because it is doing less than it looks

Zero at cost 4 does not mean there are no contradictions. It means that for every accepted item, **some** rescue exists on **some** channel, even if that rescue is expensive.

And cost 3 is not a pass. It is the study saying: to make this go away you must believe something no text states. That is available and it is not free.

There is also a structural reason zero at cost 4 was always likely, and a study that quotes the number without it is overselling.

Cost 4 requires proving a negative: that no rescue exists on any channel, not merely that none has been proposed. That is a very high bar and it gets harder the more defenders have written, because every century of apologetic literature adds candidate rescues. **An item can move from 4 to 3 forever and never move back.**

So "zero at cost 4" is partly a fact about the texts and partly a fact about how much ingenuity has been spent on them for two thousand years. **The honest reading is that the scale's top level is nearly unreachable by construction,** and the interesting line is therefore cost 3.

The honest sentence is not "the New Testament has two contradictions." It is: **of 671 allegations compiled by people trying to find the worst, 153 are distinct, 51 are real clashes, and two of those cannot be fixed without inventing an event.**

What it costs, said at the top rather than the bottom

Inerrancy.

A cost 3 means one of two texts is wrong about a fact. Judas either hanged or fell headlong, Matthew and Acts say different things, and **the scale does not remove it.** No amount of the rest of this study buys that back.

So a defender who wants to quote the number two has to also quote what the two cost, and a study that produced the number and then defended inerrancy would be contradicting its own output.

The device: what each number is and is not evidence for

NUMBER	EVIDENCE ABOUT	NOT EVIDENCE ABOUT
671	the compilations	the texts
153	the compilations, deduplicated	the texts
51	the texts, filtered by the coder's gates	how many exist
2 at cost 3	the texts, and it costs a doctrine	inerrancy surviving
0 at cost 4	that some rescue always existed	that the rescues are good

The clinical anchor

A screening result that requires confirmation.

A positive screen is not a diagnosis and a negative screen is not an all-clear. It is a number produced by a test with a known sensitivity, on a population with a known prevalence, and it means something only in that frame.

"Two" is a screening result. Two out of an unstated denominator, from a test whose specificity has just been shown to be unreliable, on a docket somebody else compiled. Every one of those qualifiers is in the study and none of them survives being quoted in a thumbnail.

The character anchor

A health inspection score in the window.

Five stars. And the score is against a defined checklist, by an inspector who came on one day, and it means the specific thing the checklist measures.

Nobody reads it that way. It gets read as "this place is good", and the gap between what the number measures and what it is taken to mean is the entire reason the checklist has to be published alongside it.

What would count against this

1. **The denominator, and it is the biggest hole in the season.** Two out of *what*? Against thousands of cross-checkable pairs it is remarkable. Against forty it is bad. **Same numerator, opposite verdicts**, and episode 17 is about this.
2. **51 depending entirely on the gates.** And the gates are about to fail their own falsifier.
3. **The core being defined by the person testing it.** "Crucified under Pilate, buried, tomb open" is a reasonable core and it was chosen by the defender. **A sceptic would define it differently.**

Handing off

Episode 15 is the falsifier firing, on the coder, by name, in the headline, and it is the episode this whole season was built to contain.

Questions

1. Two at cost 3. **Write the honest one-sentence version of that finding**, and then the dishonest one, and notice how small the edit is.
2. Why is zero at cost 4 weaker evidence than it appears?
3. **Before episode 15:** the gates rejected 102 of 153. What would make you distrust those rejections?

Collecting episode 14

51 real clashes, 2 at cost 3, 0 at cost 4, and a doctrine gone. All of it resting on 102 rejections made by one person applying nine gates.

The falsifier, written before any item was scored

F4. If A2 plus A5 rejections exceed **40%** of all rejections, the dissolution is resting on my own sense-distinctions rather than on the texts, and Number A must be reported as unreliable.

A2 is "not the same respect". A5 is "tension, not negation". Episode 11 flagged both as the gates where almost anything can be called a difference of sense, and the rulebook set a ceiling on how much of the work they were allowed to do.

It fired

A2 plus A5 are 67.6% of all rejections, against a ceiling of 40%.

Sixty-two entries dissolved by saying "**not the same respect**". Seven more by saying "**tension, not negation**." Two thirds of the dissolution rests on the two gates the author himself named as the abusable ones.

And so the results file opens like this:

The honest first sentence of this study is not a number about the New Testament. It is that I dissolved 62 entries by saying "not the same respect", and I pre-committed to treating that volume as disqualifying. It is disqualifying.

What that costs, exactly

The dissolution rate is reported as **unreliable**. The number 51 cannot be quoted as an estimate of how many real clashes exist. **The gate calls are not trustworthy in aggregate**, and the person saying so is the person who made them.

Nothing in this season is more important than that paragraph existing, and nothing in this season is easier to have quietly not written.

Think about how invisible the alternative would be. **F4 is a rule the author wrote, in a file the author controls, about a threshold the author chose.** No reviewer set it, no institution enforces it, and nobody outside the project would ever have computed the ratio.

He could have reported 51 accepted items and two hard contradictions, with a limitations paragraph noting that some rejections involved judgement, and every sentence would have been true. **The A2 and A5 counts would simply not have been added up.**

The only thing that made F4 fire was somebody doing the arithmetic on his own gates and then publishing the answer. That is what a falsifier is: not a clever test, just a number you promised in advance to compute and to say out loud whichever way it came.

The survival argument, which you have to judge

Then the file argues that the **hard count** survives, and the argument is structural rather than an assertion, which is the only kind that should be accepted here.

An A2 or A5 error puts an item back in play at cost 1 or 2, and structurally cannot create a cost 3 or 4.

Follow it. A cost 3 means no channel supplies a rescue. An item rejected at A2 is by construction an item where a channel **does** supply one, namely the sense distinction that was used to reject it. So if that rejection was wrong, the item returns as a **contested** clash with a stated sense distinction, which is cost 1 or 2 by definition.

Run as a worst case: **accepted goes from 51 to 120, cost 3 stays at 2, cost 4 stays at 0.**

This is either the best reasoning in six seasons or an escape hatch with good manners, and you have to decide which. The test is whether the argument was available before the falsifier fired, or invented to survive it.

The device: what a fired falsifier does and does not contaminate

OUTPUT	CONTAMINATED BY F4	WHY
the dissolution rate	yes, fatally	it is the thing F4 measures
the number 51	yes	it is one minus the dissolutions
27% duplication	no	arithmetic, before any gate
A8 reference failures	no	checkable by anybody
2 at cost 3	no, and the argument is directional	an A2 error cannot manufacture a 3

The clinical anchor

Auditing your own notes and finding you documented capacity in two-thirds of cases with the same sentence.

The sentence is not false. Each patient probably did have capacity. And a formulation that appears in two thirds of your notes verbatim is not an assessment, it is a habit, and the audit exists to catch the difference.

The uncomfortable part is that you cannot tell from inside which of the two it was, on any individual note, including the ones where you remember thinking about it.

The character anchor

A referee reviewing his own decisions and finding he overturned two thirds of them on one rule.

Nobody suggests he was cheating. And the plain reading is that he was applying that rule far more than the rule can bear, and every individual call felt correct at the time.

What he does next is the whole thing. The version where he publishes the figure, and the version where he notices it and does not, produce identical match reports.

What would count against this

1. **The survival argument being an escape hatch.** The strongest objection to this episode. **The check is timing:** the directional property of A2 is a consequence of the gate definitions written in advance, not a new premise. **A reader who does not accept that should discount the hard count too**, and the study would rather they did that than trust it wrongly.
2. **67.6% being expected rather than alarming.** Most alleged contradictions really are sense confusions, so a high A2 rate may be the correct answer rather than a bias. **Then the 40% ceiling was set wrong**, and setting it wrong in advance is still better than not setting it.
3. **A falsifier firing being treated as a success.** It is reported as one here, and there is a real risk that "we fired our own falsifier" becomes a credibility move rather than a cost.

Handing off

Episode 16 works the two survivors in detail, and one of them was corrected downward, against the story, by the person telling it.

Questions

1. **Decide.** Escape hatch or structural argument? Say what evidence moved you.
2. Design a falsifier for a study of your own that could fire on you. **Would you publish it?**
3. If most alleged contradictions really are sense confusions, what ceiling should F4 have used? Justify it without looking at 67.6%.

Collecting episode 15

F4 fired at 67.6% against a 40% ceiling, the dissolution rate is unreliable, and the hard count survives on a directional argument you were asked to judge.

Two items are left. They are worth working properly, because everything else in this layer is a number and these are texts.

Judas, and the correction that went the wrong way

Matthew has Judas hang himself. Acts has him fall headlong and burst open in a field. Same man, same death.

This item was originally priced at **cost 4**, which is the top of the scale and means no rescue exists on any channel. It was **revised down to 3**.

The revision went against the story. Cost 4 is the stronger position and the sharper finding, and the study gave it up.

The reason is narrow. Cost 4 means **no** rescue exists on any channel. And one does: you can say the rope broke, or the branch gave, and the body fell.

That rescue is worth nothing and it exists. No text states it. It is invented to fit. And the scale has a level for exactly that, which is 3, so the item belongs at 3 and the sharper number was wrong.

The correction was propagated everywhere it appeared, including into the ledger and the evidence file, **and the Bayes calculation gained a cost-3 case chosen before the recomputation was run**, so the correction could not be tuned.

Two words there need unpacking.

Bayes is the arithmetic of updating a belief with evidence. You start with how likely you thought something was, you observe something, and it moves you. The project runs one to ask how much the whole ledger shifts the case, **and it is the part of the study most sensitive to choices**, because you must supply the starting number yourself.

Which is why the ordering matters. Lowering Judas from 4 to 3 changes an input to that calculation. **The value the calculation would use for a cost-3 item was fixed before the recomputation ran**, so it was impossible to pick the value that made the correction painless. You correct the item, then you find out what it costs, in that order, and not the reverse.

Peter's accusers, and why the Greek closes it

The usual harmonisation says there were more challengers than any one gospel reports, so Mark names one and Luke names another and both are true. That is an A4 move, silence is not denial, and it is normally strong.

It fails here on grammatical gender, and this is a tier V finding, checked against the Greek text rather than taken from a commentary.

The second challenger is gender-marked in all three. Mark has a **definite feminine** article with the servant girl. Matthew has a **feminine** other. Luke has a **masculine** other, and Peter answers him with a masculine **vocative**, "Man".

A vocative is simply **the form a word takes when you are addressing somebody directly**. English barely has one, which is why the term is unfamiliar: we say "Man!" the same way we say "a man". Greek marks it in the word itself, so **the grammar records who is being spoken to**, not just who is being spoken about.

That is what makes this item close rather than arguable. Peter is not described as answering a man. **He addresses one, and the form of the word he uses carries the sex**, so the detail is in the grammar rather than in a narrator's summary that could be loose.

So the texts do not merely name different people. **They mark different sexes for the same challenge in the same sequence**, and multiplying challengers cannot absorb it without adding an event.

Why this is the layer's best moment

Because the check ran both ways and was recorded both ways.

Two other items were checked against the same Greek text and **cut in favour of the defence**. One rescue that had been scored as dead turned out to be alive. Another item narrowed to a smaller objection than the compilations claimed. And one check **corrected an error the coder had made himself** about which case a Greek noun was in.

A verification procedure that only ever confirms your side is not a verification procedure. This one produced two findings for, two against, and one self-correction, and all five are in the same file.

The self-correction is the most useful of the five and it is worth saying what it was. The coder had claimed a Greek noun was in one case when it is in another, and the error was found by writing the check as a **script** rather than as a sentence.

That is a transferable technique and it is the same lesson the toolchain in this series keeps teaching. **A claim written as prose cannot be run**. A claim written as code gets executed against the actual text, and it either produces the result you said it would or it does not.

Two **patristic** rescues were killed by the same script, and **writing them as code caught two further errors in the coder's own argument** on the way. A claim you can execute is a claim that can embarrass you, which is precisely what makes it worth more than a footnote.

Patristic means belonging to the church fathers, from the Latin for father. The fathers are the Christian writers of roughly the first five centuries, the ones who were arguing about these texts while the arguments were new: Origen, Chrysostom, Augustine, Eusebius, Jerome. Episodes 18 to 20 grade them, and the word turns up here because two of the rescues being scored in this layer turn out to be theirs rather than modern.

Which is itself a finding worth keeping. **Four of the rescues scored in this study are patristic**, so the common line that harmonisations were invented recently to plug holes is simply false, and it should stop being said.

The device: the two survivors, priced

ITEM	THE CLASH	THE BEST RESCUE	WHY IT COSTS 3
Judas	hanged against fell headlong	the rope or branch broke	no text states it
Peter's accusers	feminine against masculine challenger	more challengers than reported	requires an unnarrated extra challenge

The clinical anchor

The one abnormal result that changes management.

Forty numbers come back on a screen and thirty-nine are noise you will never think about again. One of them changes what you do tomorrow.

And the skill is not finding abnormal numbers, which is trivial, it is knowing which ones are load-bearing. A study that reports all forty with equal emphasis has not done the work, and a study that reports two out of fifty-one is claiming exactly this kind of discrimination and had better be right about it.

The character anchor

A plot hole that actually breaks the story, against the hundred that do not.

Every long-running show has continuity errors, and fans catalogue them by the hundred, and almost none of them matter because the story does not rest on them.

Then there is the one where a character knows something they could not know, and it matters because the plot turns on it. The difference is not how wrong the detail is. It is **what depends on it**, which is why the study reports the core separately.

What would count against this

1. **Two being suspiciously few.** It is a low number and it comes from a dissolution process whose reliability has just been disqualified. **The worst case in episode 15 puts accepted items at 120 and still leaves two at cost 3**, which is the only reason the number can be said at all.
2. **The gender argument being the coder's own philology.** It is checked against the Greek text with a reproducible script, which is better than a commentary, and it is still one person reading grammar.
3. **The Judas correction being a small mercy.** Moving 4 to 3 costs almost nothing rhetorically and buys a great deal of credibility. **That reading is available**, and the counter is that the Bayes value was fixed before the recomputation.

Handing off

Episode 17 states the hole this whole layer sits in, which the study calls its top priority and has not built.

Questions

1. **Price the rope-broke rescue yourself.** Is 3 right, or is it 4?
2. The Greek check cut both ways. **How many of your own checks have?**

3. **Before episode 17:** two hard contradictions. Out of what? Guess the denominator, then say how you would count it.

Collecting episode 16

Two items at cost 3, one corrected downward against the story, one closed by grammatical gender checked in the Greek. Now the question that decides whether either number means anything.

The hole

Two hard contradictions out of what?

Against thousands of cross-checkable claim-pairs, two is remarkable. Against forty, two is bad. **Same numerator, opposite verdicts**, and nothing in this study settles which frame you are in.

The pre-registration specifies the design for a denominator. **It has not been built**. So the defensible claim is still about the accusers' list rather than about the texts, and the study's own front page says so.

Why a denominator is genuinely hard here

It is not laziness, and understanding why is more useful than the complaint.

To count cross-checkable claim-pairs you must first decide **what a claim is**. Is "Jesus went to Jerusalem" one claim or three? Does every clause count? Every proper noun? The answer changes the denominator by an order of magnitude and there is no principled place to stop.

And the unit has to be fixed **in advance and independently**, because a denominator chosen after you know your numerator is not a denominator. It is a dial.

Watch how easily that dial turns. You have two hard contradictions. You want to know whether that is good. You start counting claim-pairs, and every time you make a granularity decision you can feel which way it moves the ratio.

Slice finely and the denominator explodes and two looks like nothing. Slice coarsely and it shrinks and two looks serious. Neither slicing is wrong, both are defensible in a methods section, and **the person choosing knows what each one does to the answer**.

This is Season 4 episode 10's problem exactly, arriving in a different project: **free parameters plus a target in view**. And the solution available there was to enumerate every choice rather than make one, which nobody has yet tried here.

This is the same problem as episode 3's NULL category, arriving at full scale. **The denominator is the count of chances**, and chances are only countable once somebody defines what a chance is.

The lead that exists and has not been used

There is a fourth-century enumeration of the cross-checkable set: the **Eusebian canon tables**, which index which passages of which gospels parallel each other. Built for a completely different purpose, by somebody with no stake in this question, seventeen centuries before it was asked.

That is a sampling frame, not a denominator. Pericopes are not proposition-pairs, and converting one to the other is the whole difficulty again. But a stake-free enumeration by a disinterested party is exactly the kind of thing you want, and it needs its own pre-registration before anybody touches it.

The distinction between a sampling frame and a denominator is worth having, since they get confused constantly and the confusion is expensive.

A denominator is the count of chances. A sampling frame is a list you can draw from so your sample is not chosen by you. The canon tables give you the second: a principled list of parallel passages, compiled by somebody with no interest in this question, from which you could draw at random.

That is genuinely valuable and it solves a different problem. **It stops the analyst choosing which passages to examine**, which is exactly the failure that sank run 1's fit ratio in episode 7. It does not tell you how many claim-pairs exist inside those passages, which is still the open problem.

Worth noticing anyway: **the fix for episode 7's disease has been sitting in a fourth-century manuscript apparatus the whole time**, and nobody in this project connected the two until the denominator question forced it.

What the study is allowed to say without one

About the accusers' list: everything above. 671 to 153 to 51 to 2, with the dissolution rate disqualified and the hard count structurally surviving.

About the New Testament: almost nothing quantitative.

A count without a denominator is not a low number or a high number. It is not a rate at all, and the honest move is to stop calling it one.

The device: the same numerator, two frames

DENOMINATOR	2 HARD CONTRADICTIONS IS	VERDICT
40 cross-checkable pairs	1 in 20	bad
400	1 in 200	unremarkable
4,000	1 in 2,000	remarkable

Nobody knows which row is true, and every public argument about biblical contradictions is conducted as though the row were obvious to whoever is speaking.

The clinical anchor

A rate with no population denominator.

Twelve cases this year. Up from eight. The room reacts.

Twelve out of how many admissions? If the catchment grew by half, twelve is a fall. And the reason this happens constantly is that the numerator is easy to count and lives in an incident system, while the denominator requires somebody to go and get the activity data, which is dull.

The easy number gets published and the hard number decides what it means.

The character anchor

A batting average with no at-bats.

Four home runs. Is that a good season? You cannot say, and everybody knows you cannot say, because in sport the denominator is published as a matter of course and nobody would accept the numerator alone.

In almost every other domain we accept the numerator alone constantly, and the only difference is that sport made the habit of publishing both.

What would count against this

1. **The denominator being unbuildable in principle.** If "claim" cannot be defined non-arbitrarily, no denominator exists and this layer is permanently a statement about lists. **That is a real possibility and the study does not rule it out.**
2. **The Eusebian route smuggling in a frame.** Eusebius indexed parallels according to his own sense of what parallels are, which is a theory of the texts.
3. **This objection applying to the accusers too.** It does, and harder. **671 is also a numerator with no denominator**, and it is quoted far more often.

Handing off

Episode 18 turns the same instrument on the defenders, grading the church fathers on the scale everybody else was graded on, and the headline is an inversion nobody expected.

Questions

1. Define "a cross-checkable claim" well enough to count one. **How long did you last?**
2. Which row of the device table do you believe, and what is your evidence?
3. **Before episode 18:** do you expect the fathers to answer these items better or worse than modern defenders? Commit.

Collecting episode 17

No denominator, so no rate, so the claims are about the accusers' list rather than about the texts. Which leaves one thing this layer can still do honestly: turn the instrument on the defence.

The design, and it is stricter than it needed to be

Twelve of the hardest items on the docket, expanded to nineteen, chosen by **mechanical criteria written to the file before any patristic source was fetched**.

That ordering matters for the same reason it always does. Choose the items after reading the fathers and you have chosen the items the fathers happen to address. Choose them first, mechanically, and **the coverage becomes a finding rather than an artifact**.

Then the patristic answer to each is graded on a fixed scale: **GOOD, PARTIAL, BAD, ABSENT, AGAINST**.

The headline, and it inverts what everybody expects

The fathers cover MORE of the hard items than modern harmonies do, and they answer them WORSE.

Both halves are surprising and the second explains the first.

The fathers engage more because they were writing in a culture where the objections were live and being pressed by people they had to answer. They answer worse because **their answers are the ones that do not survive being said aloud now**, and modern defenders have quietly stopped making them.

That is not a discovery about the fourth century. It is a discovery about the present: **the modern harmony literature is smaller than the ancient one because it has been filtered**, and the filtering removed the answers that embarrass rather than the questions that are hard.

And the filtering is not a conspiracy, which is what makes it worth understanding.

Nobody decided to drop the weak answers. What happened is that each generation repeated the answers that worked in argument and quietly stopped repeating the ones that got laughed at. **Selection by embarrassment, operating over centuries**, with no author and no moment where anything was suppressed.

The result is a literature that looks stronger than its tradition actually was, **because the failures were not refuted, they were forgotten**. Which means a modern defender is inheriting a curated case and has no way of knowing what was curated out, unless somebody goes back and grades the originals.

There is a version of this in every field, including this one. **Seasons 2 to 4 exist because a research repository kept its failures**, and almost nothing else does.

The worst finding is a silence

Augustine's *De consensu evangelistarum* works the Judas material in detail. The thirty pieces of silver, the potter's field, the misattribution of Zechariah to Jeremiah. He is thorough, and he is engaging with the hard version.

He never mentions Acts 1:18.

That is the verse with the falling headlong and the bursting open. It is the other half of the contradiction, it is the reason the item is at cost 3, and the most thorough ancient harmony of the gospels works the pericope around it and does not touch it.

Verified, not inferred from a summary.

The gate that catches everybody

Augustine's solution to the staff problem, where the disciples are told to take a staff in one gospel and not in another, is **gate A2 with no licence**: a difference of sense with nothing attested to support it.

A2 is the gate that fired F4 against this project two episodes ago.

So the same gate is abusable by the accuser, by the defender, and by the coder, **and all three have now been caught by it in this season**. That is the strongest evidence the instrument is symmetric, and it is the only kind of evidence that counts: it caught the person holding it.

And one father who helps the defence, which the headline hides

A survey reporting that the fathers answer worse should be read with suspicion, and the check is whether anything in it came out the other way. One thing did, and it is worth stating because it costs the season's headline something.

Papias, writing around 110 to 130, says Mark wrote down what he remembered accurately but "not in order."

That is a very early witness attesting, by name, that these texts were arranged **thematically rather than chronologically**, and that the arrangement was known and unremarkable at the time.

Which matters directly, because a large family of alleged contradictions is about **ordering**: this gospel puts the temple cleansing early and that one late, this healing before that sermon. Under the study's own rules, a text whose genre does not claim chronological order is not asserting one, **and a non-assertion cannot contradict**.

So Papias supplies exactly what the guard demands: **a licence from outside the argument**, cited, early, and made for reasons that have nothing to do with any modern dispute. That is a cost 1 rescue on a whole class of items, offered by a father, and it makes several of the accusers' entries weaker rather than the defence's answers.

It changed no cost in this study, because rescoring on a finding discovered mid-survey needs a new pre-registration. **It is reported against the story anyway**, which is what the season keeps asking of everybody else.

The device: what the grading found

GRADE	WHAT IT MEANS	ROUGHLY HOW MANY
GOOD	a licensed dissolution that survives	two
PARTIAL	answers part of the item	some
BAD	an answer that fails on its own terms	five
ABSENT	the item is worked and this half is not mentioned	the Judas silence
AGAINST	the father concedes the point	Origen

The clinical anchor

A textbook that omits the complication.

The current edition describes the procedure cleanly. The 1974 edition describes the same procedure and spends two pages on a complication the modern text does not mention.

The complication did not stop happening. It stopped being written about, because the answer that used to be given for it is no longer one anybody wants in print. **The older text is more honest and more wrong at the same time**, and you learn most by reading both.

The character anchor

A remake that quietly drops the hardest scene.

The original had a sequence nobody could shoot well and it shows, and everybody remembers it being clumsy. The remake simply does not have it.

The remake is better and it is also evidence that the scene was unfilmable, which is a fact about the story you can only recover by noticing an absence.

What would count against this

1. **Nineteen items being a small survey.** It is, and it is the hard tail rather than a representative sample by design.
2. **Grading ancient answers by modern standards.** Genuinely unfair in one direction, and the study's guard is symmetric on age: **age is not licence, and age is not a defect either.** No cost moved because a source was old.
3. **A survey finding only failures being suspect.** The study says this about itself in episode 20, and it is the right worry.

Handing off

Episode 19 names the defect the fathers share with the skeptic lists, and it is the same defect, mirrored.

Questions

1. **The Judas silence.** Charitable explanations exist. Give the best one, then say what would distinguish it from avoidance.
2. **Name a literature you personally rely on that has been filtered the same way**, by embarrassment rather than by refutation. How would you detect it from inside?

3. **Before episode 19:** what single property would you look for to tell a real defence from a reflex?

Collecting episode 18

The fathers cover more hard items than modern harmonies and answer them worse, and Augustine's staff solution failed the same gate that fired a falsifier against this project.

Now the property underneath all of it, which is the most transferable idea in the season.

The defect, named

Verdict flatness. An output that does not vary with its input.

If a procedure returns the same answer whatever you feed it, **it is not reading the input**, and its answer carries no information regardless of whether the answer is true.

This is the series refrain arriving in its sharpest form, and it is worth saying the whole sentence rather than pointing at it.

A check that cannot fail prints reassurance.

Verdict flatness is what that sentence looks like from outside: a stream of confident verdicts, each one individually defensible, none of them responsive to the case. **The reassurance is real and it is being printed rather than earned**, and from inside there is no difference, which is why the only reliable test is to ask what input would have changed the output.

Augustine states it, in the fourth century, as a rule

Letter 82, and it is the cleanest example anywhere in the material.

When a scriptural text appears to be in error, Augustine says, the possibilities are: the manuscript is faulty, the translator is bad, or **I have failed to understand**.

Three options, exhaustively listed. **"The author erred" is not on the list**.

That is a decision procedure whose output cannot depend on the text. Feed it a real error and it returns one of three, and feed it a scribal slip and it returns one of the same three. It cannot come back negative, and Augustine wrote it down as a rule rather than merely practising it.

And this is not a stray aside. It is the load-bearing premise of his long dispute with Jerome over Galatians 2, which is itself worth knowing about, because **Jerome and Augustine disagreed**: Jerome held Paul's rebuke of Peter was staged, Augustine held it was real. **The fathers were not one voice**, and anybody quoting "the fathers" as a bloc is quoting a selection.

The same defect on the other side, mirrored

The skeptic compilations have their own flatness, and this season already measured it. **27% of the list is the same complaint recounted**. Five entries misdescribe their own citations. Four put a question on one side, meaning the entry alleges a contradiction where one half is not an assertion at all.

A compiler who checked would have caught all of those. The list is produced by a procedure whose output does not depend on reading the verse carefully, which is the same defect wearing the opposite jersey.

The cost scale catches both, and that is the argument for having a scale at all rather than a verdict. A binary lets a flat procedure hide, because flat procedures produce confident binaries. A price list forces you to name a channel, and a flat procedure has no channel to name.

Chrysostom, and the argument that predicts nothing

Around 390, Chrysostom makes Layer 1's own argument: the evangelists' agreement proves their honesty, and their disagreement proves they did not collude.

Read that twice. Agreement confirms and disagreement confirms. As stated it cannot come back negative, on any evidence, ever.

And this is the point in the season where it should be uncomfortable, because **that is very close to the argument this study is running**. The difference is the two things Chrysostom did not have and could not have had: **a control and a denominator**. With them, "disagreement is what independence looks like" becomes a measurable prediction, because you can ask how much disagreement, of what kind, compared to what.

Use him that way, never as a cheap shot. He is not being stupid. He is running the right intuition without the equipment that turns it into a test, and the equipment took seventeen centuries to arrive and this project still lacks half of it.

The device: flat against responsive

PROCEDURE	WHAT CHANGES THE OUTPUT	FLAT?
"faulty manuscript, bad translator, or my failure"	nothing	yes
"agreement confirms, disagreement confirms"	nothing	yes
a list that never checks its own citations	nothing	yes
the cost scale	which channel supplies the rescue	no
the nine gates	which qualifier fails	no, and F4 measured how much

The clinical anchor

An observation chart where every set of numbers is identical.

Six hours of hourly obs and the pulse is 80 at every one of them. Respiratory rate 18, every time. Blood pressure to the same digit.

No patient does that. A real pulse wanders by a few beats between readings even in somebody perfectly well, and a column of identical numbers is not a stable patient, it is a chart filled in from memory at the end of the shift.

And notice what makes it detectable, because it is the whole diagnostic. **Not one of those numbers is abnormal.** Every single value is defensible, the patient may genuinely have been fine, and the chart is still worthless, because **a measurement that does not vary was not taken.**

The character anchor

A reviewer who gives everything five stars.

Their five-star review of a genuinely great film is worthless, and the film is still great. **Nothing about the flatness makes the verdict false.** It makes the verdict uninformative, and those are different failures, which is why this criticism does not touch whether Augustine was right about any particular passage.

What would count against this

1. **Flatness sometimes being correct.** If a text really is reliable, a procedure that keeps returning "reliable" is tracking reality. **The problem is that you cannot tell the two apart from the output,** which is the whole argument and also its limit.
2. **Augustine being quoted out of context.** The letter is a genuine methodological statement and it is also a letter, written to a person, about a dispute. **Tier V on the quotation, and the reading is an interpretation.**
3. **This project having its own flatness.** It should be asked. **F4 firing is evidence against it,** because a flat procedure does not produce a disqualifying number about itself.

Handing off

Episode 20 finds the exception, one father who states the criterion this whole project runs on, against his own interest, in the fourth century.

Questions

1. Test something you believe for verdict flatness. **What input would change your output?** If none, say so.
2. Chrysostom's argument needs a control and a denominator. **Design his control.**
3. Is "the author erred" a legitimate thing to exclude in advance, if you already hold the text is inspired? Argue it fairly.

Collecting episode 19

Verdict flatness on both sides, in the same season, caught by the same scale. And a survey that has found almost nothing but failures, which is a condition that should make you suspicious of the survey.

The exception

Eusebius, writing around 340, in a work answering apparent contradictions between the gospels, on the resurrection accounts.

He is discussing a charge that John conflicts with the others about time and place, and he says the charge

"really would have been justifiable... **if John had kept the same time and the same place**".

Same time. Same place. Same respect. That is Tier A. That is episode 10's definition, stated by a defender, in the fourth century, **as the condition under which the objection against him would have been correct.**

Why this is the best thing in the patristic material

Because of who says it and when.

He is not conceding the item. He is naming, in advance, **what would have made his opponents right**. He is stating his own falsification condition while arguing the other way, which is the precise habit whose absence this project has been recording on both sides for five seasons.

It acquits Eusebius of the flatness charge laid on Chrysostom. Chrysostom's argument could not come back negative. Eusebius named the exact circumstance in which his would have.

And the value of that is not that he was unusually honest, though he may have been. **It is that the sentence is checkable seventeen centuries later.**

Once a defender says "the objection would be justified if X", anybody can go and look at whether X obtains. The argument has acquired a handle. **You do not have to trust him, assess his character, or weigh his tradition against another tradition.** You go and read John and check the time and the place.

That is the entire practical value of stating a falsification condition, and it is why this series keeps returning to it. **It converts a claim about a text into an instruction anybody can follow**, and instructions outlive the people who wrote them in a way that authority does not.

And the same author, in the same work, doing the opposite

This is why the episode is not a victory lap.

Eusebius also proposes that a troublesome variant is not original, and gives as grounds that it contained **something contradictory to the evidence of the other evangelists**.

That is harmonisation driving textual criticism. The text is judged inauthentic because it disagrees, which means disagreement can never survive as evidence, which is flatness at the level of the text rather than the doctrine.

And elsewhere he answers the empty-tomb tangle by **multiplication**, deriving the number of appearances from the number of evangelists: four evangelists, therefore four occasions and four women. **The count comes from the sources rather than from the events**, which is worse than the modern harmonies that get criticised for the same move.

One author, one work, both things. The best falsifiability statement in the material and one of the worst pieces of reasoning in it. **People are not consistent, and a survey that sorts sources into good and bad has not read them.**

The warning the study writes against itself

A pattern made of one instance is not a pattern.

The study lists P6 anyway, and states why: **everything else in the survey is a failure, and a survey that finds only failures should be suspected of having looked only for them.**

That sentence is the most self-aware thing in six seasons. It is a researcher noticing that his own results have a suspiciously uniform sign, and treating the uniformity as a warning about his method rather than as a finding about his subject.

The device: Eusebius, both ways

WHAT HE DOES	WHERE	VERDICT
states the falsification condition against his own interest	Nicetas-Marinus 5	the best in the material
judges a variant inauthentic because it disagrees	<i>Ad Marimum</i>	flat, at the text level
derives four appearances from four evangelists	<i>Ad Marimum</i> 4	worse than Lindsell
the corpus contradicts itself on gospel order	flagged by its own editors	

The clinical anchor

A colleague documenting their own complication in their own notes.

Nobody made them. The note says what happened, when, and what they did about it, and it is the strongest evidence in the file precisely because it runs against the interest of the person writing.

And that same colleague will have other notes that are thin and defensive. People are not uniform, and a reviewer who decides someone is reliable and then reads everything through that decision has stopped reading.

The character anchor

An opponent conceding a point on record.

In a debate, the concession is the only moment worth quoting, because everything else either side says is what they came to say.

And it does not make them right about the rest. A person can be scrupulous in one paragraph and sloppy in the next, and the concession is evidence about the paragraph, not about the person.

What would count against this

1. **One instance being nothing.** Said by the study about itself. **Tier V on the quotation**, which is worth something: the passage was fetched from an edition after being quoted off a blog, and reading the edition corrected three claims the project had made.
2. **Translation carrying the argument.** "Same time and the same place" is a translation, and the fit with Tier A may be partly the translator's.
3. **Selection.** This is the best thing found in a survey of nineteen hard items. **How many defenders were not surveyed?** Almost all of them.

Handing off

Episode 21 leaves Layer 1 entirely for the question everybody actually wants answered, and the first thing that layer does is fence off the answer.

Questions

1. Eusebius states the criterion and violates it in the same work. **What do you do with an author like that?**
2. "A survey that finds only failures should be suspected of having looked only for them." **Apply that to this whole series.**
3. **Before episode 21:** what could a study of the resurrection case honestly conclude? Write it down before you are told.

Collecting episode 20

One father stating the falsification condition against his own interest, and the same father judging a text inauthentic because it disagreed. And a warning that a survey finding only failures may have looked only for them.

The ceiling, fixed before any scoring

Layer 2 asks whether the historical case for the resurrection is honestly built. Its rulebook fixes, in advance, what it can and cannot conclude.

It can test whether the claimed facts are supported, whether consensus claims are verifiable, whether alternatives are fairly eliminated, and whether the case is circular.

It cannot test whether a miracle occurred.

The reason is stated plainly and it is not a dodge. **Once the facts are fixed, the inference is dominated by the prior on miracles**, and no source check touches a prior. Two people who agree on every fact in the case and disagree about whether miracles happen will disagree about the conclusion, and nothing in a bibliography moves that.

The best possible outcome of Layer 2 is "the case is honestly built." Never "it happened." That is written down before any scoring, so it cannot be quietly upgraded afterwards.

The reasoning behind that ceiling deserves working, because it is not obvious and it is not a concession to sceptics.

A historical argument moves you from evidence to a conclusion. **How far it moves you depends on what you thought was possible before you started.** Somebody who holds that miracles do not occur will read any quantity of testimony as evidence that something else happened, and they are not being unreasonable, because a sufficiently improbable claim requires proportionate evidence.

And no amount of checking citations touches that. You can verify every source in the case, find every claim supported, and two readers will still land in different places, because the disagreement was never about the sources.

So the ceiling is not modesty and it is not scepticism. **It is a statement about which disagreements a bibliography can settle**, and this one is not among them. What a bibliography can settle is whether the case is built out of things that are actually there, which is worth doing and is a smaller job.

The rule that carries the whole layer

Layer 1 had structural protection: hostile compilers chose the docket, so the coder could not skip hard items.

Layer 2 has no such protection, because there is no hostile compilation of objections to a specific apologetics website. So one rule does all the work:

The apologetics site is the object under test and never the docket source.

Objections come from **named critics only**. If a session takes the objection list from the site being examined, the site has chosen which objections it will answer, and **Layer 2 is void rather than weakened**.

The key auditor named in advance is **Dale Allison**, chosen specifically because he affirms the core facts while rejecting the methodology, so both halves of him always get reported.

G1 fired, at two levels

G1 said: if the empty-tomb consensus claim cannot be verified from published, checkable data, the site's strongest citation fails and the fact rests on an unauditible premise.

Level one. The site asserts that most scholars, including non-Christian ones, accept the tomb was found empty, **and cites no published source for it**.

Level two, and this is the worse one. The underlying survey behind that consensus claim is one its own author described in 2012 as "**mostly unpublished**."

So the strongest citation in the strongest section rests on a bibliography nobody can inspect. **Cost 4**, for an unauditible warrant, which is the same discount this project applied to its own run-1 fit ratio.

Apply your own discount to yourself first, then to others, and the number is the same number.

The falsifier that signalled

G4 asks whether the surviving Layer 2 arguments need detail-level gospel accuracy. If they do, **Layers 1 and 2 are not independent and this project's whole framing is wrong**.

It signalled without firing. **Reported against the project's own interest**, and left open.

The device: what Layer 2 can say

QUESTION	ANSWERABLE	HOW
are the claimed facts supported by published sources	yes	check the citations
is a consensus claim verifiable	yes , and G1 fired on it	find the survey
are alternatives fairly eliminated	yes	read the critics directly
is the case circular	yes	trace the warrants
did it happen	no	prior on miracles dominates

The clinical anchor

A guideline citing an unpublished dataset.

The recommendation is grade A. You follow the citation and it goes to a conference abstract, or to data held by the group that wrote the guideline.

The recommendation may be entirely correct. And you cannot check it, and grade A means checkable, so the grade is wrong even if the advice is right. **Those are separate facts** and conflating them is how a whole field ends up standing on four unpublished trials.

The character anchor

A Wikipedia claim whose citation leads to a dead link.

The sentence has been there for nine years and is quoted everywhere. Follow the footnote and the page is gone, and the archived version does not say what the sentence says.

The claim might still be true. What has collapsed is the reason anybody had for believing it, and the difference between those two is invisible until somebody clicks.

What would count against this

1. **The site being a soft target.** It is one apologetics website, not the scholarly case. **Grading it does not grade Wright or Habermas properly**, and the study should not be read as though it did.
2. **The consensus claim being true anyway.** Very possibly. **G1 is about the warrant, not the fact**, and an unauditable warrant for a true claim is still unauditable.
3. **No control existing for Layer 2, at all.** Layer 1 had the apocrypha. **There is no corpus of comparably-argued cases**, and none can be built, which is a permanent ceiling rather than an unfinished task.

Handing off

Episode 22 closes the season and adds up what three fenced layers actually bought, which is less than the question deserves and more than the argument usually gets.

Questions

1. **Write the sentence Layer 2 is allowed to conclude.** Then write the one people will quote. Notice the gap.
2. G1 fired on an unauditable warrant. **Find one in something you believe.**
3. If no control can be built for Layer 2, should it have been attempted at all?

Collecting episode 21

A ceiling fixed in advance that Layer 2 cannot test whether a miracle occurred, an inversion rule doing all the protective work alone, and G1 firing at two levels on the strongest citation in the case.

The season, added up

LAYER	WHAT IT ESTABLISHED	WHAT IT COST
1a	the apocrypha embellish about 11× more, robustly	magnitude failed twice, direction only
1a control	restraint holds in direction	H3 and H5 fired , four units carry the verdict
1b	2 at cost 3, 0 at cost 4 , from a hostile docket of 671	F4 fired , dissolution unreliable, inerrancy gone
patristic	the fathers cover more and answer worse	one father states our own criterion
2	the strongest citation rests on an unpublished bibliography	cannot test whether it happened

Four falsifiers fired across one season, and every one of them fired on the person running the study rather than on his opponents.

Set that against the first four seasons and the contrast is the finding. In Seasons 2 to 4 the falsifiers were about the instrument: a null was badly built, a round was under-powered, a metric measured the wrong half. **Here they are about the person**. The dissolution rate rests on his own sense-distinctions. The verdict rests on four sentences he coded. The blinding does not blind him.

An instrument you can fix. A coder you cannot. Which is why this season's honest ceiling is lower than Season 4's despite the study being better designed, and why the phrase "strong structured pilot" appears on every headline rather than in a limitations section.

The one sentence worth keeping

The apocrypha demonstrably embellish and hard-contradict far more than the canon, which is robust. The canon appears to interlock more, which is real in direction and inflated in degree by who did the coding. And there is no denominator, so none of it is a rate.

Three clauses, three different confidences, nothing thrown away and nothing promoted. **That is the whole method in one sentence**, and it took five seasons to be able to write it.

What this season answers that Season 4 could not

Season 4 ended unable to interpret its own best result, because four partisan documents about one week had nothing to be compared against.

This season had a control, and the answer is genuinely informative: **partisan documents about Jesus, written by believers in the same tradition, do not all behave the same way**. The apocrypha embellish eleven times more. That comparison was

unavailable in every one of the first four seasons, and it is why the numbers here mean more despite the instrument being cruder.

A cruder instrument with a control beats a refined instrument without one, and Seasons 2 through 4 are 50 episodes of refining an instrument without one.

That sentence is the most useful thing in six seasons and it should be uncomfortable, because it means most of the effort went to the wrong place.

Season 2 fixed the tokeniser, de-noised the lexicon, audited the corpus and chose a metric for principled reasons. Every one of those was a real improvement and **none of them addressed the question of what the measurement should be compared against**. The instrument got sharper and the comparison stayed absent.

Refinement is visible work and controls are invisible work. You can see a better lexicon. A control looks like a folder of texts nobody is studying, and it is the folder that decides whether any number means anything.

Which is the practical lesson for anybody running a study of their own: **build the control first, while it is still cheap and while you do not yet know which way it will cut**.

What is still missing, said plainly

- **The denominator**. The top priority, specified and not built. Without it every count in Layer 1b is a statement about a list.
- **A second coder**. The ceiling on everything, stated on every headline.
- **The apocryphal control on the contradiction count**. F3 is registered and **untested**. Nobody has run the nine gates on the Gospel of Peter, and until somebody does, "51 accepted" has no comparison.
- **190 items unscored**, the New Testament against the Old.
- **32 of 51 audits are tier K**, meaning real defenders and real arguments with citations unverified. A lead, not a fact.

The clinical anchor

Three tests answering three different questions, and the harm when they merge.

The troponin, the echo and the angiogram answer different questions and a consultant who lets a normal troponin reassure them about the valve has merged three results into one feeling.

The fences are the skill. Not the tests, which anybody can order.

The character anchor

Three locked rooms, and the temptation to say you searched the house.

You searched one room thoroughly, one partially, and looked through the keyhole of the third. That is a real search and it is a specific claim, and the sentence "I searched the house" is available, shorter, and false.

What would count against this

1. **One coder, and it is stated here for the last time.** Nothing in this season survives a reader who declines to trust the coding, and the season has said so in every episode by design.
2. **Four fired falsifiers being a bad sign rather than a good one.** A study whose own checks keep firing may simply be a badly designed study. **That reading is available and the counter is thin:** the alternative design is one where the checks were never set.
3. **The whole season being about method rather than about God.** Deliberately. And somebody could fairly say that a study which fences off every interesting question has bought its rigour by declining to be about anything.

What comes next

Season 5b takes one allegation per episode and works it to the floor: what both sides say, what the Greek actually says, and what the rescue costs. It is the most repeatable format in the series and it is where the price list stops being a scale and becomes a set of worked examples.

Questions

1. **Four falsifiers fired.** Good study or bad study? Defend it.
2. Write the denominator's pre-registration. **One page, before you know the number.**
3. Nobody has run the gates on the Gospel of Peter. **Predict the result,** then say what you would conclude from either outcome.

