The Myths Repo, Forensic Record

Every instrument built in `~/Myths`, every prediction made, every result, every piece of the mathematics, and every error the project caught itself committing.

Compiled 14 August 2026 from the repository at commit `930d8e3`. Nothing in this document is remembered. Every number was read off a file on disk or recomputed from the data. Where a claim comes from a file, the file is named.

> *This is a record, not an argument. It does not tell you what the work means, whether it was worth doing, or what to say about it on camera. Parts VII, VIII and IX are where you take over. The conclusion is deliberately absent.*

0.1 What is actually in the repository

Four instruments were pre-registered. **Three** were run.

#	INSTRUMENT	LOCATION	WHAT IT TESTS	ROUNDS RUN
1	Four-bucket scapegoat detector	`~/Myths` root	Girard's four stereotypes of persecution, as a lexical measure over Greek and Roman texts	4
2	Conversion instrument	`~/Myths/conversion`	Girard's three-part novelistic conversion, as a rubric over novel endings	1, partial
3	Interlock instrument	`~/Myths/interlock`	Gospel transmission reliability, in three layers	3
4	**Restraint control**	`interlock/resurrection`	whether the canonical resurrection narratives show the features critics' own models of legendary growth predict	**0. Never run.** See Part IV

279 tracked files. 28 commits. **Seven** pre-registrations in total: four on the myth detector, one on conversion, and in the interlock tree one for the instrument itself, one for the contradiction count, one for Layer 2, and one for the restraint control.

> *Instrument 4 is the one nobody has mentioned in any summary of this project, including mine. It is arguably the best-designed file in the repository, and nothing was ever coded against it. Part IV is about that.*

The rule shared by all three, quoted from `interlock/DESIGN-RATIONALE.md`:

> **The one rule under all of them** Never tune the instrument, and never tune the sample, to pass. Everything above is machinery for making that rule enforceable by someone who does not trust the person who wrote it, including when that person is me.

0.2 The seven days

The whole repository was built on **seven calendar days**. Not months. This is the true figure and it is checkable with `git log`.

DATE	WHAT WAS BUILT	COMMITS
21 July 2026	corpus, v1 scorer, round 1, lexicon v2, round 2	6
22 July 2026	Correction 1, round 3 (localization), NEXT	4
24 July 2026	interlock run 1, NT contradiction count, F4 fires	4
25 July 2026	verdict, registry, Greek checks, Layer 2, design rationale	8
5 August 2026	pre-registration C (restraint control)	1
11 August 2026	round 4, the dissent test	2
12 August 2026	conversion instrument, and own novella as Set D	3

Elapsed 21 July to 12 August is 22 days. Working days, seven. Most sessions ran between midnight and 03:30.

The full commit record, which is the only part of this a stranger can check.

#	DATE, TIME	HASH	WHAT LANDED
1	07-21 01:37	`6a252e4`	corpus, scorers, findings before myth-level scoring
2	07-21 01:37	`dd32cd0`	**PRE-REGISTRATION 1**, Ovid tier assignments
3	07-21 01:39	`8986df5`	RESULTS 1, the MIN prediction failed
4	07-21 03:00	`00a9556`	**PRE-REGISTRATION 2**, lexicon v2 and six predictions
5	07-21 03:14	`e1f5253`	RESULTS 2, 4 of 6, density confound identified
6	07-21 03:24	`4d4afa6`	NEXT.md handoff, session transcripts archived
7	07-22 00:55	`ab8fcc6`	**CORRECTION 1**, truncated text, speaker diagnostic
8	07-22 00:55	`fc094bb`	**PRE-REGISTRATION 3**, localization
9	07-22 03:22	`4188308`	RESULTS 3, 1 of 5, the one vacuous
10	07-22 03:25	`1dc15d7`	NEXT.md, the locator reframe falsified
11	07-24 22:53	`83f5f55`	**PRE-REGISTRATION NT**, definition, gates, falsifiers
12	07-24 23:09	`c148944`	RESULTS NT, **F4 fires**
13	07-24 23:42	`3248616`	rescue-source pass, no licensed cost changes
14	07-24 23:57	`0724875`	round 2b, Shamoun sites, primary harmony scored
15	07-25 00:06	`188696b`	VERDICT.md, all 153 items with reasons
16	07-25 00:26	`4457b48`	REGISTRY, README, **Judas 4 → 3**, cost-3 Bayes case
17	07-25 00:37	`ac2f784`	all 51 accepted items get a defender audit, two tiers
18	07-25 00:49	`772664c`	five rows K → V against MorphGNT, one own note corrected
19	07-25 00:56	`0fda18f`	**PRE-REGISTRATION R**, Layer 2
20	07-25 00:58	`9cfc1fa`	Layer 2 round 1, **G1 fires at two levels**
21	07-25 01:04	`ce95dd1`	Layer 2 handoff, README restructured for three layers
22	07-25 01:30	`96b435c`	DESIGN-RATIONALE, 20 decisions
23	08-05 03:17	`4c614fc`	**PRE-REGISTRATION C**, the restraint control. **Nothing follows it.**
24	08-11 02:07	`9310c30`	**PRE-REGISTRATION 4**, the dissent test
25	08-11 02:17	`0c765d0`	RESULTS 4, nothing survives its own controls
26	08-12 16:50	`d671056`	**PRE-REGISTRATION**, the conversion instrument
27	08-12 17:12	`cf4bfa0`	RESULTS 1, conversion fails P1 and P4
28	08-12 17:25	`930d8e3`	addendum, own novella scored as Set D

Two things are visible in that table before any content is read. **Fourteen of the 28 commits fall on 24 and 25 July**, one continuous stretch, which is where the entire interlock instrument was built. And **commit 23 is an orphan**: a pre-registration with no run behind it, and two further studies begun after it.

0.3 The scoreboard

Every named prediction ever committed in this repository, with its outcome.

ROUND	PREDICTION	OUTCOME
R1	MIN separates HIGH from LOW on Ovid myths	**FAIL**
R2	P1 genre ordering TRAGEDY > OVID ≈ EPIC ≈ MYTHOGRAPHY > CONTROL	**partial FAIL**
R2	P2/P3 no named LOW outscores any named HIGH	PASS
R2	P4 Bacchae crimes > 0.5 and violence in top quartile	**FAIL**
R2	P5 two Agamemnon translations within 25%	PASS (17.2%)
R2	P6 MIN non-zero for at least 75% of texts	PASS (90.3%)
R2	null test, buckets beat frequency-matched random	PASS, $p = 0.0045$
R8	U1 GOSPEL dissent co-locates with violence	**PASS**, $U = 0.4171$, $p = 0.0015$
R8	U2 ordinal series GOSPEL < TRAGEDY < MYTH	**FAIL**, tragedy not intermediate
R8	U3 the *Bacchae*, fourth test	**FAIL**, $p = 0.414$
R8	U4 gap ≥ 0.05 and significant	**FAIL**, gap +0.066 but $p = 0.123$
R8	N3 core reverses the tier order	**FIRED**
R7	L1 at least 3 of 4 Gospels co-locate	**FAIL, 0 of 4**
R7	L2 ordinal series on z	**FAIL**
R7	L3 the *Bacchae*, third test	**FAIL**, $p = 0.888$
R7	L4 above chance overall	pass, **vacuously**, 4 of 66 vs 3.3
R7	M1 myth pass rate not below gospel	**FIRED**
R7	M3 under-powered, median z within 0.2 of 0	**FIRED**
R6	P2 phrases beat a same-length frequency-matched n-gram null	**PASS**, 0 of 2000
R6	P1 full lexicon beats the same null	**PASS**, $p = 0.0075$
R6	P3 KJV phrase rate within 25% of ASV	**PASS**, 20.4%
R6	P4 GOSPEL > MYTH survives the translation swap	**PASS**, $p = 0.00064$
R5	reconstruct the round-2 null across all 48 defensible configurations	**UNSTABLE**, 24 of 48 under $p=0.05$; original procedure identified
R3	P1 at least 7 of 9 named HIGH localize	**FAIL, 0 of 9**
R3	P2 Bacchae at $p < 0.01$, peak in second half	**FAIL, $p = 0.66$**
R3	P3 at most 1 of 6 named LOW localizes	pass, **vacuously**
R3	P4 TRAGEDY z exceeds EPIC z by at least 1.0	**FAIL, gap 0.30**
R3	P5 at least 6 of 9 HIGH beat Null B	**FAIL, 0 of 9**
R4	D1 ordinal medians MYTH < TRAGEDY < GOSPEL	PASS nominally
R4	D2 GOSPEL > MYTH	PASS nominally, $p = 0.00034$
R4	D3 TRAGEDY > MYTH	PASS nominally, $p = 0.032$
R4	D4 frequency-matched null	**FAIL, $p = 0.296$**
R4	D5 Prometheus Bound top 3 of 23 tragedies	**FAIL, rank 13**
R4	D6 no myth text above the lowest Gospel	**FAIL, 7 of 71**
R4	S1 sensitivity, core only	**destroys D2** (p 0.00034 → 0.293)
R4	S2 sensitivity, epic excluded	**destroys D3** (p 0.032 → 0.388)
Conv	P1 Set A T = 1 in ≥70%, sum ≥2 in ≥70%	**FAIL, 50% and 50%**
Conv	P2 the three bits co-occur above independence	**NOT RUN**
Conv	P3 Set A mean bit-sum exceeds Set B	**NOT RUN**

ROUND	PREDICTION	OUTCOME
Conv	P4 The Bell Jar scores (1, 0, 0)	**FAIL, observed (0, 0, 0)**
Int	P1 canon interlock ratio beats apocrypha	PASS
Int	P2 apocrypha hard-collide more	PASS
Int	P3 apocrypha embellish at least 3× more	PASS, ~11×
Int	P4 the advantage does not depend on Synoptic copying	PASS
Int	P5 canon interlock ratio ≥ 0.60	**FAIL, 0.50**
NT	F1 any cost 4 on the core	did not fire
NT	F2 cost ≥3 exceeds 10% of admitted	did not fire
NT	F3 apocryphal control discriminates	**NOT RUN, untested**
NT	F4 A2 + A5 exceed 40% of rejections	**FIRED at 67.6%**
L2	G1 empty-tomb consensus verifiable	**FIRED, at two levels**
L2	G2 any alternative dismissed at cost 4	not yet scored
L2	G3 case circular at some premise	not yet scored
L2	G4 Layer 2 depends on Layer 1 detail accuracy	**signalled, not fired**
C	H1 canonical restraint ratio above 0.40	did not fire, 0.352
C	H2 control set below 0.50, instrument invalid	did not fire, **0.577**
C	H3 gap under 0.25	**FIRED, 0.225**
C	H4 over 30% unstatable judgement calls	did not fire, 9.8%
C	H5 result depends on any single unit	**FIRED, 4 of 22 units**

Count, as of 14 August. Of 51 committed predictions and falsifiers: **14 clean passes, 22 failures, 3 never run, 7 falsifiers fired, 2 vacuous passes**, and 2 nominal passes destroyed by their own registered sensitivity analyses.

That count changed today. It used to read *eight never run*, five of them belonging to a study that was designed and abandoned. **The restraint control was run on 14 August and two of its five falsifiers fired.** What remains untested is the conversion instrument's missing control set (P2, P3) and the apocryphal control the NT count still needs (F3).

One surviving positive result in the entire repository. Round 2's null test, p = 0.0045. Part VII shows why that one is also the weakest-protected number in the project.

1.1 The claim under test

Girard, *The Scapegoat*, ch. 2. Persecution texts carry four stereotypes:

1. **Crisis** — the sacrificial crisis, loss of distinctions. Plague, famine, sterility, the doubling of antagonists into indistinguishable rivals.
2. **Crimes** — what the victim is *accused of*. Crimes that erase difference: parricide and regicide, incest, sacrilege, poisoning and sorcery.
3. **Marks** — the victim's marks of selection. Stranger, cripple, king, monster.
4. **Violence** — the collective violence itself.

The standard objection to Girardian reading is that it is unfalsifiable: any text can be read as confirming it after the fact. The project's entire design answers that objection by stating what the theory predicts **before looking**.

> *Note the load-bearing distinction, because a later error turns on it.* `crimes` *is what the victim is* ***accused of****. The sparagmos, the tearing apart, is what the* ***mob does****. They are different buckets.*

1.2 What was built before the corpus existed

`make_scorers.py`. Four word-buckets, **66 words total**, written before any text was downloaded. Corpus: 19 Greek and Roman texts from Theoi, plus Homer's *Odyssey* in Butler's translation split into 24 books.

1.3 The mathematics of the whole instrument

Everything in rounds 1 to 3 is built from one quantity.

$$r\_b(T) = 1000 \cdot h\_b(T) / N(T)$$

where `h_b(T)` is the number of hits of bucket *b* in text *T*, and `N(T)` is the number of tokens. Hits per thousand words. Tokens come from `re.findall(r"[a-z]+", text.lower())`, so punctuation is stripped and case is flattened. Phrases are matched against the token stream re-joined by spaces, so that "father's murder" becomes "father s murder" and punctuation can never block a match the word-level pass would have made.

Two ways to combine four buckets into one score:

$$SUM(T) = \Sigma\_b \ r\_b(T) \quad MIN(T) = \min\_b \ r\_b(T)$$

SUM is accumulation. MIN is conjunction. This is the whole argument. Girard does not claim persecution texts are *saturated* with these four vocabularies. He claims the four appear **together**. MIN is the algebraic form of "together": a text scores only as high as its weakest stereotype, so a text drenched in violence with no accusation scores zero.

1.4 The Odyssey book 7 false positive

This is why MIN was chosen, and it is the cleanest demonstration in the project.

BOOK	SCENE	CRISIS	CRIMES	MARKS	VIOLENCE	SUM
7	Odysseus welcomed at Alcinous's palace	0.30	0.00	**6.21**	**0.00**	6.51
18	Irus the beggar	0.00	0.00	5.49	2.15	7.64
22	**the slaughter of the suitors**	1.09	0.22	0.22	4.14	5.67

Book 7 is a peaceful hospitality scene with **zero violence**, and it outscores every tragedy in the corpus on `marks`, because "king" and "queen" saturate the passage. The actual massacre ranks 7th.

Under MIN, book 7 and book 18 both score **0.00** and disappear. On the tragedy corpus MIN separated cleanly: every Greek tragedy at MIN ≥ 0.48, every control, epic and didactic text at MIN ≤ 0.32.

So MIN was fixed as the primary metric in `PREREGISTRATION.md`.

1.5 Round 1: the pre-registration, and the failure

`PREREGISTRATION.md`, committed `dd32cd0`, 21 July 01:37.

Ovid's *Metamorphoses* was re-cut into 31 myth-level episode files. 17 were tiered in advance on Girardian grounds only, never on any score. 12 excluded on a stated rule (minimum 800 words). Narcissus and Phaethon were declared controls in `get_myths.py` **before the corpus existed**.

HIGH: Pentheus, Actaeon, Lycaon, Cadmus and the dragon. LOW (controls): Narcissus, Phaethon, the creation, the four ages, Daphne, Ceres, Aglauros.

The falsification conditions, written in advance:

- any LOW text scoring above any HIGH text on MIN
- HIGH and LOW overlapping so much that no threshold separates them
- the null test matching the Girard buckets' separation

Result, on the pre-registered metric:

TIER	N	MEAN MIN
HIGH	3	0.22
MID	6	0.19
LOW	6	**0.25**

No separation, and LOW is the highest. Two falsification conditions fired at once:

- `ovid2_aglauros_and_mercury`, a **declared control**, scores MIN **1.27**, beating every HIGH text.
- `ovid3_pentheus_and_bacchus`, **Girard's single most important persecution myth**, scores MIN **0.00**.

The unregistered metric, SUM, ordered the tiers exactly as predicted (HIGH 9.08, MID 6.06, LOW 3.52). `RESULTS.md` refuses to claim it, in its own words:

> "Switching metrics after seeing the result is precisely the move pre-registration exists to prevent."

1.6 The diagnosis, and it is arithmetic

MIN failed for a reason that can be written down exactly.

Model bucket hits as approximately Poisson with rate λ_b per thousand words. In a text of N tokens the expected count is λ_b·N/1000, and

$$P(\text{zero hits in } T) = \exp(-\lambda\_b \cdot N / 1000)$$

Recomputed from `results_corpus2_v2.csv`, 76 texts, lexicon v2:

BUCKET	MEAN RATE PER 1000
violence	4.757
marks	2.896
crisis	2.069
crimes	**0.949**

`crimes` is the rare bucket. Under this model:

TEXT LENGTH	EXPECTED `CRIMES` HITS	P(ZERO)
1,500 words (an Ovid episode)	1.42	**0.24**
10,000 words (a tragedy)	9.49	~0.000
17,180 words (the *Bacchae*)	16.31	~0.000

And this is the **repaired** lexicon. Under v1, `crimes` caught "murder" 138 times out of 297 real occurrences, so its effective rate was far lower and the zero probability far higher.

MIN takes the weakest bucket. If the weakest bucket is zero with high probability at episode length, **MIN saturates at 0.00 and carries no information at all.**

> **The finding, stated as a rule** The metric is scale-dependent. It was validated on texts six times longer than the ones it was then applied to. Any metric must be validated at the length it will actually be used on.

Pentheus is the sharp case, and it is a second, separate defect: `crimes 0.00` because Ovid describes the *sparagmos* as tearing and rending, and the bucket contained `murder`, `parricide`, `pollution`. **None of the words Ovid uses.** That is a lexicon coverage failure, not a failure of Girard's theory.

1.7 The repair, and the category error that was refused

Three defects were found in lexicon v1.

Defect 1, undercount. Exact matching missed inflections. `murder` caught 138 of 297; `kill` 105 of 393; `defilement` 4 of 54. The miss rate was **uneven**: common base forms (blood, king, death) caught at ~85%, rare Latinate nouns under 10%. So the v1 headline finding "each bucket is dominated by one generic word" was **partly an artifact of this bug**.

Defect 2, overcount. `cast` (422 hits) is mostly "cast a glance". `stone` (211) is mostly the material. `drive` and `driven` (327) are mostly wind and desire. Girard's stereotype is *expulsion*, which English builds phrasally, so these now require a particle.

Defect 3, the category error, and this is the one that matters.

`RESULTS.md`'s own next-action list said: fix the crimes lexicon, add the dismemberment vocabulary so it catches Pentheus.

That was refused, and the reason is quoted from `lexicon.py`:

> In Girard the crime stereotype is what the victim is ACCUSED OF; the sparagmos is what the MOB DOES. Tearing belongs in `violence`. Putting it in `crimes` would have corrupted the bucket to pass a test.

The tearing vocabulary went into `violence`. **The instrument was one line of reasoning away from being repaired to make its own failing example pass, and the thing that stopped it was a definitional distinction, not a statistical one.**

How the repair was made auditable. Defects 1 and 2 are corrections to *matching* and add no concepts. Genuinely new concepts are tagged `v2` in `lexicon.py` and can be excluded with `score.py --no-new`, so the effect of the revision stays separable from the effect of the bug fixes. Three tags: `v1` (original), `inf` (an inflection of a v1 word, no new concept), `v2` (new concept).

No stemming anywhere. Every concept lists its surface forms explicitly, because prefix matching is a disaster here:

PREFIX	WHAT IT CATCHES BY MISTAKE
`tear-`	"tears" (441), which is weeping, not rending
`rend-`	"render, rendered, rendering" (28)
`torn-`	"tornado"
`stone-`	the material, not the punishment

The cost is that adding a concept means writing out its forms. `lexicon.py` says that cost **is the point**: every match is auditable.

1.7b The anatomy of the repaired lexicon

Worth having in front of you, because every result in rounds 1 to 3 is a function of these word lists and nothing else. Sourced to *The Scapegoat* ch. 2.

Bucket 1, CRISIS. Loss of distinctions. Three sub-groups, and the third is the one people miss.

SUB-GROUP	FORMS
epidemic and sterility	plague, pestilence, contagion, famine, dearth, starvation, drought, blight, barren, sterile, sickness, disease, malady, death
undifferentiation proper	twin, double, likeness, semblance, indistinguishable, undistinguished
collapse of order	discord, strife, faction, lawless, civil war, civil strife

Bucket 2, CRIMES. What the victim is **accused of**, following Girard's four named categories.

CATEGORY	FORMS
(a) violence against those most criminal to attack	parricide, patricide, regicide, infanticide, murder, homicide, manslayer, bloodguilt
(b) sexual transgression	incest, rape, ravish, adultery, bestiality
(c) religious crime	sacrilege, blasphemy, profane, desecrate, impiety, godless
(d) contagion and "scalable" crime	poison, sorcery, witchcraft, treason, traitor, perfidy, betrayal
difference-erasing quality	pollution, defilement, abomination, unclean, taint, monstrous

> **The v1 failure in this bucket is instructive on its own.** It held criminological **labels**, and `parricide`, `patricide` and `regicide` fired **zero times in 200,000 words**. Narrative texts do not say "he committed parricide". They say "slew his father". So v2 adds bigrams: `slew his father`, `killed the king`, `slew her children`, `devoured his children`, `blood of his father`. And per `RESULTS-2.md`, **38 of 112 crime forms still never fire**, mostly those bigrams. They were written to catch what narrative actually says instead of the Latinate label, and they mostly did not.

Bucket 3, MARKS. Signs that select the victim, including the paradox that being too **high** is a mark as much as being too low: stranger, foreigner, alien, barbarian, outlander; cripple, lame, blind, maimed, deformed, misshapen; beggar, orphan, outcast; **king, queen, tyrant**; monster, scapegoat, accursed, cursed, marked, branded, stigma.

Bucket 4, VIOLENCE. All-against-one, and what the crowd does.

SUB-GROUP	FORMS
the crowd	crowd, mob, throng, multitude, rabble, populace
its unanimity	unanimous, unanimity, with one voice, one accord, as one man, all together, one and all, whole city, all the people
stoning	stoned, stoning, cast stones, hurled stones, pelted
expulsion, phrasal	expel, banish, exile, cast out, cast forth, cast down, drive out, drove out, driven from, thrust out, hunted out
killing	slay, slew, kill, slaughter, butcher, massacre
blood	blood, bloody, bloodshed, bloodstained, gore
the hunt	hunt, hunted, hunter, huntress, hound, hounded
sparagmos, the v2 Pentheus fix	tore, torn, tear, tearing, rend, rent, rending, dismember, asunder, sunder, mangle, limb from limb, in pieces, to pieces, piecemeal, torn apart

The exclusions, with their hit counts, which are the real design work

Each of these was **removed** or made phrasal, and each removal makes the instrument score **lower**, not higher.

EXCLUDED	HITS IT WAS CONTRIBUTING	WHY
bare `stone`, `stones`	**211**	nearly all the material. A stone wall, a stone tomb
bare `cast`, `drive`, `driven`	**749 combined**	"cast a glance", "driven by the wind". Girard's stereotype is expulsion, which English builds phrasally
`tears`, `tearful`	**462**	weeping, not rending
`render`, `rendered`, `rendering`	28	not rending
`ravishing`	n/a	almost always "ravishing beauty" in this corpus, not the act

> **This is the part of the instrument that is actual craft**, and it runs against the maker's interest every time. Nine hundred and sixty junk hits were deleted from a lexicon whose author wanted it to score high. The comment in `lexicon.py` on why there is no stemming anywhere ends: "Explicit lists mean every match is auditable. **The cost is that adding a concept means writing out its forms; that cost is the point.**"

1.8 Round 2: six predictions on 71 unscored texts

`PREREGISTRATION-2.md`, committed `00a9556`, 21 July 03:00.

The clean-test logic: the v2 lexicon was revised **after** seeing round 1 fail, so no result on `corpus/` or `myths_ovid/` can confirm it. Those are the texts that motivated the change. `corpus2/`, 71 scoreable texts never scored with any lexicon, is the clean test.

Disclosure written into the pre-registration itself: the v2 lexicon was smoke-tested on `corpus/` to check the code runs, and that output was seen before the file was written. No `corpus2/` file was scored.

Primary metric switched to **SUM**, stated in advance this time, because round 1 established MIN is scale-dependent and `corpus2` is length-heterogeneous (4,800 to 20,800 words).

Results:

	PREDICTION	RESULT
P1	genre ordering	**partial fail**
P2/P3	no named LOW outscores any named HIGH	pass
P4	Bacchae crimes > 0.5 and violence in top quartile	**fail**
P5	two Agamemnon translations within 25%	pass (17.2%)
P6	MIN non-zero for ≥75% of texts	pass (90.3%)
null	buckets beat frequency-matched random	pass, p = 0.0045

P1, the partial fail, and the confound it exposed.

GENRE	N	MEAN	MEDIAN	SD
MYTHOGRAPHY	8	**13.95**	13.91	4.70
TRAGEDY	12	**13.88**	14.97	4.31
OVID	11	9.85	9.47	2.28
EPIC	36	9.48	8.62	3.54
CONTROL	5	7.77	8.25	1.47

The floor held. Tragedy clears epic, Ovid and control. But **mythography ties tragedy**, and it was predicted to sit in the middle group.

The diagnosis: Apollodorus and Diodorus are bare plot summaries. No description, no lyric, no filler. Every word is event. **A rate per thousand words rewards narrative density, not persecution structure.** `hyginus_fabulae_1` ranks 1st of 71 at 1,208 words, the most compressed text in the corpus, and is treated as an artifact rather than a finding.

This is the same confound as the v1 "measures tragic register" result, in a new dress. **It is still unresolved.**

P4, the Bacchae test, and the failure that reframed the project.

`euripides_bacchae`: crimes **0.52** against a 0.5 threshold, a pass by 0.02, which is noise. violence **4.71** against a top-quartile cut of **6.43**, a clear fail. Overall rank **35 of 72**. Dead centre.

But **the lexicon fix worked.** The sparagmos vocabulary fires exactly as designed: `tore` 7, `torn` 7, `tear` 3, `rent` 2, `rending` 2, `asunder` 2, plus `limb from limb` and `in pieces`. Roughly 24 sparagmos hits, none of which v1 could see.

The failure is dilution. The *Bacchae* is 17,180 words and the lynching is one messenger speech. A whole-text rate spreads one collective murder across seventeen thousand words. Meanwhile `quintus_smyrnaeus_2` scores violence 8.32 because it is continuous battle narrative, with diffuse violence everywhere, no victim, no crowd and no mechanism.

> **The instrument prefers a uniformly violent text to a text containing one lynching. That is precisely backwards from Girard**, whose mechanism is a localized *event*. Rate per thousand cannot represent an event.

P5, the translation control, and why it is the most load-bearing number in the project.

Same play, two translators. `aeschylus_agamemnon_murray` (Gutenberg, Murray) scores 14.82. `corpus/aeschylus_agamemnon` (Theoi, a different translator) scores 12.64. **17.2% apart**, inside the 25% bound.

Had this failed, every cross-text comparison in the entire project would have been void, because the instrument would have been measuring translators' word choices rather than texts. It has been run on **exactly one paired play**.

Bucket composition, re-derived after the bug fix:

BUCKET	TOP WORD	SHARE, V2	WAS, V1	DEAD FORMS
crisis	`death`	60%	68%	9 of 50
marks	`king`	44%	54%	10 of 45
violence	`blood`	21%	36%	6 of 83
crimes	`murder`	12%	43%	38 of 112

Fixing the matching materially reduced single-word dominance in every bucket. But `crisis` is still 60% the word "death" and `marks` is still 44% the word "king". **`marks` largely detects whether a text is about royalty.** Sophocles' *Oedipus the King* scores marks 9.58, the highest in the corpus, and the title character is addressed as "king" throughout.

1.9 The null test, and the only surviving positive result

TRAGEDY versus CONTROL, groups fixed in the pre-registration before scoring, 2000 frequency-matched random word-sets:

```
Girard buckets:  +6.10 per 1000 words
random sets:     mean -0.16,  sd 2.95
beat Girard:     8 / 2000        p = 0.0045
```

What makes this one citable and the two earlier ones not. The first null test chose its "persecution texts" list *after* seeing the grid. The second used total rate, not the pre-registered MIN, switching metrics after seeing the result. This one had **both the groups and the metric committed in advance.**

What it means, in the pre-registration's own words, and this framing was written before the result:

> A confirmation here means only that the four stereotypes co-occur more in persecution-shaped texts than in others, **the first half of Girard's claim, and the uncontroversial half.**

1.10 Correction 1: the truncated play, and the page that was excluded

Found 22 July, by accident, while running a speaker diagnostic for something else.

The defect. `get_corpus2.py` fetched Theoi pages but never followed their pagination. Theoi splits long pages and links the remainder as `CONTINUED`. One text hit that split:

TEXT	AS SCORED	ACTUAL	MISSING
`seneca_hercules_oetaeus`	11,119 words	20,621	**46%**

The missing half is not arbitrary. The file broke off after Hyllus' speech, so Deianira's death, Hercules on the pyre and the apotheosis were all absent. Hercules spoke exactly one line. Alcmena and Philoctetes never appeared. **The scorer was reading a play with its sacrifice removed.**

> *Why this is worse than a missing text: a missing text is visible. A half-play still scores, and the score looks unremarkable. It sat mid-table in `RESULTS-2.md` and nothing flagged it.*

How it was nearly missed. The speaker diagnostic classified the string `CONTINUED` as a phantom one-off speaker and dropped it into a list of ignored singletons. The pagination marker was in the very first diagnostic output and **was read past twice** before anyone noticed what it meant.

A first pass claimed five truncated files. That was wrong: the search was case-insensitive and matched the ordinary word "continued" in prose. Only one text was affected. `CORRECTION-1.md` records the overstatement, on the ground that a correction which overstates its own scope is not a correction.

Effect on the result. Rescored, the text moves from 15.12 to **14.67**. The full text scores *lower* than the truncated one. The expectation going in was the opposite. The restored half is largely lament and apotheosis, which dilutes the rate. TRAGEDY mean 13.91 → 13.88. P1 was already a partial fail on a 0.04 margin and remains a fail on a 0.07 margin. **No verdict changed.**

And then the part that matters

The same audit found `DiodorusSiculus4D`, a real Theoi page absent from the fetcher's page list. It was fetched, and it is on disk, and it is **excluded**.

Because Diodorus is MYTHOGRAPHY, the group P1 loses to. Folding it in:

```
MYTHOGRAPHY  13.95  →  13.67       vs  TRAGEDY  13.88
P1                FAIL  →  PASS
```

> A text discovered after seeing the results, added to the comparison group, flips a failed pre-registered prediction into a passing one. **It would have been entirely defensible as a bug fix**, and it is exactly the move pre-registration exists to make impossible.

The rule this produced, and it is the second discipline rule of the project:

> **The sample is part of the registration, not just the lexicon.**

`diodorus_siculus_4d` is available to any future round whose predictions are committed before it is scored. It has no bearing on that one.

> *Hold onto this. Part VII shows that this rule was broken two weeks later, by the same project, without anyone noticing.*

1.11 Round 3: the localization test

`PREREGISTRATION-3.md`, committed `fc094bb`, 22 July 00:55.

The reasoning. Rounds 1 and 2 measured **saturation**. Girard does not claim saturation. He claims the mechanism is an **event**: a single episode where the crisis, the accusation, the marks and the collective violence all converge.

The statistic

Slide a window over the token stream, score the four buckets inside it, take the **min** of the four (co-occurrence), then take the **max** over all windows (does such a passage exist anywhere).

$$peak\_min(T) = max\_i\ [\ min\_b\ c\_b(W\_i)\ ]$$

with W_i the window starting at token *i*, window length 1000, stride 250. The window is exactly 1000 tokens, so a raw count *is* a rate per thousand.

Implemented in `nulltest.py` with a cumulative sum so the cost is linear in text length rather than quadratic:

$$c[b, j] = \Sigma\_\{k<j\}\ ind[b, k]\ ;\ counts = c[:, ends] - c[:, starts]$$

The null, which is the actual contribution

A high `peak_min` on its own is uninterpretable. A text with dense persecution vocabulary throughout will produce one by chance. So:

Null A, localization, primary. Take a uniform random permutation π of the token positions and recompute the statistic on the permuted indicator array. 500 shuffles per text, seed `20260722`.

$$z = (\ obs - mean(null)\ )\ /\ sd(null)$$

$$p = (\ \#\{\ s : peak\_min\text{^}(s) \geq obs\ \} + 1\ )\ /\ (\ 500 + 1\ )$$

The plus one in both numerator and denominator is the standard correction that keeps an empirical p from ever being exactly zero.

> **What Null A holds constant and what it destroys.** It preserves the text's entire vocabulary exactly, so every bucket's total count and whole-text rate are unchanged. It destroys position and nothing else. So it isolates **locality**, which is precisely the thing round 2's density confound could not be separated from. A text can be drenched in all four stereotypes and still fail it.

One implementation detail that is load-bearing. Multi-word forms are marked at the position of their **first token**, so a permutation moves each phrase hit as a unit. Shuffling raw words instead would dissolve every phrase and make the null trivially easy to beat.

Null B, word identity, secondary. Four random lexicons matched to the Girard buckets on surface-form count and corpus frequency, scored on the unshuffled text.

The rule attached to the round

`window.py` was written after seeing round 2 fail, so its settings (window 1000, stride 250, min then max) are **post-hoc**. They were frozen in the pre-registration and the file says:

> If 1000/250 gives a null result, that is the result. A window size chosen because it made the *Bacchae* pass would void this round exactly as tuning the lexicon would.

1.12 Round 3 result: total failure, and a bad null design

`RESULTS-3.md`, run `4188308`, 22 July 03:22.

	PREDICTION	RESULT
P1	≥7 of 9 named HIGH localize	**fail, 0 of 9**
P2	Bacchae at p < 0.01, peak in second half	**fail, p = 0.66**
P3	≤1 of 6 named LOW localizes	pass, 0 of 6, **vacuously**
P4	TRAGEDY z exceeds EPIC z by ≥1.0	**fail, gap 0.30**
P5	≥6 of 9 HIGH beat Null B	**fail, 0 of 9**

Across all 71 texts:

```
z          mean +0.07    sd 1.33    median 0.00
p ≤ 0.05   4 of 71 (5.6%)      expected by chance  3.6
```

The distribution is centred on zero. The four texts that clear the threshold are exactly the number chance predicts, and **not one of them is a named HIGH text.**

> `peak_min` is therefore not detecting episodes. **It is reporting how much bucket vocabulary a text contains. Word order is irrelevant to it.**

P3 is not a success and the file says so. The named LOW texts fail to localize, but so does everything. A prediction that "these texts show no effect" is worthless in a corpus where no text shows the effect.

The Bacchae, twice. Round 2 ranked it 35 of 72 on whole-text rate. The diagnosis was that the metric was wrong. Asked at episode scale with a null, it comes out at **p = 0.66**. Its peak window sits at **54% of the text**, which is where the sparagmos is, so the locator is pointing at the right passage. That passage simply does not contain a convergence beyond what the play's vocabulary produces by chance.

> Girard's central example has now failed two pre-registered tests on two different operationalizations. The objection this project exists to answer is that the exemplars are selected by interpretation rather than by any property the text carries, and **after two rounds that objection is stronger, not weaker.**

The power problem, stated as a limitation and not an excuse

`peak_min` is a minimum over four counts in a 1000-word window. Observed values run 0 to 5, **median 2**. For 64 of 71 texts the weakest bucket contributes three hits or fewer in the best window.

This is a statistic built on counting to two. Whether the null reflects an absence of localization or a statistic too coarse to detect it is not settled by this round. The file's own words: the test as specified had little power, and **this was foreseeable before running it and was not foreseen.**

The rule that follows: no re-running with a larger window until something passes. A different operationalization is a new pre-registration with its own commit.

Null B was a bad design, and the reason is structural

Frequency-matched random lexicons scored **higher** than Girard's buckets. For the *Bacchae*: observed 2.0, null mean **4.41**.

That is not evidence about Girard, and the mechanism is worth understanding because it generalises:

> A **minimum** over four sets is raised by sets that are uniformly spread and lowered by sets that are clumpy. Random words matched on individual frequency spread across topics and so appear at more uniform rates, which raises a minimum. **Any thematically coherent lexicon is clumpier and will lose this comparison.** Null B would fail identically for any four topical word sets. It tests nothing and its result should not be cited.

Recorded in `RESULTS-3.md` as "my error, in the pre-registration."

1.13 Round 4: the dissent test

`PREREGISTRATION-4.md`, committed `9310c30`, 11 August 02:07. Outstanding since July, called by `FINDINGS.md` §6 "the most important unbuilt piece of the project."

Why the four buckets could never do this

`FINDINGS.md` §5, and this was known from the beginning:

> Girard **predicts** the Passion scores high on all four. The Gospels are built from the same persecution material as myth. So a high score there confirms the first half of his claim and discriminates nothing. **The four buckets are structurally incapable of separating Gospel from myth. This is the design, not a limitation to patch later.**

The discriminating claim, in Girard's own words

I See Satan Fall Like Lightning, Introduction, p. 3:

> "not merely the innocence of the victims versus their guilt, but the fact that, in mythology, **no one ever questions this guilt**. In the Gospels, the revealing account of scapegoating emanates not from the unanimous crowd but from a dissenting few. In mythology no dissenting voice is ever heard."

Same book, p. 121:

> "a rebellious minority, a small group of dissidents that separates from the collective violence of the crowd and **destroys its unanimity**. This dissident minority has no equivalent in the myths."

The differentia is not innocence versus guilt. It is whether the text contains a voice that breaks the unanimity. And "no dissenting voice is ever heard" is a **universal negative**, so a single genuine counterexample falsifies it. That is the falsifiable shape the project had wanted since round 1, and it is Girard's own formulation rather than one imposed on him.

Predicted ordinal series:

$$\textit{myth} < \textit{tragedy} < \textit{Gospel (in dissent / broken unanimity)}$$

The instrument, and the tag that carries the round

A fifth bucket, `dissent.py`, same architecture. Two provenance tags:

TAG	SOURCED TO	CONCEPTS
`core`	p. 3, questioning of guilt	innocent, guiltless, no_fault, done_nothing, unjust, wrongly, without_cause, false_witness, slander, protest, object, denounce, defend, plead, speak_for
`unan`	p. 121, broken unanimity	division, dissent, some_others, not_all, withheld, spare, minority

Exclusions, each of which makes the hypothesis harder to confirm:

EXCLUDED	WHY
`have mercy`	frequent in the Gospels and almost always a healing petition, not dissent
`rebuke`	dominant sense here is rebuking demons, wind, disciples
bare `witness`	"bear witness" is Johannine testimony, not objection. Only `false witness` kept
`condemn`	that is the crowd's act, not the dissenter's
bare `refuse`	too generic. Only explicit "consented not" forms kept

The sample

`tiers.py`, committed with the pre-registration.

TIER	N	WHAT IT IS
GOSPEL	4	Matthew, Mark, Luke, John, ASV 1901
TRAGEDY	23	Greek and Senecan verse drama
MYTH	71	mythography 34, epic 37

Translation matching, decided in advance. The corpus is almost entirely 1900 to 1921 English. The ASV is 1901, literal, public domain. The KJV was rejected as three centuries off register. The pre-registration states this choice "is made now and is not revisited after seeing results."

Corpus checksum recorded: `asv.txt` sha256 `aecff150eb9c…`. Word counts extracted: Matthew 23,427, Mark 14,855, Luke 25,655, John 19,005. The file states that a different checksum in future is a changed sample and needs saying out loud.

The contamination disclosure, written before scoring

> I knew before writing `dissent.py` that John repeats a division formula ("there was a division among them") and that the Passion contains "I find no fault in this man." Those are the two most obvious surface forms for the sourced construct, so `no_fault` and `division` are at risk of being motivated rather than derived.

Mitigations declared: `unan` is tagged separately, every result is reported core-only, and D4's frequency-matched null absorbs a lexicon that only works because its words happen to be frequent in one tier.

1.14 The mathematics of round 4

Mann-Whitney U with midranks. Pool the two groups, rank them, average the ranks of tied values, then

$$U\_a = R\_a - n\_a(n\_a + 1) / 2$$

where R_a is the sum of group a's ranks in the pooled sample.

Significance by label permutation, not by normal approximation.

$$p = (\#\{ perm : R\_a^perm \geq R\_a^obs \} + 1) / (NPERM + 1), NPERM = 200{,}000$$

`score_dissent.py` states the reason, and it is not a workaround:

> scipy is not installed in .venv, so Mann-Whitney significance is obtained by label permutation rather than a normal approximation. **At n=4 for GOSPEL a normal approximation would be wrong anyway, so this is the better test.**

The vectorisation, which is exact and not an approximation. Ranks are computed once on the pooled sample, then labels are permuted **over the ranks**. Since the rank vector does not depend on which labels are assigned, this is algebraically identical to re-ranking every permutation, and about three orders of magnitude faster.

The D4 null. Bin the corpus vocabulary by $\lfloor \log_2(\text{count}) \rfloor$. For each dissent unigram, draw a replacement from the same frequency bin, excluding real lexicon words. 2000 draws. The statistic being nulled:

$$separation = median(GOSPEL\ rates) - median(MYTH\ rates)$$

1.15 Round 4 result: nominally 3 of 4, actually nothing

`RESULTS-4.md`, run `0c765d0`, 11 August 02:17.

TIER	N	MEDIAN RATE /1000	MEAN	HITS	WORDS
GOSPEL	4	**1.387**	1.349	112	83,180
TRAGEDY	23	0.667	0.795	236	320,887
MYTH	71	0.547	0.594	365	612,835

	RESULT	
D1 ordinal medians	0.547 < 0.667 < 1.387	**PASS**
D2 GOSPEL > MYTH	U = 268, p = 0.00034	**PASS**
D3 TRAGEDY > MYTH	U = 1027, p = 0.032	**PASS**
D4 frequency-matched null	591/2000 beat it, p = 0.296	**FAIL**
D5 Prometheus top 3 of 23	rank **13**, rate 0.656	**FAIL**
D6 no myth above lowest Gospel	7 of 71 exceed it	**FAIL**

That table read alone says 3 of 4 on the primary. It is not 3 of 4, because two sensitivity analyses were registered in advance and both break a passing prediction.

S1, core only. Dropping the `unan` tag, the half disclosed in advance as at higher risk of being motivated:

```
core only:   GOSPEL median 0.402   TRAGEDY 0.503   MYTH 0.325
             ordinal BREAKS (tragedy now above Gospel)
             GOSPEL > MYTH  p = 0.293      ← D2 gone
             TRAGEDY > MYTH p = 0.005      ← survives
```

> The concepts that state Girard's actual claim, innocence, protest, denounce, defence, plead, **carry none of the Gospel effect. The contaminated half carries all of it. This is the disclosure coming true.**

S2, epic excluded. Restricting MYTH to mythography, n = 34:

```
MYTHOGRAPHY median 0.674    EPIC median 0.382    TRAGEDY median 0.667
GOSPEL  > MYTHOGRAPHY  p = 0.005
TRAGEDY > MYTHOGRAPHY  p = 0.388      ← D3 gone
```

Mythography scores above tragedy. D3 passed only because 37 epic texts drag the myth tier down. The myth < tragedy half of the ordinal series does not exist. What exists is **epic < everything**.

Why D4 failed, and it is not what it looks like

The null matched **unigrams only**, 45 of the 64 lexicon unigrams present in the corpus. Splitting the instrument in half explains everything:

HALF	GOSPEL MEAN	TRAGEDY MEAN	MYTH MEAN
unigrams	0.367	**0.639**	0.486
phrases	**0.982**	0.156	0.108

The unigram half runs backwards. On unigrams the Gospels score *lowest* of the three tiers. The observed unigram separation is **−0.075 per 1000**, which is why 591 random word-sets beat it.

> D4 did not fail because the instrument is weak. It failed because **the null tested the half of the instrument that carries no signal**, and the entire effect lives in the phrase half, which **has never been tested against a null at all.**

Recorded as a defect in the test design. The corrected version, a phrase-aware frequency-matched null, is **round 5 with its own commit**, not a rerun folded into this one.

What is actually firing

Top six discriminators by GOSPEL-minus-MYTH mean rate:

CONCEPT	TAG	FORMS DOING THE WORK	G−M
some_others	`unan`	"others said", "some said", "some of them"	+0.328
spare	`unan`	"let him go", "release him", "loose him"	+0.229
false_witness	`core`		+0.130
division	`unan`		+0.118
not_all	`unan`		+0.082
no_fault	`core`	"no crime"	+0.076

Four of the top five are `unan`. And `defend` runs the wrong way: myth 0.084, Gospel 0.000.

John, the highest-scoring Gospel: `others said` ×7, `division` ×3, `some of them` ×3, `not all` ×3, `no crime` ×3, plus one each of `let him go`, `loose him`, `release him`, `without a cause`. **Sixteen of its 26 hits are the "some said / others said" narrative formula.**

Whether that formula is Girard's broken unanimity or **a translator's reporting convention** is exactly what a whole-text rate cannot distinguish.

1.16 The D6 counterexamples, and the escape hatch that was refused

The pre-registration committed that any myth outranking the Gospels gets read and named. Seven do.

`orphic_hymns_2`, 5,812 words, rate 2.24, the corpus maximum. Firing: `defend` ×5, `blameless` ×4, `unjust` ×3, `spare the` ×1. **Every one is a hymnic epithet or a petition addressed to a god.** "Blameless" is a standard Orphic epithet. "Defend" is a prayer. None of it is a voice questioning a victim's guilt. Same class of error as `king` inflating `marks` in round 1.

`ovid_metamorphoses_9`, 9,633 words, rate 1.76. `plead` ×3, `object` ×3, `unjust` ×2, `divided` ×2. Ordinary rhetorical verbs inside speeches, and `divided` in its physical sense.

`hyginus_fabulae_2`, 613 words with **one hit**. A rate artifact on a short text and nothing else.

The trap, named in the file so it would not be walked into

> The obvious next sentence is "so these are false positives, therefore Girard's universal negative survives." **That is the escape hatch and it is refused.** What the inspection shows is that the instrument produces counterexamples that are not counterexamples, which means **it cannot test a universal negative at all, in either direction.** It could not have found a real counterexample and distinguished it from an Orphic epithet. Girard's claim is neither supported nor falsified here. **It is untested.**

1.17 One root cause for two rounds

Round 3 failed because `peak_min` is invariant to word order, so it could not locate episodes.

Round 4 fails the same way one level up. Girard's claim is not "Gospels contain more division-language than myths." It is that **at the moment of collective violence**, a voice separates from the crowd. John's three `division` hits are at 7:43, 9:16 and 10:19, and they are disputes about **who Jesus is**, not dissent from a killing. A whole-text rate cannot tell those apart, because position in the narrative is exactly the information a rate discards.

> **Two failed rounds, one cause: whole-text rates cannot test claims about narrative structure.** That is the finding of round 4 and it is worth more than a pass would have been.

Also settled by D5: **`Prometheus Bound` does not replicate.** `FINDINGS.md` §2 had flagged it as the one tragedy scoring 0.00 on the four-bucket MIN and "the one tragedy that unambiguously sides with its victim", and forbade the claim until tested. Tested, rank 13 of 23. The anomaly was four-bucket vocabulary, not dissent. That note is closed.

1.18 Round 5: the lost null test, reconstructed

Run **14 August 2026**, after this document's first draft named the missing script as the project's highest-value outstanding item. Spec `6d11a56`, results `4e81d36`.

The discipline, applied to the reconstruction itself

The hazard is specific. The reconstruction has free parameters the committed record does not fix, and there is a target number in view. **With several free parameters and a known answer, a reconstruction can be walked anywhere, one defensible choice at a time.** That is Correction 1 again, moved from the sample to the procedure.

So `NULLTEST-2-SPEC.md` enumerates every free parameter and every defensible resolution of each, fixes the verdict rule in advance, and commits to running the **whole grid** rather than choosing afterwards. Committed before `nulltest2.py` existed on disk, and the manifest proves it:

COMMIT	CONTAINS
`6d11a56` spec	`NULLTEST-2-SPEC.md` **only**
`4e81d36` results	`RESULTS-NULL2.md`, **`nulltest2.py`**, `results_nulltest2.csv`

Finding zero, recorded before any null ran

The **deterministic** half should have been trivially reproducible. It is not. Eight candidate definitions of "separation" were computed against the committed CSV, and the search is declared in the spec rather than reported as a lucky hit. Mean of per-text rates, TRAGEDY minus CONTROL, is the only near candidate at **+6.1067**. Restoring the pre-correction truncated text gives **+6.1442**.

> ***The committed +6.10 matches neither.*** *And `CORRECTION-1.md` states the correction moves the effect size "6.10 → ~6.07", where the CSV says the movement is **6.144** → **6.107**. Direction right, both endpoints displaced by about 0.037. Small, checkable, and still unexplained.*

The five free parameters

	PARAMETER	RESOLUTIONS RUN
F1	corpus state	`corrected` / `truncated` (pre-Correction-1)
F2	matching procedure	`band` (round 3's) / `log2` (round 4's)
F3	phrase handling	`phrase-aware` / `unigrams-only`
F4	frequency pool	`full` 76 texts / `groups` 17 texts
F5	seed	three

2 × 2 × 2 × 2 × 3 = 48 configurations, each of 2000 draws.

The verdict, by the rule fixed in advance

> ***UNSTABLE.*** *24 of 48 give p < 0.05, 24 give p ≥ 0.05. The spec says that when the grid straddles, "p = 0.0045 cannot be cited, whatever the majority of configurations say," and no configuration may be selected afterwards.*

But the grid does not scatter. It splits, on one parameter

ARM	N	P RANGE	NULL MEAN	NULL SD
`unigrams-only`	24	**0.00050 to 0.00700**	−1.93 to +0.40	**2.46 to 3.58**
`phrase-aware`	24	**0.15892 to 0.54123**	−5.28 to +7.14	10.48 to 13.09

Corpus state, matching procedure, pool and seed move p between 0.0005 and 0.0070 and do nothing else. That is striking stability across four free parameters.

The lost script is identified, and not by its p-value

This is the part that matters. `RESULTS-2.md` committed **three** numbers in July and two of them have nothing to do with the p-value:

COMMITTED	RECONSTRUCTED, UNIGRAMS-ONLY	INSIDE?
null mean **−0.16**	−1.93 to +0.40	**yes**
null sd **2.95**	2.46 to 3.58	**yes**
beat count **8 of 2000**	0 to 13	**yes**
p **0.0045**	0.0005 to 0.0070	**yes**

Against the phrase-aware arm the committed sd of 2.95 is not merely outside the range, it is **four times smaller than anything that arm produces.**

> **So the original matched unigrams only, and under that procedure the committed result reproduces.** The identification rests on a mean and a standard deviation committed three weeks earlier, which could not have been fitted here.

The truncation account also checks out. The sanity reconstruction gives 15.11 against `CORRECTION-1.md`'s committed 15.12 truncated, and 14.67 against its 14.67 corrected.

Why the phrase-aware arm does not count against round 2

Two measurements, **neither of which involves a null distribution or a p-value.**

Phrases carry 0.86% of round 2's effect. Observed separation is +6.1043 with phrases and +6.0518 without. The four-bucket lexicon is **99.1% unigram-driven**, so a unigram-matched null tests essentially the whole instrument.

> **This is the precise inverse of round 4**, where the phrase half scored GOSPEL 0.982 against a unigram half at 0.367 that ran **backwards**. There a unigram-only null tested the half with no signal, and `RESULTS-4.md` correctly called that a defect. **The same procedure is appropriate in one round and defective in the other, and what decides it is a property of the lexicon, not of the null.**

And the phrase-aware arm is not a frequency-matched null. Inherited from round 3's `build_pseudo`, it matches a multi-word form on the corpus frequency of its **first token**:

PHRASE	OCCURRENCES	FIRST-TOKEN OCCURRENCES	INFLATION
`as one man`	0	4,397	4397×
`with one voice`	3	8,294	**2765×**
`to pieces`	16	16,734	1046×
`in pieces`	13	11,199	861×
`one accord`	5	1,924	385×

Median inflation across all 68 phrases: 113×. In aggregate a draw replaces **206 phrase-occurrences with words totalling 65,137**, a 316-fold inflation. Which is visible directly in the output: the null sd jumps from about 3 to about 11 the moment the arm is switched on.

> ***Third instance of the same null-design failure in this repository.*** *Round 3's Null B favoured uniformly-spread sets over clumpy ones by construction. Round 4's D4 nulled the half of the instrument with no signal. Round 5's phrase-aware arm matches a 3-occurrence phrase to an 8,294-occurrence word.* ***Every one is a mismatch between what the null holds constant and what the statistic is sensitive to.*** *Three times now.*

The part where the discipline gets tested, and reported rather than dodged

The verdict rule was written before running. The literal answer is UNSTABLE. And the argument above discounts half the grid, produced **after** seeing which half was which. On its face that is the excluded Diodorus page one level up.

`RESULTS-NULL2.md` §4 states the three things that make it admissible, so a reader can check rather than trust:

1. **Both measurements are independent of every p-value.** The 0.86% and the 316× come from the observed corpus. Neither uses a null draw. **Both would read identically if the p-values had come out reversed.**
2. **The argument is directional, not selective.** It does not say the unigrams-only arm is right because it passed. It says the phrase-aware arm is an invalid null **whatever it produces**.
3. **It does not restore the claim.**

> **And the honest concession, recorded in the results rather than patched into the spec: the verdict rule was written badly.** It anticipated that the arms would be *different valid choices* and provided for a scatter. It did not anticipate that one arm would be an **invalid instrument**, and gave no way to say so. That is a defect in my own pre-registration, discovered by running it.

What round 5 establishes, and what it does not

Established. The lost script is identified. Under that procedure the committed result reproduces at p between 0.0005 and 0.0070 across twelve combinations of the remaining parameters, against a committed 0.0045. Round 2's unigram matching was appropriate for round 2's lexicon. Correction 1's truncation account is faithful to two decimals.

Not established. A genuinely phrase-aware null has still never been built, for any round of this project. The correct construction matches each phrase against *other phrases* of comparable corpus frequency, which no script here does. Until it exists, the 0.9% of round 2 that lives in phrases is untested and the **100% of round 4** that lives in phrases is untested. The spec's own verdict rule returns UNSTABLE. And the committed effect size is still not exactly reproducible.

Round 6 is that null, and it goes to **round 4 first**, because that is where the phrase half carries the entire effect and where the answer could still go either way.

> **Built and run the same day. See §1.19.** The phrase half passed at **0 of 2000**, and a translation control that had never been applied to this construct passed at 20.4%. So round 5's diagnosis was right and the thing it said was untested is now tested. **What it did not change: S1 still destroys the tier result, so the effect remains entirely inside the contaminated half.**

> And what none of it touches. `PREREGISTRATION-2.md` said it before the result existed: a confirmation here means only that the four stereotypes co-occur more in persecution-shaped texts than in others, **the first half of Girard's claim, and the uncontroversial half.** Rounds 3 and 4 tested the systematic half and both failed. **Restoring one number does not move that.**

1.19 Round 6: the phrase-aware null, and the translation control

Run **14 August 2026**, immediately after round 5 identified the omission and sent it here. Pre-registration `4d9f8cf`, containing that file and nothing else. Results `d364559`, where `score6.py` arrives.

Four passes, no falsifier fired. The first clean round since round 2. And the reason it changes less than that sounds is fixed in the pre-registration, before anything ran.

What round 4 left open, exactly

HALF OF THE INSTRUMENT	GOSPEL	TRAGEDY	MYTH
unigrams	0.367	**0.639**	0.486
phrases	**0.982**	0.156	0.108

D4 matched **unigrams only**, so it nulled the half that runs backwards, and 591 of 2000 random sets beat an observed separation of **−0.075**. Round 4's actual finding, the roughly ninefold phrase separation, had **never faced a null**.

N1, the null that should have been built three rounds ago

Every one of the 53 lexicon phrases (34 bigrams, 19 trigrams) replaced by a real n-gram **of the same word length**, matched on corpus frequency within a 1.5× band, drawn from `corpus3`'s own inventory of **378,412 distinct bigrams and 808,967 distinct trigrams**. 2000 draws.

	OBSERVED	NULL MEAN	NULL SD	BEAT	P
P1 full lexicon	+0.8396	−0.1203	0.2836	14 / 2000	**0.0075**
P2 phrases only	**+0.9184**	+0.0722	0.1330	**0 / 2000**	**0.0005**

F3 did not fire at 0.0%. All 53 phrases found a match, so this null is not under-powered in the way round 3's `peak_min` was, where the statistic was a minimum over four counts with median 2.

Not one draw of two thousand beat it. Random phrases of the same shape and the same rarity do not concentrate in the Gospel tier. These do.

N2, the control that answers the question round 4 could not

`NEXT.md` had named register as "the largest threat to this test," and `RESULTS-4.md` ended on it: is "some said / others said" Girard's broken unanimity, or **a translator's reporting convention?**

A second translation of the same four gospels has been in this repository since Layer 1. `interlock/corpus/canonical/` holds the KJV, 84,025 words against the ASV's 83,180. Same text, different translator, **290 years apart**.

GOSPEL	ASV /1000	KJV /1000	DELTA
Matthew	0.724	0.674	−6.9%
Mark	1.141	0.856	**−25.0%**
Luke	0.854	0.731	−14.4%
John	1.207	1.255	+3.9%
median	**0.998**	**0.794**	**−20.4%**

Inside the 25% bound, well inside the 40% falsifier. For scale, round 2's paired *Agamemnon* translations differed by **17.2%**, and `RESULTS-2.md` calls that control "the most load-bearing" in the project. And GOSPEL > MYTH survives the whole substitution at **p = 0.00064**.

> **So the phrase family is not the ASV's private habit.** Translators working independently, in a different English, three centuries apart, put comparable amounts of it into the same four texts. That was a live and serious explanation of round 4's entire result and it is now substantially weakened.

Though it is translation-sensitive, and Mark says so. Mark alone drops 25.0%, exactly at the bound, and the median moves 20.4% in one direction. A pass, not a clean bill of health.

What has not changed, which is the section that matters

The round-4 sensitivity analyses were rerun on the KJV substitution:

```
S1  core only,      KJV GOSPEL > MYTH ........... p = 0.254
S2  epic excluded,  KJV GOSPEL > MYTHOGRAPHY ..... p = 0.011
```

Round 4's S1 took GOSPEL > MYTH from p = 0.0003 to p = 0.293. On the KJV it goes to **0.254. The finding is unchanged.**

> ***Everything round 6 established is about the unan half****, the half* `PREREGISTRATION-4.md` *§7 disclosed in advance as at higher risk of being motivated, because I already knew John repeats a division formula before the lexicon was written.*
>
> *The* `core` *concepts are the ones that state Girard's actual claim: innocent, guiltless, no fault, done nothing, unjust, false witness, protest, object, denounce, defend, plead.* ***They carry none of the effect, in either translation.*** *What survives two nulls and a translation swap is division, some-said, not-all, spare-him and withheld consent.*

And the ceiling fixed before running holds. `PREREGISTRATION-6.md` §5 stated that N1 cannot distinguish Girardian dissent from a translation convention, because a convention is also a set of rare n-grams concentrated in one tier; that N2 removes one translator's habit and not the class; and that neither test touches round 4's root cause, since a whole-text rate discards position and both tests here are whole-text rates.

> There was a clean sweep of P1 to P4. **That finding stands. Round 4's verdict is unchanged**, because S1 and S2 were its controls and S1 still breaks it.

The exact amendment round 4 now earns

`RESULTS-4.md` §6 said of the phrase half: "**Untested against a null and confounded with translation idiom. Not a result. A thread.**"

- **Tested against a null.** 0 of 2000.
- **Not confounded with translation idiom**, on the one alternative translation available, at 20.4%.
- **Still not a result about dissent.** It is a result about a phrase family that lives entirely in the tag disclosed as contaminated, and that a whole-text rate cannot locate at the moment of collective violence.

One deviation, declared rather than absorbed

2000 rescorings of 98 texts is unrunnable if every text is rejoined into a string on each draw, so `score6.py` pre-counts unigrams, bigrams and trigrams per text.

The two counters disagree on **exactly one text of 98**. `euripides_alcestis` contains "not all" twice in succession; `str.count` is non-overlapping and consumes the trailing space, so it counts one where the n-gram counter counts two. Reference 0.6040 against 0.6589, a gap of 0.055 per 1000.

The n-gram counter is the more correct count. The reference is what round 4 committed. Decision applied throughout: **observed values use the reference scorer**, so they stay exactly comparable with round 4, and **null draws use the fast path**, where random n-grams cannot produce that adjacency. The script verifies the tier medians against `RESULTS-4.md`'s committed **1.387 / 0.667 / 0.547** and reproduces them exactly before drawing anything. No tier median moves either way.

1.20 Rounds 7 and 8: localization, tested twice because the first null was wrong

The bottleneck since round 3, and the thing rounds 3 and 4 both died for lack of. Run **14 August 2026**, in two rounds, because the first one used a null that could not work.

Round 7: the statistic, and why it has no window

For each dissent hit, the distance in tokens to the nearest **violence** hit; the text's value is the median of those. Girard's claim predicts it is **small**: the dissenting voice speaks at the lynching, not three chapters away.

No window anywhere. A windowed version has round 3's disease. At a 1000-token window the median text holds **0.63** dissent hits, which is worse than counting to two. A hit-level distance uses every hit and needs no binning. That calculation was run **before** the pre-registration was written, which is the thing `RESULTS-3.md` says was not done in round 3.

Round 7 failed, and two falsifiers fired

	PREDICTION	RESULT
L1	≥3 of 4 Gospels co-locate	**FAIL, 0 of 4**
L2	ordinal series on z	**FAIL**
L3	the *Bacchae*	**FAIL**, p = 0.888
L4	above chance overall	pass, **vacuously**, 4 of 66 against 3.3, none a Gospel

M3 fired, a power falsifier that exists precisely because round 3 had none. Median z **−0.012**, inside the ±0.2 band fixed in advance, so the round is reported **under-powered** rather than as an absence. **M1 fired**: myth passed at 7.5% against the Gospels' 0.0%.

And the disclosed expectation came true. `PREREGISTRATION-7.md` §8 stated before scoring that I expected L1 to fail, because `RESULTS-4.md` had already placed John's three `division` hits at 7:43, 9:16 and 10:19, chapters from the passion. **John came out at p = 0.876.** Written down first.

The reason it failed was the null, and the reason is measurable

Mean z was **positive in every tier**: GOSPEL +0.276, TRAGEDY +0.665, MYTH +0.326. Positive means dissent sits *farther* from violence than chance. Every tier repelling is not a Girardian result, not an anti-Girardian result, and not plausible as a fact about narrative.

Round 7's null permuted **all** tokens, which destroys the arrangement of the **violence** hits as well as the dissent hits. So:

TIER	N	MEDIAN INDEX OF DISPERSION OF VIOLENCE-HIT GAPS
GOSPEL	4	**466.9**
TRAGEDY	23	**263.3**
MYTH	61	**350.9**

Uniform scatter is 1.00. **All 88 measurable texts are above it, median 323.** Violence vocabulary is concentrated in battle scenes and executions by a factor of a few hundred. A uniform-shuffle null spreads it evenly, shrinks the nearest-violence distance under the null, and pushes z positive regardless of where dissent sits.

> ***Fourth instance of the same null-design failure in this repository.***
>
> *| round | null | the mismatch | |---|---|---| | 3 | Null B | a **minimum** favours uniformly-spread sets over clumpy ones by construction | | 4 | D4 | nulled the **unigram half** when the whole effect lived in phrases | | 5 | phrase-aware arm | matched a phrase to its **first token's** frequency, inflating occurrences 316-fold | | 7 | Null A | **destroyed the violence clustering the statistic depends on** |*
>
> *Every one is a mismatch between what the null holds constant and what the statistic is sensitive to.* ***This is the most transferable thing in the repository and it is worth more than any of the p-values.***

Round 8: one line changed, and the power problem solved

Hold the violence positions exactly as they are. Permute only the dissent positions.

And the statistic was rebuilt to pool. For each dissent hit, *u* = the fraction of that text's token positions lying at least as close to a violence hit. **Under random placement* u *is exactly uniform on [0,1], for any text length and however clustered the violence is, so the clustering that broke round 7 is absorbed into the definition instead of being destroyed by the null. The tier statistic is the mean of* u* over every hit in the tier: **0.5 under the null, below 0.5 under Girard.**

Pooling gives **112 / 235 / 334 hits** instead of round 7's per-text median over five. Power solved rather than declared.

TIER	TEXTS	HITS	U	NULL MEAN	NULL SD	P
GOSPEL	4	112	**0.4171**	0.5018	0.0263	**0.0015**
TRAGEDY	23	235	0.4838	0.5025	0.0184	0.161
MYTH	50	334	0.4832	0.5018	0.0159	0.124

> **U1 passed. In the four Gospels, dissent vocabulary sits closer to violence vocabulary than chance places it, at p = 0.0015.** That is the first positive localization result in eight rounds, and it is the shape of the claim Girard actually makes rather than the frequency claim rounds 2 and 4 tested.

And then N3 fired

TIER	FULL LEXICON	CORE ONLY	CORE P
GOSPEL	**0.4171**	**0.5143**	0.602
TRAGEDY	0.4838	0.4838	0.196
MYTH	0.4832	0.5053	0.568

The Gospels go from lowest to highest. On the concepts that state Girard's actual claim — innocent, guiltless, no fault, protest, denounce, defend, plead — the Gospels show no co-location at all, and sit slightly *farther* from violence than chance.

> ***The entire U1 effect belongs to `unan`****: division, some said, others said, not all, spare him, withheld consent.*
>
> ***This is the fourth consecutive round to find the same split.*** *Round 4's S1, round 6's S1 on both translations, round 7's core sensitivity, and now round 8.* ***The `core` half has never carried anything, on any measure, in any round.***

U2 failed: tragedy 0.4838 and myth 0.4832 are the same number to three decimals, so there is no ordinal series, only a Gospel effect and a flat background. **U4 failed instructively**: the gap is **+0.0661**, which *clears* the 0.05 margin the pre-registration demanded, but tier-label permutation gives p = 0.123. **Magnitude without significance**, the mirror image of the restraint control's H3 and of interlock P5.

The Bacchae, a fourth time

ROUND	OPERATIONALIZATION	RESULT
2	whole-text rate, four buckets	rank **35 of 72**
3	`peak_min` against a word-shuffle null	**p = 0.66**
7	distance to nearest violence, uniform null	**p = 0.888**
8	co-location, violence held fixed	**p = 0.414**

Four pre-registered tests, four operationalizations, four failures.

And a falsifier I wrote correctly and coded wrongly

`score8.py` printed "core order matches full lexicon? yes" and did not flag N3. The check was `corder == u2`, comparing **two booleans rather than two orderings**. Both orderings failed their own strict test, both booleans were `False`, so the comparison returned `True` and printed reassurance.

> ***N3 fires because I read the table, not because the script caught it.*** *A falsifier stated correctly and coded incorrectly is worse than no falsifier, because it prints reassurance. It was caught only because the numbers are printed alongside. Recorded in the same class as round 3's Null B and round 5's badly written verdict rule.*

What survives, in one sentence

> **The Gospels put crowd-division language near violence language more than chance does. Myth and tragedy do not, and the concepts that state Girard's claim do not, in any tier.**

And the ceiling a pass does not lift: this measures co-location of two vocabularies. It does not show the dissenting words belong to a person separating from a crowd, that the violence is collective, or that the Gospel effect is Girard's mechanism rather than passion narratives putting speech and killing in the same chapters for ordinary narrative reasons.

No round 9 on this construct. `PREREGISTRATION-8.md` §8 committed that a third null would be the move the rules forbid.

1.21 The diagnostic: why `core` never carried anything, and what does

Not a round. Exploratory, post-hoc, computed on already-scored texts, and §4 below is why it cannot be claimed.

The standing puzzle

Four rounds, four different measures, the same result:

ROUND	MEASURE	WHAT `CORE` DID
4	whole-text rate, ASV	S1 took GOSPEL > MYTH from p = 0.0003 to **0.293**
6	whole-text rate, KJV	S1 gave **p = 0.254**
7	distance to nearest violence	every tier positive, **Gospels worst**
8	co-location, violence fixed	GOSPEL **0.5143**, highest of three tiers

"The effect lives in `unan`" was reported four times. It is true and it is not an explanation.

`core` is three different things, and they cancel

SUB-FAMILY	CONCEPTS	GOSPEL	TRAGEDY	MYTH	G/M
A. the accusation is EMPTY	no_fault, false_witness, done_nothing, without_cause, wrongly, unjust	**0.373**	0.162	0.101	**3.68×**
B. someone ADVOCATES	protest, object, denounce, defend, plead, speak_for	**0.012**	0.212	0.201	**0.06×**
C. the victim is INNOCENT	innocent, guiltless, not_guilty	0.060	0.165	0.106	0.57×
	all `core`, as scored in rounds 4-8	0.445	0.542	0.413	**1.08×**

> ***A runs 3.68× toward the Gospels. B runs 17× away. They cancel to 1.08×.***
>
> ***Six of sixteen `core` concepts fire zero times in the Gospels****: `not_guilty`, `slander`, `protest`, `denounce`, `defend`, `plead`, `speak_for`. And the two dominating the myth tier, `defend` at 0.109 and `plead` at 0.052, are exactly the forms that fired on hymnic epithets in `orphic_hymns_2` and produced round 4's false counterexamples.*
>
> *`core` has not been measuring nothing.* ***It has been measuring two opposite things at once and reporting their difference.***

And the split is not arbitrary. The file said so itself.

`dissent.py`'s header, written before any of this was scored:

> "So the construct is **not innocence**. It is a voice that **questions the verdict** and **breaks the unanimity**. Two different things, and the second is the one Girard says is absent from myth."

And Girard, quoted in the same header:

> "**not merely the innocence of the victims versus their guilt**, but the fact that, in mythology, no one ever questions this guilt."

> **The header states that innocence is not the differentia. The bucket then leads with `innocent`, `guiltless`, `not_guilty` and fills out with `defend`, `plead`, `protest`, `denounce`, `speak_for`.** Sub-families B and C together are ten of the sixteen concepts, and neither is what the sourcing note says the construct is. **The bucket contradicts its own docstring**, and that is checkable in the file rather than a matter of interpretation.

What every surviving result in this project has in common

WHAT SEPARATES THE GOSPELS	WHAT DOES NOT
the accusation is empty (no fault, false witness, done nothing)	the victim is innocent (innocent, guiltless)
the crowd is divided (division, some said, others said, not all)	someone advocates for him (defend, plead, protest)
the crowd is asked to release him (spare him, let him go)	

Both surviving columns are about the collective's verdict. Both failing columns are about the victim's character or his supporters. That is a sharper construct than the project has been testing, and it is closer to the source passage than the lexicon built from it.

Why it cannot be claimed, and what would test it

This split was found by looking at corpus3 **after eight rounds of scoring it**. Re-scoring corpus3 with `core` split three ways would be Correction 1's excluded Diodorus page, moved from the sample to the lexicon. It is the most tempting move available and it is forbidden.

And there is no held-out sample on disk. One text in `corpus/` or `corpus2/` but not in the scored tier set, and it is a second translation of a text already scored. Six apocrypha never scored on this lexicon, but they are Jesus-narratives and cannot be the myth tier of a tier contrast.

Round 9 needs new texts, with the fetch list frozen before anything is scored, the three-way split fixed by the sourcing note rather than by the numbers, and B's backwards run reported as part of the test.

The other structural idea, and why this corpus cannot carry it

Girard's claim is about **a voice**, and no instrument here has ever distinguished a character's speech from the narrator's. `orphic_hymns_2` firing on hymnic epithets is exactly that confusion.

SAMPLE	QUOTE MARKS	SPEECH VERBS	USABLE
`gospel_john`	**0**	454	speech verbs only
`euripides_bacchae`	94	17	quote marks
`homer_iliad_1`	70	42	quote marks
`apollodorus_2`	1	30	**neither**
`orphic_hymns_2`	**0**	**0**	**neither**

> *A speech-versus-narration instrument on this corpus would measure **typographic convention across four source sites and a dozen translators**, and the tier contrast would track the convention rather than the texts. That is the same mismatch that produced all four null-design failures. **Building it anyway would be the fifth.***

2.1 The claim under test

`conversion/PREREGISTRATION.md`, frozen `d671056`, 12 August 16:50.

Girard, *Deceit, Desire and the Novel*, p. 294. Novelistic conversion is a **complete inversion of the experiential register**, given as three pairs:

1. **deception becomes truth** — the subject can name the mediation
2. **agitation becomes repose** — the perpetual motion of the triangle stops
3. **hatred becomes love** — the rival can be looked at without measuring

Girard treats these three as **one event with three faces**, and uses the presence of that event to separate *novelistic* works from *romantic* ones.

That is two separable empirical claims:

- **(i)** the three components co-occur; they are a system, not a list
- **(ii)** they are concentrated in the canon Girard names

(i) is the primary test, because it is what makes "conversion" a natural kind rather than a word covering three unrelated properties.

2.2 The unit, and the rule attached to it

The ending window: the last 1,500 words of the novel, after Gutenberg boilerplate, afterwords and translator's notes are stripped.

Fixed in advance, not adjustable per text. The rule, quoted:

> If a text's conversion is famously outside its last 1,500 words, that text scores 0 on the affected bit and the mismatch is reported as a limitation of the window, **never fixed by widening it for that text. Widening the window is a new pre-registration with its own commit.**

2.3 The rubric: three bits, and the evidence rule

A bit scores 1 only if a quoted span from inside the window is recorded in the results file. No quote, no point. Same rule as the interlock instrument's FIT criterion: the text supplies the evidence, not the coder.

T, deception becomes truth. Scores 1 if the protagonist refers to their own former wanting as derived, borrowed, illusory, or belonging to another, or names a specific person or model as its source. Does **not** score for general regret, moral repentance, shame about conduct, or the narrator explaining the protagonist from outside in a third-person text.

R, agitation becomes repose. Scores 1 if **both**: (a) a quoted span shows the protagonist stopping, renouncing, releasing or coming to rest, **and** (b) no new object of desire is introduced after that span, inside the window. Death alone does not score R. Dying is termination, not repose.

> *Disclosure written into the rubric: clause (b) was suggested by the Plath case that motivated the study. It is retained because it is independently derivable from Girard's own definition of repose as the stopping of the triangle's motion, which requires that no new triangle start. Flagged so the effect stays visible.*

L, hatred becomes love. Scores 1 if a former rival or model is present or named and the protagonist's attitude shows no comparison or measurement. **Score NA if the novel has no rival or model figure at all. NA is not 0.**

2.4 The sample design

Set A, Girard's canon. Cervantes, Stendhal, Flaubert, Proust, Dostoevsky. Every novel by these five retrievable in English from Gutenberg. **No selection by me among them.**

Set B, control. Taken from Gutenberg's "Top 100 EBooks last 30 days", an externally generated ranking, applied mechanically in list order: single-author prose fiction of at least 40,000 words, nothing by a Set A author, no anthologies or plays or poetry or non-fiction.

Set C, the target. *The Bell Jar*. **Scored last, after Sets A and B are scored and committed.** It motivated the instrument, so it cannot calibrate it.

2.5 The mathematics

P2, the primary test. Let M be the n×3 matrix of bits. Define

$$E(M) = \#\{ \text{ rows at } (0,0,0) \text{ or } (1,1,1) \}$$

Null hypothesis: the three bits are independent with the observed marginals. Generate it by **permuting each column independently**, 10,000 times. This preserves each bit's marginal rate exactly and destroys any association between them.

$$p = (\#\{ \text{null} \geq \text{obs} \} + 1) / (NPERM + 1)$$

> **Why this is the right test for "one event with three faces."** If T, R and L are three faces of one thing, endings pile up at the extreme cells. If they are three unrelated properties, the cells fill at the independence rate.

P3. Permutation test on the group labels, difference of mean bit-sums, 10,000 shuffles.

Pairwise association. Phi coefficient for each pair, with a permutation null on the second variable, two-sided.

2.6 What ran, and what did not

`conversion/RESULTS.md`, `cf4bfa0`, 12 August 17:12.

> **The fetch was killed partway through, so Set B does not exist and P2 and P3 were never run.**

Not retrieved, therefore not scored: **all Dostoevsky, all Proust.** Their absence is a hole in Set A, not a result about them. Dostoevsky is where Girard's strongest conversion claims live.

All five scored windows are **tier V**: the last 1,500 words were cut from a text on disk and read. No score comes from memory.

#	TEXT	T	R	L	SUM
1	*Don Quixote* (Ormsby, PG 996)	1	1	1	**3**
2	*The Charterhouse of Parma* (PG 57638)	1	1	0	**2**
3	*The Red and the Black* (PG 44747)	0	1	0	**1**
4	*Madame Bovary* (PG 2413)	0	0	0	**0**
5	*The Bell Jar* (Set C, scored last)	0	0	0	**0**

2.7 The quoted spans

Quixote T. *"I was mad, now I am in my senses; I was Don Quixote of La Mancha, I am now, as I said, Alonso Quixano the Good."* And the mediation named as transmitted: *"Forgive me, my friend, that I led thee to seem as mad as myself, making thee fall into the same error I myself fell into."*

Quixote R. Sancho offers a **new** project, the shepherd's life, and Quixote refuses it: *"in last year's nests there are no birds this year."* Clause (b) satisfied in the text's own words.

Quixote L. The rival is Avellaneda, the false author: *"they beg of him on my behalf as earnestly as they can to forgive me."*

Charterhouse T. *"ever since the moment when I had the happiness of being locked up by Barbone, ambition has been to me an empty word."*

Charterhouse L = 0. Fabrizio has his rival, the Marchese Crescenzi, kidnapped **inside the window**. Rivalry is operative, not dissolved.

Red and Black T = 0. Julien names a mediator, but somebody else's: *"This idea of going to Paris is not your own. Tell me the name of the intriguing woman who suggested it to you."* The rubric requires the protagonist to name the source of **their own** wanting. His does not appear in the window; the prison conversion sits just outside it.

Red and Black L = 0. The measuring is still running: *"the Maslons, the Valenods, and the thousand other people who are worth more than they."*

Bovary, all three = 0. The protagonist is not in her own ending window. Charles refuses to name any mechanism (*"It is the fault of fatality!"*, which Flaubert marks as *"a fine phrase, the only one he ever made"*), and metaphysical desire is still arriving rather than departing: *"He would have liked to have been this man."* The last sentence is a triangle paying out: *"He has just received the cross of the Legion of Honour."*

2.8 P1 and P4 both failed

P1. Predicted T = 1 in ≥70% of Set A and sum ≥2 in ≥70%. Observed **2 of 4 (50%)** on both. F4 also fired (Set A under 6 texts), so this is descriptive and carries no p-value. But the direction is not marginal. **Girard's canon does not pass Girard's own definition when the definition is applied to a fixed window.**

P4. Predicted, in advance and in conversation before the rubric existed: *The Bell Jar* scores T = 1, R = 0, L = 0. Observed **(0, 0, 0)**.

R and L came out as predicted, with the quotes the prediction was built on. R = 0 because the renunciation is there (*"'Never,' I said, and hung up with a resolute click. I was perfectly free."*) but clause (b) breaks: *"my red wool suit flamboyant as my plans"*, the unanswered *"I wonder who you'll marry now"*, and the last line is an entrance.

T failed for the same reason Red and Black failed. Esther's naming of the mechanism is real and quotable, *"for the price of their care and influence, have me resemble them"*, but it is in **chapter XVIII**, roughly **nine thousand words before the end**. The window cannot see it.

Not fixed. Widening the window to reach it would be tuning the instrument to make a text pass.

One operationalization disclosed during scoring. The L rubric as written scores 1 when the attitude "shows no comparison or measurement", which lets silence score. Applied literally, Plath's funeral scene would score L = 1 on an absence. The stricter reading was applied throughout: **L requires positive regard toward the rival, not merely the absence of measuring.** Quixote scores under it, Plath does not. Recorded because it changed a score, and it was applied to every text, not to the one that needed it.

2.9 What survived, and neither part was predicted

5.1 The canon spreads across the whole range: 3, 2, 1, 0. If conversion were one event with three faces, Girard's own novels should pile up near 3. Across the four retrieved they occupy **four different scores.** Even granting that the window mislocates two of them, *Don Quixote* at 3 and *Madame Bovary* at 0 are both correctly scored and both are core cases. P2 was designed to test exactly this and never ran, but the qualitative pattern is already pointing **away from the bundle.**

5.2 Bovary scores 0 because the conversion in Flaubert is not in Emma.

Girard classes *Madame Bovary* as **novelistic**, the highest category he has, and its ending scores the same as a novel he would call romantic. The reason is visible in the text: **the lucidity in Flaubert belongs to the narration, not to any character.** Nobody in that book ever sees the triangle. The author does, continuously, and hands it to the reader over the characters' heads.

> So the three pairs describe **what happens to a character**, while *novelistic* names **where the author is standing**. The instrument, applied honestly, measures the wrong object, and **it took a failing score on Flaubert to show it.**

2.10 Set D: your own novella, scored after the results were committed

`930d8e3`, 12 August 17:25.

TEXT	T	R	L	SUM
The Crowning (19,146 words)	1	0	NA	1 of 2 applicable

T = 1, and it lands on the last page: *"It was warm because I did."* Supported by *"I carried it down seven floors, held it toward a gate that did not want it, put it back on my own head, and called the act restraint."*

R = 0, and the results file calls it the least defensible score in the run. The cessation is explicit (*"I stopped walking", "I do not close the distance."*) and then the final sentence is *"I am going to open my mouth now."* Clause (b) breaks: a want opens after the stop. The contrary argument, that asking plainly in a way that can be refused is the dissolution of the triangle rather than a new one, **is real and is not answered.** Recorded as contested.

L = NA, the first NA in the corpus. No rival figure appears in the window at all. Structural consequence: the novella runs mediation without rivalry.

Cross-text observation. *The Crowning* and *The Bell Jar* both score R = 0 for the same reason. Both end **one beat before an ask, on forward motion.** They differ on T placement: last page versus nine thousand words upstream.

3.1 Three questions, kept strictly apart

LAYER	QUESTION	WHERE
1a	Do the NT texts interlock the way independent reports do, or collide the way legends do?	`RESULTS.md`
1b	When someone alleges a contradiction, is it one, and what does the rescue cost?	`ntcount/`, `rescue/`
2	Is the historical case for the resurrection **honestly built**?	`resurrection/`

Status on every headline in every file: strong structured pilot, not a confirmatory result. One coder, not blinded, who recognises every text on sight. No amount of process moves it past that.

Boundary stated once. Nothing here reaches whether the resurrection happened. Nothing here argues for God. A clean result is compatible with an independent record of a shared legend; only dating constrains that, and dating is not in these files.

3.2 Layer 1a: the interlock run

The claim, and the blinding tooling that almost nobody knows is there

`interlock/PREREGISTRATION-2.md`, written after two pilots. The claim: the canonical gospels **interlock** (one account's dangling detail is incidentally completed by another) more than the legendary apocrypha, and contradict and free-invent less.

This run has real blinding machinery, and it is stronger than I have been giving it credit for. `build_run.py` emits each pericope as a packet of witnesses under **opaque codes** W1, W2, W3 in a fixed shuffled order, with the source identities sealed in `run/key.json`. Gaps are logged and committed **per witness, before the key is opened**. Confirmed on disk: `key.json` maps crucifixion W1 to gospel-of-peter, W2 to mark, W3 to luke, W4 to matthew.

> **And §7 states the limit in the same breath, which is why the file is trustworthy.** "An LLM coder recognises the canonical text by content, so it cannot be truly blinded to canon-vs-apocrypha the way a naive reader could, and one coder yields no inter-coder agreement statistic. So the subjectivity of the classes goes **unquantified**."

The two scans, and the fact that one of them barely worked

Two independent passes per pericope. The **gap scan** logs every instance of a closed typology before any other account is opened: **G1** unintroduced referent, **G2** unexplained motive, **G3** abrupt statement, **G4** missing place or time, **G5** uncaused emotion, **G6** raised-then-dropped detail. Then the **flat-collision scan** lists every concrete detail asserted by two or more accounts, because Pilot 1 found the gap scan alone under-counts collisions about fivefold.

Recomputed from `run/coding_sheet.csv`, 39 items:

GAP TYPE	ITEMS TAGGED
G2 unexplained motive	3
G6 raised-then-dropped	2
G1 unintroduced referent	1
G3 abrupt statement	1
untagged (came from the flat scan)	**32**

> ***The elaborate six-type gap typology produced seven items out of 39.*** *Four fifths of the run came from the flat-collision scan the pilots added as a patch. The typology was the theoretically motivated half and it did almost no work. This is not recorded anywhere in the repository.*

Class and pair distribution, recomputed:

CLASS	N	PAIR	N
AGREEMENT	11	IND (independent)	30
COLLISION-SOFT	10	DEP (Matt/Luke on Mark)	9
NULL-SILENCE	9		
FIT	**5**	pericopes	5
COLLISION-HARD	4		

> ***The headline statistic rests on a numerator of five.*** *The interlock ratio is FIT over FIT plus collisions, and there are five FITs in the whole run.*

39 coded items, 5 pericopes. **A difference is a FIT only if a text supplies the reconciling fact, not the coder.**

Aggregate over independent-pair items: canon 5 fits / 1 hard / 4 soft, ratio **0.50**. Apocrypha 0 fits / 3 hard / 1 soft, ratio **0.00**. Embellishment items: canon 1 in 4,555 words = 0.22 per 1000; apocrypha 8 in 3,152 words = 2.54 per 1000. About **11 times**.

	PREDICTION	RESULT
P1	canon ratio beats apocrypha	holds, both matched comparisons
P2	apocrypha hard-collide more	holds, 0.95 vs 0.22 per 1000
P3	apocrypha embellish ≥3× more	holds, ~11×
P4	advantage does not depend on Synoptic copying	holds, all five canon fits are independent-pair
P5	canon ratio ≥ 0.60	**FAILS at 0.50**

The part of this the file tells you not to trust, and it is its own

> **The fit-ratio half (P1, P5) is the least trustworthy thing here, and I produced it.** The fits I logged for the canon are the *famous* undesigned coincidences I already knew to look for: Bethsaida/Philip, Passover/green grass, John's "we". **That is textbook selection.** A blind coder scanning every gap without knowing the apologetic literature would very likely log fewer, and the canon ratio would fall. The numerator is exactly where a knowing coder's thumb rests, and there was only my thumb.

The objective half (P2, P3) is trustworthy because it does not depend on charity or on knowing the hypothesis: a schoolmaster dies in one infancy gospel and merely blushes in the other, a walking cross speaks in the Gospel of Peter and nowhere else. Those are facts on the page.

3.3 Layer 1b: the contradiction count

`ntcount/PREREGISTRATION-NT.md`, frozen `83f5f55`, 24 July 22:53.

Why every existing count is worthless, on two grounds

> **No definition.** None of them states what makes something a contradiction. A count nobody can reproduce cannot be checked, and a claim that cannot be wrong is not a measurement.
>
> **No denominator.** 439 out of *what*? A corpus with four accounts of one life has vastly more contradiction opportunities than a corpus with one. **Counting raw hits punishes multiple attestation.** That is backwards, and correcting it is the whole point of this study.

The definition, frozen

> Two texts contradict when they assert P and not-P about the **same subject, in the same respect, at the same time**, such that no reading licensed on the channels in `RUBRIC.md` makes both true.

The three qualifiers are Aristotelian, which matters, because the obvious objection is "you defined contradiction so that nothing counts." They are what the law of non-contradiction has always required. Dropping them is what lets a list reach 439.

Note what the definition does **not** say. It does not say "no reading I can imagine makes both true." The licensing requirement binds both ways.

The nine gates

Applied first, to every docket item. Rejections record **which** gate failed, because the histogram of failure reasons is a primary result.

CODE	NAME	TEST
A0	not an assertion	one side is a question, command, wish, curse, hypothetical, or the speech of a character the text does not endorse
A1	not the same subject	different events, persons, places or occasions
A2	not the same respect	different sense of a shared word; rhetorical quantifier read as strict; figurative against literal; legal against biological; role against person
A3	not the same time	different period or dispensation; before against after
A4	omission only	silence is not denial. Naming one of two is not asserting only one
A5	tension, not negation	two emphases in tension without P and not-P
A6	text against doctrine	the clash is with a later dogma, not between two texts
A7	designed juxtaposition	an intentional literary pairing (Prov 26:4-5)
A8	reference failure	the cited verses do not say what the entry claims
ADM	admitted	survives all gates

The first failing gate is recorded, in fixed order. Many items fail several. If the coder may choose among applicable gates, the histogram becomes a record of preference rather than of the texts.

The gates with real items attached, so you can judge them yourself

Abstract gates are unfalsifiable. Here is what each actually did.

A0, not an assertion (4). All four are worth knowing because they are the cleanest wins on the whole docket.

ITEM	THE ALLEGATION	WHY IT IS NOT A CLAIM
C0104	Did Jesus ask God to save him from crucifixion?	John 12:27 is a rhetorical question **Jesus immediately rejects**: "And what shall I say? Father, save me from this hour? No, it was for this that I came"
C0209	Did Herod think Jesus was John the Baptist?	Luke 9:9 is perplexity: "John I beheaded, but who is this?"
C0406	Was Jesus the son of David?	Matt 22:45 is a riddle Jesus poses
C0454	Is God's will always done in heaven?	Matt 6:10 is a **petition**, "your will be done", read as an assertion that it always is

A1, not the same subject (7), and the abuse risk stated in the file itself.

> "**The abuse risk is real**: a harmoniser can manufacture this gate by multiplying events ('there were two temple cleansings'). So it fires only where the texts themselves mark different occasions, never where I would need them to. Where the texts do not mark it, the item was admitted instead: see C0421, the temple cleansing, **admitted at cost 2**."

That last clause is the check working. The single most famous multiply-the-events move available was refused, and the item was let through.

A2, not the same respect (62). The gate F4 fired on. A representative sample, chosen to show both the strong and the weak end:

ITEM	THE SENSE DISTINCTION CLAIMED	HOW STRONG IT LOOKS
C0051	Did Jesus baptize anyone?	**Strong.** The author corrects his own phrasing in the same book: John 4:2, "although in fact it was not Jesus who baptized, but his disciples"
C0075	Should we bear each other's burdens?	**Strong.** Different Greek words three verses apart: *barē* (v2, crushing burdens borne together) against *phortion* (v5, one's own proper load)
C0250	Was Joseph the father of Jesus?	**Strong.** Luke 3:23 marks the distinction itself with *hōs enomizeto*, "as was supposed"
C0318	Are those who believe Jesus is the Christ of God?	**Strong.** James 2:19 draws the distinction the entry ignores: "even the demons believe, and shudder"
C0433	Was it ok to touch the risen Jesus?	**Strong.** *mē mou haptou* is a present imperative, "stop clinging", and he invites touch nine verses later at 20:27
C0099	What colour was Jesus' robe?	**Weaker.** *porphyra* covers a range of red-purple and *kokkinē* a crimson, "a faded military cloak falls in the overlap"
C0140	Does Luke contain everything Jesus did?	**Weaker.** "All that Jesus began to do and teach" called a rhetorical quantifier
C0353	Ransom for many or for all?	**Weaker.** "Many" as the Semitic inclusive idiom from Isa 53
C0517	How long was Jesus in the tomb?	**Weaker.** Inclusive reckoning, where a part-day counts as a day

> *Look at that spread and you can see F4's problem without any statistics. The top five cite a lexeme or a verse in the same passage. The bottom four cite a convention.* ***The gate is one gate and it is doing two different jobs****, and nothing in the coding separates them. That is a defect the repository does not name, and it is the concrete form of what F4 detected as a number.*

A4, omission only (12), and this is the gate that quietly does the most work.

> "Silence is not denial. Naming one figure at the tomb does not assert that there was exactly one; to convert an omission into a contradiction you need a text that says *only*, and none does. **This single gate dissolves the whole one-angel-against-two class that appears on every popular list, and it requires no reconstruction at all, which is why it is stronger than the harmonisers' answer to the same items.**"

Items it takes: the donkey at the entry, the blind men at Jericho, who buried Jesus, Annas before Caiaphas, the centurion's words, the demons' request, the kiss of Judas, the last words from the cross, the possessed men, the rich young man's commandments, how many women came, how many beatitudes.

A8, reference failure (5), fully checkable by anyone with a Bible. The compilers cited Luke 8:5 and 8:7 for a claim about prayer; those verses are the parable of the sower. They set 1 Cor 9:1 against Acts 9:8 to claim Paul did not see Jesus, when Acts 9:8 says he saw nothing **after opening his eyes, having been blinded**. Matt 3:7 says the Pharisees came to the baptism and were rebuked; it never says they were baptised.

A7, designed juxtaposition, fired zero times, and `VERDICT.md` records why that is worth printing: "a gate that never fires is evidence it was not used as a free pass."

The cost scale, inherited unchanged

COST	NAME	MEANING
0	dissolved	the plain text shows no contradiction at all
1	licensed	a rescue on some channel is directly cited and uncontested
2	contested	a real rescue exists but rests on a debated reading
3	ad hoc	reconciliation needs an event no text states, or a text emendation
4	hard	no non-ad-hoc rescue on any channel. A genuine contradiction

The five rescue channels, each requiring a citation: GREEK (must cite the lexeme), ACCOUNT (must cite the verse), HISTORY (must cite the fact), HISTORIAN (must cite the source), GENRE (must name the convention, not merely assert it).

The Ehrman guard, and its symmetry is the whole point:

> A rescue rates cost 1 only if it cites a lexeme, a verse, an attested custom, a named source, or a named convention. **It stops the apologist's invented harmony and it stops me dissolving an item because it feels obviously fine. Both are the same move**: supplying from outside the text what the text does not supply.

Consequence accepted: some real reconciliations get priced at 2 or 3 because nobody wrote them down. That is the correct cost of a rule that binds both ways.

Deduplication and granularity, frozen before scoring

> Granularity is a **free variable in every published count**, and it is worth more than any argument. Resurrection morning can be one contradiction or eight depending on whether you split time, participants, angels, posture and sequence. **A compiler can multiply or divide a total by five without touching a word of the text, and no published list declares its rule.**

Rule fixed: **one item equals one pair of incompatible propositions.** Merges and splits logged individually, before Tier A, before any cost is assigned.

3.4 The funnel

Docket is hostile-sourced and complete: BibViz aggregation of the Skeptic's Annotated Bible (474), Infidels.org (56) and EvilBible.com (141). Raw total 671.

STAGE	COUNT
raw entries claimed, NT-internal	210
after exact-match deduplication	198
after granularity merge	**153**
admitted as genuine contradictions	51 (33.3%)
cost ≥3	**2** (3.9% of admitted)
cost 4	**0**
cost 4 on the core	**0**

The entire hard residue, both items:

- **C0264, cost 3** — how Judas died. The hanged-then-fell reconstruction needs a rope or branch failure and an interval **no text states**.
- **C0120, cost 3** — to whom Peter denied Jesus. Mark's servant girl against Luke's "a man" at the second denial. Harmonising requires positing more challengers than any text states.

Both peripheral. Neither touches crucified-under-Pilate, buried, or tomb-found-empty.

The 51 accepted items, by cost

COST	NAME	N
1	Licensed	**11**
2	Contested	**38**
3	Ad hoc	**2**
4	Hard	**0**

> ***Read that distribution before reading anything else about this study.*** *Three quarters of what survived the gates landed at* ***contested****, which the rubric defines as "banked for neither side." The study's real output is not "the contradictions dissolve." It is "****most of them are live difficulties with real but disputed answers****", which is a much duller and much more defensible sentence.*

The eleven at cost 1 are the ones with a citable licence, and they are the ones worth being able to say out loud:

ITEM	THE LICENCE, CITED
C0165 potter's field	Acts uses *ektēsato* (*ktaomai*, "acquire"), the verb Luke uses of Paul acquiring citizenship at Acts 22:28, **not** *agorazō*, "buy"
C0041 ascension timing	The same author supplies the interval. Acts 1:3 states forty days. An author who writes forty days in volume two is not asserting same-day in volume one
C0095 answer to the high priest	*su eipas* is an affirmative Semitic idiom, used the same way to Judas at Matt 26:25
C0109 who carried the cross	The condemned carried the *patibulum* out, and Mark 15:21 has Simon pressed in "on his way in from the country". Roman *angareia*, the same verb as Matt 5:41
C0117 arrival at the tomb	One journey, not rival claims. Matt 28:1 "toward the dawn" and Luke 24:1 "early dawn" supply the gradient between John's dark and Mark's sunrise
C0169 first appearance	Mark 16:9 alone asserts "first" (*prōton*) and **sits in the longer ending, absent from Sinaiticus and Vaticanus**
C0392 the sign over the cross	The four differ only by inclusion, no word contradicts another, and John 19:20 states the inscription stood in three languages
C0065 the best seats	Matthew itself replies to the sons in the plural at 20:22, so both mother and sons are involved in Matthew's own telling
C0089 the centurion	Agency convention: an act done through delegates is attributed to the principal, and Luke supplies the delegation
C0143 did the eleven believe	Luke records ongoing disbelief in the same scene: 24:41, "while they still disbelieved for joy"
C0343 after the baptism	John never narrates the baptism, only refers back with the perfect *tetheamai*, "I have seen"

And one item where the study corrected itself against its own advantage. C0035, Galilee against Jerusalem, carried a `core` flag from the earlier rescue ledger. `VERDICT.md` removes it:

> "**CORE FLAG REMOVED**: the frozen definition in `PREREGISTRATION-NT` §3 is crucified/buried/tomb empty, and appearances are not in it. The rescue ledger used a wider notion of core; **that discrepancy is reported, not silently reconciled.**"

Removing the flag makes the headline claim ("nothing hard touches the core") easier to satisfy, so this is a change in the study's favour. It is recorded as a discrepancy rather than absorbed, which is the correct handling and worth noting because the direction is the tempting one.

The items that drift, named by the file

`VERDICT.md` flags where a cost-2 is holding on by its fingernails. C0423, on whether the women told the disciples, rests on reading Mark 16:8's "said nothing to anyone" as silence in flight rather than permanent silence, and the file says plainly that "no text supplies the 'on the way' qualifier, **so it drifts toward 3 if Mark's absolute phrasing is pressed.**"

That is a study marking its own weakest joint in public.

Judas was moved from cost 4 to cost 3 on 24 July. Cost 4 means *no* rescue exists; an ad hoc one does. The correction went **against** the first instinct and **against** the rhetorically stronger position, since a lone hard contradiction is a cleaner story than a cost 3. The scale outranks the instinct and outranks the story.

3.5 F4 fired, and the only valid form of "what survives"

> **F4 FIRED.** A2 plus A5 rejections are **67.6%** of all rejections against a pre-declared ceiling of 40%. Per §9 that means the dissolution is resting on my own sense-distinctions rather than on the texts, and Number A must be reported as unreliable.

The results file's own first sentence:

> So the honest first sentence of this study is not a number about the New Testament. It is that I dissolved 62 entries by saying "not the same respect" and 7 more by saying "tension, not negation," and I pre-committed to treating that volume as disqualifying. **It is disqualifying. The 66.7% rejection rate is not a finding.** I set that threshold precisely because A2 is the gate where a motivated coder's thumb would rest. It was.

CODE	REASON	N	% OF ENTRIES
A2	not the same respect	62	**40.5%**
A4	omission only	12	7.8%
A1	not the same subject	7	4.6%
A5	tension, not negation	7	4.6%
A8	reference failure	5	3.3%
A0	not an assertion	4	2.6%
A3	not the same time	4	2.6%
A6	text against doctrine	1	0.7%
A7	designed juxtaposition	0	0.0%
	rejected	**102**	**66.7%**
ADM	**admitted**	**51**	**33.3%**

The structural argument, and why it is not an escape

F4 says the rejections are untrustworthy. It does not automatically contaminate every output, **and the reason is directional**:

> An A2 or A5 error puts an item back in play as a *contested* clash, at cost 1 or 2. **It cannot manufacture a cost 3 or 4**, because a cost 3 requires that no channel supplies a rescue, and an item rejected for having a sense-distinction available is **by construction** an item where a channel does.

Run as a worst case, with every A2 and A5 call assumed wrong:

	BEFORE	AFTER
admitted	51	**120** (78.4% of entries)
cost ≥3	2	**2** (1.7% of admitted)
cost 4	0	**0**
F2 (cost ≥3 over 10%)	not fired	still not fired

Untrustworthy: the 33.3% admitted rate, the 66.7% rejection rate, and any claim that the NT "mostly dissolves." Treat 33.3% as a floor and 78.4% as a ceiling.

Robust: the count of hard contradictions is 2 and the count at cost 4 is 0, across the entire hostile docket.

> **The discipline this encodes:** report the fired falsifier at the top, then show what survives and **why structurally**, never merely assert survival.

3.6 What is objective regardless of anyone's judgement

These do not run through Tier A at all and anyone can check them against `docket.csv`, `dedupe_log.md` and `granularity_log.md`.

- **57 of 210 raw entries (27%) are restatements** of an entry already on the list. The cock crow appears four times. Matt 5:16 against Matt 6:1 appears three times. The potter's field appears three times. **That is counting, not adjudication.**
- **Five entries misdescribe their own citations.** The clearest cites Luke 8:5 and 8:7 for an importunity-in-prayer claim; those verses are the parable of the sower and say nothing about prayer. Another sets 1 Cor 9:1 against Acts 9:8 to claim Paul did not see Jesus, when Acts 9:8 says he saw nothing *after opening his eyes, having been blinded*.
- **Four entries put an interrogative on one side.** Jesus asking "shall I say, Father save me from this hour?" and then rejecting it. Herod asking "who is this?". A question is not a claim.

3.7 The Bayes lever, and its stated boundary

`rescue/bayes.py`. Two hypotheses, stated precisely because they are the whole meaning:

- **H_multi** — the material derives from MULTIPLE uncoordinated sources reporting the same underlying events.
- **H_single** — the material derives from ONE fabricated source, later copied.

For each observation X:

$$B(X) = P(X \mid H\_multi) \,/\, P(X \mid H\_single)$$

$$B\_total = \exp(\ \alpha \cdot \Sigma\_i \log B\_i\ ),\ \alpha \in \{\ 1.0, 0.5, 0.33\ \}$$

The **α exponent is a correlation discount** for the fact that the items are not independent of each other. Each cell in the frozen table is given as (low, default, high) so the run reports robustness rather than a single fragile number. Result at default and $\alpha = 0.5$: **B ≈ 72.7**.

Posterior from prior odds:

$$odds\_post = odds\_prior \times B\ ;\ P = odds\_post \,/\, (1 + odds\_post)$$

The boundary, printed on every run:

> This decides independence versus coordinated fabrication. **It does NOT reach truth.** Independent witnesses to a shared *legend* also satisfy H_multi; only DATING constrains that, and dating is not in this file. So a large B kills "some men made it up together" and nothing more.

The cost-3 amendment, and why it is Bayes-neutral by construction. The frozen table had no cost-3 case, because nothing had ever scored 3 until Judas moved. The value was written down **before running it**:

> peripheral cost 3 gets identical treatment to peripheral cost 4, because the reasoning that earns the factor, "a fabricator would have smoothed this", does not care whether an ad hoc rescue happens to exist, so the factor must not change either.

Total held at 72.7. **That is what a non-tuned amendment looks like.**

3.8 Layer 2: the rule that inverts

`resurrection/PREREGISTRATION-R.md`, frozen `0fda18f`, 25 July 00:56.

The ceiling, fixed before any scoring

> Once the facts are fixed, the step from "the tomb was empty and the disciples had experiences they took to be encounters" to "God raised him" is dominated by the prior on whether such an event can occur. **No amount of source checking touches a prior.** Hume's objection lives there and this instrument cannot reach it in either direction.

Best possible outcome: **"the historical case is honestly constructed."** Not "the resurrection happened."

The inversion, and it is the only protection Layer 2 has

In Layer 1 the hostile compilers picked the docket. Here the temptation runs the other way.

> `isjesusalive.com` is the **object under test** and never supplies the docket. The docket of objections comes from **critics**: Ehrman, Lüdemann, Goulder, Carrier, and above all **Dale Allison**.

Why Allison is the key auditor. He **affirms the minimal facts while rejecting the methodology.** That splits the two questions Layer 2 must keep apart: *are the facts right* and *is the argument valid*. A source friendly to the conclusion and hostile to the argument is the ideal auditor of an argument's structure. **Both halves always get reported.**

> Layer 1 had structural protection. Layer 2 has this rule and nothing else, and the motivation risk is the highest in the project: a believer auditing a case he wants to be true, on the question that matters most to him, with a single non-blind coder. **If relaxed: Layer 2 is void. Not weakened. Void.**

G1 fired, at two levels

Level one, the site. `isjesusalive.com/empty-tomb/` asserts "Most scholars, not just Christian scholars, accept that the tomb was found empty." **It cites no published source for this at all.** The only modern scholar named on the page is William Lane Craig, quoted on a different point.

Level two, the underlying survey. The claim traces to Gary Habermas's bibliography of resurrection scholarship, source of the "roughly 75 percent" figure. **Habermas stated in 2012 that "most of this material is unpublished."** That is his own description of his own data.

Verdict: the empty-tomb consensus premise scores cost 4. Not because the empty tomb is false, and not because the majority claim is false. **Because the cited support cannot be examined.**

> **This is the same failure as our own run-1 fit ratio**, where the numerator was coder-selected and we discounted it ourselves. It gets the same discount here. **Consistency requires that, and the study would be worthless if it applied the rule only to the other side.**

The independent methodological problem, which stands without trusting Carrier: the figure is not 75% *of scholars*, it is 75% of **writers who published articles arguing specifically for or against the empty tomb**. People who publish on a contested question are self-selected. The inference is invalid regardless of who raises it. (Carrier's separate claim that Habermas refused data requests is tagged **tier K**, a lead and not a finding.)

Enemy attestation is not what it is called: cost 3

The page's strongest-sounding argument cites Matthew 28:12-13, then Justin Martyr and Tertullian. Three problems in order of severity:

1. **Matthew reporting his opponents' explanation is not his opponents attesting.** It is one author's account of what his adversaries said. This is the reverse of the criterion's normal use.
2. **The passage is single-source**, and belongs to the guards-at-the-tomb material critics standardly identify as apologetic development. Building the strongest leg on the pericope most suspected of being defensive is poor ground.
3. **Justin and Tertullian may be downstream of Matthew.** Both are a century or more later and both had Matthew.

G4 signalled, not fired. The enemy-attestation argument leans on a specific single-source pericope, which is a detail-level dependence. If round 2 finds the *surviving* arguments also require detail-level accuracy, G4 fires and the project's two-layer framing was wrong.

What this does not show

- **It does not show the tomb was not empty.** Allison, who rejects the method, nonetheless makes a relatively strong case for the historicity of the empty tomb on other grounds. The premise being unauditable is a fact about the argument, not about the tomb.
- **It does not show most scholars reject the empty tomb.** That would need the same published data nobody has. The honest position is that **the state of scholarly opinion on this question is not currently establishable from published sources**, which is uncomfortable for both sides.

Practical conclusion recorded in the file: if you argue the empty tomb in public, **drop the consensus percentage and drop the enemy-attestation move.** Argue the women and the Jerusalem proclamation, and cite Allison.

3.9 The patristic stratum: the Hard Twelve, graded

Fathers enter only as **named sources**, never as a scoring round. The Ehrman guard stays symmetric, so age is not licence. **Zero cost changes** across the whole survey, which is the expected and safe result.

How "hardest" was chosen, and when

> "**The ranking below was written and saved to this file before any source was fetched for this round.** That ordering is the only thing stopping the list from being quietly reshaped by which fathers turned out to have good answers."

Seven mechanical criteria, read off `scores.tsv` rather than off a sense of which items feel bad:

#	CRITERION	WHY IT MARKS A WEAK POINT
D1	cost 3	top of the observed scale; the rescue invents an event
D2	verified-dead rescue	a tier V row where the defence was checked and failed is worse than a contested one, **because the contest is over**
D3	no defence offered	an item a developed harmony skips is usually the one without an answer
D4	clustering	several accepted items on one pericope compound; a reader does not meet them singly
D5	proximity to the resurrection claim	nothing is core-flagged, but some items sit next to the core
D6	the rescue concedes something expensive	a defence that works by admitting the text was altered has a price defenders rarely state
D7	not a harmony problem at all	an item falsified by events rather than reconciled by exegesis is a different and worse kind of weak

The scorecard

#	ITEM	FATHER	GRADE	COST MOVED
1	Judas's death	Papias / Augustine	**ABSENT + AGAINST**	no
2	Peter's accusers	Augustine	PARTIAL, bundled	no
3	Galilee	Augustine	**GOOD**, bundled	no
4	empty tomb	Augustine / Origen	**BAD** / PARTIAL	no
4	empty tomb	Eusebius	**BAD** solution, **GOOD** method	no
5	staff on the mission	Augustine	**BAD**	no
6	Paul's companions' hearing	Chrysostom	**BAD**, verified dead	no
7	the hour of crucifixion	Augustine	**BAD**	no
8	the genealogies	Africanus / Origen	PARTIAL / **AGAINST**	no
9	temple cleansing	Origen	**AGAINST**	no
10	rendered speech	Augustine	**GOOD**, honest about the price	no
11	James against Paul	Eusebius	not a rescue: **canonical doubt**	no
12	"this generation"	Augustine	**BAD**, evasive, tier K	no

Three GOOD, three PARTIAL, six BAD, three AGAINST, one ABSENT.

> **The headline: the fathers cover MORE of the hard items than the modern harmonies do, and their answers are worse.** They walk straight into items a modern defender skips, and then rescue them with moves that fail the guard.
>
> **The distribution is the point.** A defender who says "the church has always had answers to these" is right that answers exist and wrong about what they are worth. A critic who says "they never noticed" is simply wrong. **Both of those moves die on this table**, which is why the table is the deliverable and not the individual verdicts.

The worst finding is a silence

Augustine's *De consensu evangelistarum* is four books long and exists precisely to answer people alleging the evangelists disagree. Book 3 chapter 7 works the Judas pericope: the thirty pieces of silver, Judas's repentance, the priests refusing to bank the money, the potter's field, and the prophecy Matthew attributes to Jeremiah that is actually in Zechariah.

He never mentions Acts 1:18. Checked directly. No falling headlong, no bursting, no attempt to square the hanging with the bowels.

And the one father who **does** describe the death gives a **third** account. Papias, via Apollinaris:

> "Judas walked about in this world a sad example of impiety; for his body having swollen to such an extent that he could not pass where a chariot could pass easily, he was crushed by the chariot, so that his bowels gushed out."

No rope. No tree. No fall. A swelling and a chariot.

> The modern rescue asks you to believe there was one event, hanging followed by a fall, told in two partial ways. **The earliest patristic witness to that event knows a version with neither the hanging nor the fall. That is what a multiform tradition looks like, not what one event reported twice looks like.**

Papias cannot move the score, because the docket is NT-internal by pre-registration and a patristic text cannot supply the P or the not-P. Recorded for what it shows about the shape of the tradition.

The best patristic answer, and the bundle attached to it

Augustine, *De consensu* 3.25, on Galilee against Jerusalem:

> "the words are simply these: Behold, He goes before you into Galilee; there shall you see Him; and **there is no statement of the precise time** at which that meeting was to take place."

That is a dissolution, not a harmonisation, licensed by the text's own silence, and it gives an opponent nothing to grip. Made in AD 400.

Then he bundles it with two figurative readings of "Galilee", transmigration and revelation, neither licensed and neither needed, plus an ordered list of ten appearances which is a reconstruction attackable at every joint. **Grade: GOOD, with a detachable bad half.**

Eusebius: the worst solution and the best method, in one text

To Marinus 4, read from the Pearse edition directly, tier V. He states the empty-tomb problem himself, and his statement is a better enumeration of the clash than the compilers' is. His solution:

> "there are four evangelists, and also a corresponding number of sightings **There are four occasions, and four seen**"

The Greek is explicit that the count is **derived rather than observed**: *tessarōn ontōn tōn euangelistōn, isarithmoi toutois*, "there being four evangelists, the sightings are **equal in number to them**."

> ***Grade: BAD, the most extreme multiplication in the file.*** *Lindsell's six denials are modest beside it. The number of events is being read off the number of sources,* ***which guarantees in advance that no two sources can ever conflict.***

And the corpus contradicts itself on order, flagged by the edition's own editors: *To Marinus* 4.1 opens "The incident in Matthew comes first"; Nicetas-Marinus 5, answering the same question, opens "I take it that the narrative in John comes before that in Matthew."

Then, in the same fragment, he does the one genuinely excellent thing in the whole patristic survey. He states what would have made the accusation stick:

> "it really would have been justifiable to accuse them of discrepancy **if John had kept the same time and the same place**, and had said that two of them were seen, sitting on the stone late on the sabbath. Correspondingly, if Matthew had maintained the time as early morning, as in John, and he too had said that one angel, not two, had been seen inside the tomb, he would plausibly have been regarded as writing a contradictory account."

> **That is a falsifiability criterion, and it is ours.** Same time, same place, same respect. Gates A1, A2 and A3 of `PREREGISTRATION-NT.md`, stated in the fourth century, **by a defender, against his own interest.**
>
> It also acquits him of the charge laid against Chrysostom, whose argument confirms on agreement and confirms on disagreement and therefore predicts nothing. **The man was not running a flat verdict.** The honest qualifier: he states the criterion and then satisfies it by multiplying events until the times and places differ. **Naming what would falsify you is necessary and it is not sufficient.**

Two article moves, two centuries apart, both dead

PAT1, a check written into `greek/check.py` and run against MorphGNT.

Augustine makes Mark's *hē paidiskē* a **different** girl at the second denial. But Mark 14:66 introduces her as *mia tōn paidiskōn*, one of the servant girls, and 14:69 has *hē paidiskē*, **singular and articular**, pointing back at that one. He settles it by preferring Matthew's *allē* and reading Mark against Mark's own article.

Chrysostom makes Acts 9:7's *tēs phōnēs* **Paul's** voice, so the companions heard Paul answering someone they could not see. But Acts 9:4, the voice Paul hears, is *φωνὴν*, **anarthrous**, first mention; 9:7 is *τῆς φωνῆς*, **articular**. There is no other voice in the pericope for that article to point at.

> **Two fathers, two centuries apart, two unrelated items, one illegitimate move.** Which suggests it is not carelessness but the natural shape of the pressure: when two texts collide, the cheapest exit is to make one of them refer to something else, **and the article is what stands in the way.**

And writing the check as a script caught two errors in the claim it was written to support: a postpositive δέ broke the article test, and Mark 14:66 is *mia tōn paidiskōn*, not anarthrous. Both left recorded in the code comments.

What survives on that item is neither rescue, and it was sitting unused in the text: Acts 22:9 is also articular and adds *tou lalountos moi*, "of the one speaking **to me**." The verse specifies whose speech is denied. That is an ACCOUNT-channel rescue with a cited licence, needing no case argument and no second voice. **Cost stays 2 on that, and would be 3 without it.**

The A2 gate, abusable in three directions, and all three now caught

Augustine on the staff, *De consensu* 2.30: Mark's staff is spiritual authority, Matthew's and Luke's staff is a stick.

> **Grade: BAD, and the most instructive BAD in the file.** The theological point is not arbitrary, and he is not pretending the problem is absent. But nothing in either text marks *rhabdos* as figurative in one gospel and literal in another; the word sits in a list with bread, bag and money in both. **This is the A2 move performed with no licence, and A2 is the exact gate that fired F4 against this project at 40.5%.**
>
> So the gate cuts three ways. **A2 is abusable by the accuser** (call everything a contradiction), **by the defender** (call every clash a difference of sense), **and by the coder** (dissolve what feels fine). This is the strongest argument in the project for the gate structure being worth having at all.

Origen, the outlier, who priced the trade first

Commentary on John X, on the temple cleansing:

> "I conceive it to be impossible for those who admit nothing more than the history in their interpretation to show that these discrepant statements are in harmony with each other."
>
> "**The spiritual truth was often preserved, as one might say, in the material falsehood.**"

The earliest sustained commentator on John, working this exact passage, **refuses the historical harmony** and does not propose two cleansings.

Guards, both directions. This does not make the item a cost 4: the scale prices rescues, not authorities. It also does not hand the item to the accuser, because Origen is not saying the evangelists erred; he is saying they composed with a purpose ranked above chronology, which is nearer the GENRE channel than a confession of error. **What it destroys is the inerrantist reading and only that.**

> He is the only father here who declines to harmonise, and he pays for it with a sentence no inerrantist can use. **Everyone else who wants to keep both the history and the harmony has to buy an event, a sense, or a preference. That is the trade the whole ledger is measuring, and Origen priced it first.**

The conclusion protected at the level of what counts as evidence

Augustine, *Letter* 82, and the file corrects its own earlier reading of it:

> "if in these writings I am perplexed by anything which appears to me opposed to truth, I do not hesitate to suppose that either **the manuscript is faulty**, or **the translator has not caught the meaning**, or **I myself have failed to understand it**"

> It is **not** a stray pious aside. It is the load-bearing premise of an argument: because no error can be admitted anywhere without the whole authority collapsing, every apparent error must exit through one of three doors. "The author erred" is not on the list.

Eusebius does the same thing one level down, in textual criticism: a variant is extraneous "most particularly if it contained something **contradictory to the evidence of the other evangelists**." Harmonisation driving the choice of text.

> **This is what separates the patristic version from the modern one. The fathers announce that the conclusion is fixed and then reason from it. The modern defender performs the identical operation and presents the output as a finding. Announcing it is more honest, and it is still the thing that makes the enterprise unfalsifiable.**

And the fathers were not one voice

Galatians 2:11-14, Paul rebuking Peter. **Jerome held the rebuke was staged**, a piece of theatre for weak Jewish believers, so Peter was not really wrong and Paul was not really opposing him. **Augustine held it was real**, and fought a long correspondence over it, arguing that admitting a single expedient falsehood undermines the whole authority of Scripture.

Two fathers, one text, incompatible readings, neither able to settle it.

Eusebius, HE _2.23.25, on James: it "is disputed; at least, not many of the ancients have mentioned it." James sat in the_ antilegomena. That kills "Luther was an outlier against settled tradition." It does **not** show the early church saw a Paul clash, because Eusebius gives no reason for the dispute, and reading one in would be supplying a motive the text does not supply, **which is the accuser's version of the same fallacy this file grades the defenders down for.**

The flag raised against the project itself

> **P8.** Eusebius spent three of his four *Problems to Marinus* on the Matthew-against-John Magdalene tangle. **Our docket rejected C0202 in one line at gate A3**, and A3 is a near neighbour of the gates that fired F4. Our rejection may well be right, and Eusebius in fact arrives where we do. But he needed a multi-visit reconstruction to get there and we needed a sentence. **When a fourth-century harmonist treats an item as one of the four hardest things in the resurrection narratives, and our gate disposes of it in a line, that is evidence bearing on whether the gate calls are too quick.**

What is missing, listed so the absences stay visible

1. **_To Stephanus_ 1-16 is unread**, on the infancy narratives and genealogies, bearing directly on four docket items. Already on disk. Highest-value unread source in the project.

2. *Augustine* Letter *199 unfetched*, so the "this generation" grade is provisional and tier K.
3. **Chrysostom under-read.** Two homilies fetched of ninety on Matthew alone.
4. **Nothing from the Greek fathers after Chrysostom**, nothing Syriac, and nothing on Ephrem's commentary on the Diatessaron, which would be the obvious place to see how a harmony handles what it cannot fit.
5. **No father found on C0264's actual clash, on C0397 outside Augustine, or defending C0512 head-on.** Those three absences are findings.

3.9b The apologists, scored by the same guard

`rescue/apologist-responses.md`. **Declared bias, written before reading:** "I am now sampling rescues from one side only. This pass can therefore only push costs *down*. That asymmetry is a property of the procedure, not a finding."

Outcome: no licensed cost changes. Twice, across two passes.

Seven failure modes shared across every source surveyed

#	MODE	WHAT IT LOOKS LIKE
F-A	**verdict flatness**	every item returns "dissolved / evaporates / really bad argument"
F-B	**bundling**	a licensed rescue presented alongside an invented one as interchangeable. The critic attacks the invented one and the licensed one dies with it
F-C	**unfalsifiability critiques that cut both ways**	correctly nailing an opponent for whom agreement and disagreement both confirm, while running the mirror structure oneself
F-D	strawmanning the thesis	answering a claim the opponent explicitly disclaims
F-E	speculation as the simpler explanation	"he could have", "maybe he got a copy", offered minutes after condemning the opponent's hypotheticals
F-F	ad hominem	in place of answering the argument
F-G	answering the popularisation, not the scholarship	engaging the video summary of a scholar rather than the scholar

The weak points, ranked by how badly they will be exploited

#	THE MOVE	WHY IT FAILS	COST
W1	Emmaus, "60 stadia is the round trip"	Luke 24:13 reads *apechousan stadious hexēkonta apo Ierousalēm*, "distant sixty stadia **from** Jerusalem". That syntax states a distance from a point	**3**
W2	Theudas	Rebuts the **naming** objection well. Gamaliel's problem is **chronological**: the speech is set c. 35 CE and Theudas revolted 44-46	**3**
W3	the fiction/hoax strawman	The opponent's stated thesis, in the clip played on screen, is "a gifted writer and a serious historian". The answer is built around what it would take for Acts to be *fiction*	n/a
W5	the Pervo remark	A smear standing in place of answering a thesis that is directly threatening: that Acts uses Paul's letters, which would generate **the appearance of undesigned coincidences**	n/a
W7	"the location of Qumran, which nobody got right except Luke"	**Luke does not mention Qumran.** One checkable error inside a list discredits the list	n/a
W9	leaning on undesigned coincidences	He cites Philip/Bethsaida, **which is literally one of the five fits logged in our own run 1**, in the section our own file flags as selection-inflated	

> ***W9 is the one that matters for this project.*** *"We cannot cite in public a metric our own instrument flags as selection-inflated." The rule bit the project's own strongest-sounding public argument.*

What holds, and should be kept

The Jericho climb is the strongest item: Luke's own text uses the directional verbs, a man "going **down** (*katabainō*) from Jerusalem to Jericho" (10:30) and Jesus "going **up** (*anabainō*) to Jerusalem" (18:31, 19:28). Cited, internal, checkable in ten seconds. Also *polis*/*kōmē* flexibility with Josephus *Ant.* 18.28, where Philip "advanced the village Bethsaida into the dignity of a polis", and Samaria as a city rebuilt by Herod as Sebaste, *Ant.* 15.296.

The harmony scored item by item

A full published resurrection harmony, against the five empty-tomb cluster items:

ITEM	THE HARMONY'S MOVE	LICENSED?
women at the tomb	John's plural "we" (20:2) implies unnamed company	**yes**, same licence already in our docket
number of angels	"unless the angels spoke simultaneously, one was the spokesperson"	**no**. No text says one spoke for the pair
inside or outside	Matthew reports the initial descent witnessed by the guards	**partly**. The guards are cited; the sequencing is not
sitting or standing	**not addressed at all**	n/a
the stone	cites Mark 16:4, that the stone was already removed	**no**. That cites the side already agreed and never engages Matt 28:2
arrival time	Matthew "reflects Bethany's perspective", Mark "John's Jerusalem household"	**no. Cost 3.** Source communities in different villages, invented to explain a time gap

> **On three of six items a simpler move was available that needs no unstated event, and the harmony passed it over.** Ours on the angels: naming one is not asserting only one. Theirs: an unattested spokesperson convention. Ours on the time: two texts supply the gradient across one journey, cost 1. Theirs: two source communities, cost 3.
>
> And it builds a **named roster of five women including Susanna**, whom no text places at the tomb at all, plus an uncited plausibility flourish that is doing work John's "we" already does for free.

Even the harmonisers' own camp is not united. Sympathetic reviewers of the strongest harmony on the market, reviewers who accept "reconcilable variation" as a category, still flag an incomplete incorporation of 1 Cor 15:3-8, an inadequate handling of Mark's ending, and a **speculative** sequencing of the women. On the exact empty-tomb cluster where our docket's cost-2 items concentrate. Which is what run 1 found independently: the empty-tomb pericope is the collision-dense one.

- **Origen, Comm. John X** concedes the temple cleansing cannot be harmonised as history: *"the spiritual truth was often preserved in the material falsehood."* Kills the inerrantist reading only, not the texts.
- **Augustine, Ep. 82.1.3** is verdict flatness in its earliest clean form: faulty manuscript, bad translator, or my failure to understand, and **"the author erred" is not on the list.** Same defect as the skeptic lists, mirrored.
- **Chrysostom, Hom. Matt 1.5-6** is Layer 1's own argument c. 390, and **as he states it, agreement confirms and disagreement confirms, so it predicts nothing.** Repairable only with a control and a denominator, which we lack.
- **Worst finding is a silence.** Augustine's *De consensu* 3.7 works the Judas pericope, the silver, the potter's field, the Jeremiah-for-Zechariah problem, and **never mentions Acts 1:18.** Verified.
- **PAT1**, a check written into `greek/check.py`: two fathers rescue two unrelated items by reading a definite article as non-anaphoric, and **both are dead.** Writing it as a script **caught two errors in my own claim**, both left recorded in the code comments.
- **Best patristic answer**, Augustine *De consensu* 3.25 on Galilee: "there is no statement of the precise time" is a licensed dissolution. Use that minimal form, never his ten-appearance list.
- **Correction to our own registry.** Africanus, now verified, makes **Matthew natural and Luke legal**, the reverse of the Machen defence. Two rows were bundling incompatible rescues with an "or."
- **Eusebius, Nicetas-Marinus 5** states a falsifiability criterion **against his own interest**: the charge "really would have been justifiable if John had kept the same time and the same place." **Same time, same place, same respect. Our Tier A, stated by a defender in the fourth century.**

3.10 Two findings that generalise beyond this project

Verdict flatness. Across the census videos, the Judas videos, the Acts livestream and the resurrection harmony, every item returned "dissolved." Across the skeptic lists, every item returned "contradiction."

> **A function whose output does not vary with its input is not reading its input.** This is why the cost scale exists at all, and why it catches both sides with one test.

Dissolution beats harmonisation, tactically. On three of six empty-tomb items the published harmony *reconstructed* where *dissolution* was available, and its rescues came out weaker than ours.

> "Naming one angel is not asserting only one" gives an opponent nothing to grip. "One angel was the spokesman, and these five women including Susanna came from Bethany while John's household came from Jerusalem" gives him **five joints**, and one broken joint makes the whole sequence look manufactured. The harmoniser buys attack surface for no evidential gain. **And harmonisation concedes the frame**, since producing one reconciled sequence accepts the inerrantist premise that the accounts must reduce to a single narrative.

Bundling, the third recurring failure: good cited arguments shipped welded to invented ones at equal confidence. A critic breaks the weak half and the good half dies with it.

3.11 What is missing, and why no number here settles anything

1. **No denominator.** Two hard contradictions out of *what*? Against thousands of cross-checkable claim-pairs that is remarkable; against forty it is bad. **Same numerator, opposite conclusions.** Not built. **The top priority.**
2. **No control.** The apocryphal corpus is on disk and the procedure has not been run on it, so **F3 is untested** and there is no evidence the instrument discriminates canon from legend at all.
3. **32 of 51 audits are tier K**, citations unverified. Only 19 are tier V.
4. **Stratum 2 (190 NT-vs-OT items) unscored. Number B, the independent sweep, not attempted.** So this is a ceiling on the accusers' case, never a floor on the truth.

A denominator may already exist and be sixteen centuries old. The Eusebian canon tables index every gospel section by which gospels carry it: an enumeration of the cross-checkable set made by someone with no stake in the result. It is a sampling frame, not a denominator, and building on it needs its own pre-registration.

`interlock/resurrection/PREREGISTRATION-C.md`, committed `4c614fc`, 5 August 03:17. There is no `RUN-C.md` and no `RESULTS-C.md`. The directory contains four files and this is the only one with nothing downstream of it.

Nine days later, two further studies were built and run. So it was not blocked. It was passed over.

> *This section exists because a record that lists only what was done is a biased record.* ***The best-designed file in the repository produced nothing.***

4.1 What it was going to test

`NEXT-R.md` had said no control exists for Layer 2 and none can be built, and called that permanent. **This file narrows its own project's earlier claim:**

> "That is too strong, and this file narrows it. It is true for the case as a whole. But **one specific sub-argument inside Axis 2 does have a control, and the corpus for it has been sitting in `../corpus/apocrypha/` since Layer 1.** The sub-argument is the restraint claim: *the canonical resurrection narratives do not look like what legendary development in this genre produces.* That is a comparative claim. A comparative claim needs a comparison set. We have one."

4.2 Why the docket rule is the anti-cheat again, in a new place

> "If I read both corpora and then decide which features count as 'embellishment,' I will pick the features the canonical texts happen to lack, and **the shaping will be invisible from the inside**.
>
> **So the feature list is not mine.**"

Sources: Bart Ehrman on legendary growth, and the form-critical tendencies of the tradition as stated in the literature. **A feature enters only if a named source states it as a tendency of legendary development.** Wanting to add a feature because you noticed it in the texts is a new pre-registration, not an amendment.

4.3 The frozen feature list

ID	FEATURE, AS THE CRITICS STATE IT
E1	the resurrection event itself is **directly narrated** rather than left off-stage
E2	opponents or authorities are made **eyewitnesses** to the vindication
E3	previously anonymous figures are **given names**
E4	the miraculous is heightened by **physical scale or spectacle**
E5	**numbers, times or measurements** are supplied where earlier sources have none
E6	**doubt, fear or failure** of the followers is removed or resolved
E7	**explicit proof** is offered to satisfy a stated objection

Test set: the resurrection material in the four canonical gospels. Control set: the Gospel of Peter and the Acts of Pilate, **both already on disk from Layer 1 and neither selected for this study.**

And it has the thing Layer 1 never had. A denominator: features present out of features possible per unit, with a written-in-advance NA rule "so it cannot be used to shed inconvenient units."

4.4 Five falsifiers, with the numbers fixed before any text was read

	FIRES IF	CONSEQUENCE
H1	the canonical set scores **above 0.40**	the restraint claim has no empirical content and comes out of the argument entirely
H2	the control set scores **below 0.50**	the apocrypha are not exhibiting the features the critics' own model predicts, so the feature list measures something else and **the instrument is invalid regardless of how the canon scores**
H3	the gap is **under 0.25**	direction without magnitude, which is what failed P5 in run 1
H4	more than **30%** of calls require a judgement not statable as a rule someone else could apply	this is the A2 problem that fired F4, in a new place
H5	the result depends on **any single narrative unit**	leave-one-out is run **before** the headline number is written, not after

> "Numbers chosen now, on principle, before any text is read. 0.40, 0.50, 0.25 and 30% are stipulated here so they cannot be moved later to accommodate a result. **Moving one afterwards voids the run.**"

Note **H2**, which is the design's sharpest feature: a falsifier that fires when the study's *opponent* fails to behave as predicted. Most instruments have no way to detect that their yardstick is broken. This one does.

4.5 Six confounds, declared against itself

> "These are not caveats to add at the end. They are limits on what any positive result can mean, and they are written now so I cannot discover them later and soften them."

	CONFOUND	STATEMENT
C1	**date**	the apocryphal texts are later, so a difference is consistent with "later texts in this tradition embellish more". The study cannot separate genre from date and will not claim to
C2	**survivorship, the worst one**	"The canonical four are the texts that **won**. If restraint was itself a criterion of canonization, then finding restraint in the canon is partly **circular**, because I would be measuring the selection rule rather than the composition process. **No result may be reported without this stated in the same paragraph.**"
C3	sample size	four canonical narratives and two apocryphal. "This is a descriptive comparison and **the word 'significant' does not appear in the results**"
C4	Gospel of Peter's date is disputed	if an early stratum is right, the control is contaminated toward the test set, which works **against** a positive finding. Record the direction either way
C5	sole coder, not blind	strong pilot, never confirmatory
C6	genre may not be matched	Acts of Pilate has a different purpose from a gospel. If a control unit is not attempting what the test units attempt, the feature is recorded **NA rather than absent**

4.6 The paragraph it licenses, written before the data existed

> "On seven features drawn from critics' own accounts of legendary development, the canonical resurrection narratives score X and two apocryphal narratives score Y. The canonical texts do not narrate the resurrection at all, and one apocryphal text narrates it with the stone moving of its own accord, figures whose heads reach heaven, and a speaking cross, with the guards and the elders watching. **The comparison is confounded by date and by survivorship, and it says nothing about whether the events occurred.**"

And the sentence that closes the file:

> "It does not answer the vision hypothesis, which is the alternative that matters most and which **predicts restrained early testimony just as comfortably as the resurrection does**. Lüdemann and Goulder are untouched by this entire study. **If a future session reports this result as evidence for the resurrection, it has broken the file.**"

4.7 The order of operations, and where it stopped

STEP	STATUS
1. Commit this file. Nothing before this step	**done**, `4c614fc`
2. Declare pericope boundaries for both corpora in `RUN-C.md`. Commit	**never done**
3. Code the **control set first**, so the feature list is exercised on the corpus I care less about before it touches the one I want a result from	never done
4. Code the test set	never done
5. Run leave-one-out and the H4 audit **before** computing the headline ratio	never done
6. Write results, falsifiers at the top, confounds in the same paragraph as any headline number	never done

> *Step 3 is the detail worth sitting with.* ***Code the control first, so the instrument is exercised on the corpus you care less about before it touches the one you want a result from.*** *That is a procedural safeguard against a bias nobody can feel operating, and it costs nothing, and it was written down, and it was never used.*

4.8 It was run on 14 August, and two falsifiers fired

Written after the rest of Part IV, which is left standing above because the question of what a nine-day-old unexecuted pre-registration is worth does not disappear once it is finally executed.

Order followed exactly as §7 specified. `RUN-C.md` declaring the 22 pericope boundaries and the applicability rule was committed at `0514fea` **containing only that file**. The control set was coded before the test set. Leave-one-out and the rule audit ran **before** the headline ratio was computed. Results `3725bcc`.

FALSIFIER	THRESHOLD, FIXED 5 AUG	OBSERVED	
H1 canonical ratio above 0.40	0.40	**0.352**	did not fire
H2 control below 0.50, instrument invalid	0.50	**0.577**	did not fire
H3 gap under 0.25	0.25	**0.225**	**FIRED**
H4 unstatable calls over 30%	30%	**9.8%**	did not fire
H5 result turns on any single unit	any	**4 of 22**	**FIRED**

H2 not firing is the important one. It means the apocryphal narratives do exhibit the features that critics' own model of legendary growth predicts, at 0.577, so the feature list is measuring something real and **the instrument is not invalid.** That was the outcome most likely to sink the study and it did not happen.

H3 firing is the result. The gap is 0.225 against a 0.25 threshold fixed before any text was read. Direction without magnitude, which is the same failure as P5 in interlock run 1.

Where the difference lives, and where it runs backwards

FEATURE	CANONICAL	APOCRYPHAL	
E1 resurrection directly narrated	**0 / 5 = 0.00**	**3 / 5 = 0.60**	categorical
E4 miraculous heightened by scale	0.15	0.56	strong
E3 anonymous figures named	0.08	0.33	as predicted
E5 numbers and measurements supplied	0.38	0.67	as predicted
E7 explicit proof offered	0.54	0.78	weak
E2 opponents made eyewitnesses	0.67	0.71	**tied**
E6 doubt or failure resolved	**0.73**	**0.25**	**BACKWARDS**

E1 is the one thing worth saying in public. Not one of the five canonical units that could have narrated the resurrection does so. Matthew comes closest and stops short: the angel rolls the stone back and sits on it, and the tomb is *already* empty, because "He is not here: for he is risen." Peter §9 opens the heavens, walks two men down, has the stone roll of itself, and brings three men out with a cross behind them. **It is categorical rather than a rate, it needs no threshold, and it survives every leave-one-out.**

> ***E6 runs the other way and it is not a small effect.*** *The canonical texts resolve the followers' doubt and failure at nearly three times the apocryphal rate. Thomas puts his hand in the side. The Emmaus pair go from "we trusted that it had been he" to burning hearts. He eats a piece of broiled fish in front of them. Meanwhile the Gospel of Peter leaves the women fleeing in fear and the twelve grieving and going home.*
>
> *Confound **C6**, declared on 5 August, anticipated the mechanism: the apocryphal units are about opponents and cosmic events, so E6 is applicable 11 times in the canon and only 4 times in the control.* ***The feature is measuring what each corpus is interested in, not how each corpus develops.***

The advance prediction that was confirmed, and cost the study its result

`RUN-C.md` recorded, **before coding**, that Mark 16:9-20 is later disputed material, that including it should **raise** the canonical ratio and **work against** the restraint hypothesis, and that it was kept for exactly that reason.

It scores **4 of 5, the highest of any canonical unit**, and removing it takes the gap to 0.259, one of only two ways to stop H3 firing.

> **The conservative choice was correctly identified in advance, made anyway, and it is what cost the study its magnitude test.** That is what a pre-registration is for, and it is the cleanest instance of the discipline paying its own bill anywhere in this repository.

The two cleanest canonical units are the two earliest. Mark 16:1-8 and Luke 24:1-12 both score **0 of 6**. Nothing fires on any feature in either.

And the control's result is carried by one unit. Dropping the Descent into Hell, a perfect 6 of 6 and the latest stratum in the Acts of Pilate, takes the control from 0.577 to 0.522 and the gap to 0.170.

What it licenses, and the sentence the file forbids forgetting

Drop "the canonical accounts are strikingly restrained compared with the apocrypha" as a general claim. Three of seven features carry the whole result and one runs the other way.

Keep only E1, stated exactly.

> "It does not answer the vision hypothesis, which is the alternative that matters most and which **predicts restrained early testimony just as comfortably as the resurrection does**. Lüdemann and Goulder are untouched by this entire study. **If a future session reports this result as evidence for the resurrection, it has broken the file.**"

And **confound C2 is untouched**, which is the worst of them. If restraint was itself a criterion of canonization, this measures the selection rule rather than the composition process. No study on this corpus can separate them.

4.9 What the existence of this file means

It cuts two ways and both should be said.

For the project. This is the most disciplined pre-registration in the repository. It has a control, a denominator, five falsifiers with numbers fixed in advance, six confounds declared against itself including one that admits possible circularity, an ordering rule that front-loads the unfavourable corpus, and a closing sentence forbidding its own future misuse. Nobody writes that file to launder a result.

Against the project. It was also the cheapest study in the repository. Both corpora were already on disk. Seven binary features. Twenty-two narrative units. It took one working session, and the two studies run instead of it both failed. **A pre-registration with nothing behind it is a promise, and the value of the discipline is in the running, not in the writing.**

And running it settled the question the section opens with. The file was worth writing. H2 not firing means the feature list detects what it claims to, which the writing alone could not have shown. H5 firing means the result turns on two units out of twenty-two, which no amount of design could have anticipated. And the advance note about Mark's longer ending was right, and cost the study its magnitude test. **None of that existed on 13 August. All of it is checkable now.**

Ten statistics. For each: the formula, what it holds constant, what it can see, what it cannot, and where it broke.

4.1 Rate per thousand

$$r\_b(T) = 1000 \cdot h\_b(T) / N(T)$$

Holds constant: text length, by dividing it out. **Can see:** how saturated a text is with a vocabulary. **Cannot see:** where anything is. Position is discarded entirely. **Broke:** everywhere. This single property is the root cause of rounds 3 and 4 both failing, and of round 2's Bacchae failure.

4.2 SUM

$$SUM(T) = \Sigma\_b\ r\_b(T)$$

Can see: accumulation across four vocabularies. **Cannot see:** whether they occur together. **Broke:** Odyssey book 7, a hospitality scene with zero violence, outscores every tragedy on `marks` because "king" and "queen" saturate it.

4.3 MIN

$$MIN(T) = min\_b\ r\_b(T)$$

Can see: conjunction. A text scores only as high as its weakest stereotype. This is the algebraic form of Girard's actual claim. **Cannot see:** anything, once any bucket is likely to be zero. **Broke:** scale-dependence. With P(zero) = exp(−λ_b·N/1000) and λ_crimes ≈ 0.95, a 1,500-word episode has a 24% chance of zero crimes hits even under the repaired lexicon. MIN saturates at 0.00 and carries no information.

> **The transferable rule.** A metric validated at one scale is not validated at another. MIN was built on 10,000-word tragedies and applied to 1,500-word episodes.

4.4 peak_min

$$peak\_min(T) = max\_i\ [\ min\_b\ c\_b(W\_i)\ ],\ |W| = 1000,\ stride\ 250$$

Can see: in principle, whether a text contains a passage where all four converge. **Cannot see:** in practice, nothing that the whole-text rate does not already say. Round 3 proved this empirically. **Broke:** two ways at once. **Invariance**, since the null showed word order is irrelevant to it. And **power**, since it is a minimum over four counts with median 2. A statistic built on counting to two.

Computed in linear time by prefix sums: `c[b,j] = Σ_{k<j} ind[b,k]`, then `counts = c[:,ends] − c[:,starts]`.

4.5 The shuffle null (Null A)

$$z = (obs - mean\_s) / sd\_s \; ; \; p = (\#\{s : stat\_s \geq obs\} + 1)/(S + 1)$$

Holds constant: the exact multiset of tokens. Every bucket's total count and whole-text rate are unchanged. **Destroys:** position, and nothing else. **Therefore isolates:** locality. This is the cleanest null in the project and it is immune to the density confound. **Result:** median z = 0.00 across 71 texts. 4 of 71 significant against 3.6 expected by chance. None a named HIGH text.

Implementation detail that is load-bearing: phrase hits are marked at their first token so a permutation moves each phrase as a unit. Shuffling raw words would dissolve every phrase and make the null trivially beatable.

4.6 The frequency-matched lexicon null (Null B, and D4)

Draw replacement words matched on corpus frequency. Round 3 matched on surface count and frequency; round 4 binned by $\lfloor \log_2(\text{count}) \rfloor$.

Round 3, Null B, was a bad design and the reason is structural:

> A **minimum** over four sets is raised by uniformly-spread sets and lowered by clumpy ones. Frequency-matched random draws spread across topics, so they are less clumpy, so they beat any thematically coherent lexicon on a minimum. Null B would fail identically for any four topical word sets. **It tests nothing.**

Round 4, D4, failed for a different and more instructive reason: the null matched **unigrams only**, and the entire effect lives in the **phrases**. On unigrams the Gospels score lowest of the three tiers and the observed separation is negative (−0.075 per 1000). So 591 of 2000 random sets beat it. The half of the instrument that produces the whole result **has never been null-tested at all.**

4.7 Mann-Whitney U by label permutation

$$U\_a = R\_a - n\_a(n\_a+1)/2 \; ; \; p = (\#\{perm : R\_a\text{^}perm \geq R\_a\text{^}obs\} + 1)/(NPERM+1)$$

Midranks for ties. NPERM = 200,000.

Why permutation rather than the normal approximation: n = 4 for GOSPEL. The normal approximation to U is invalid there. This is the better test, not a workaround for a missing scipy.

The vectorisation is exact. Ranks are computed once on the pooled sample and labels are permuted over the ranks. Since the rank vector does not depend on the labels, this is algebraically identical and about a thousand times faster.

4.8 The extreme-cell test (conversion P2)

$$E(M) = \#\{ \text{ rows at } (0,0,0) \text{ or } (1,1,1) \}$$

Null by permuting **each column independently**, 10,000 times.

Holds constant: each bit's marginal rate. **Destroys:** all association between the bits. **Therefore isolates:** whether the three properties travel together. Which is exactly claim (i): is conversion one event with three faces, or three unrelated properties under one word. **Status: never ran.** No control set was ever fetched.

4.9 The Bayes lever

$$B\_total = \exp(\ \alpha \cdot \Sigma\_i \log B\_i\ )\ ;\ odds\_post = odds\_prior \times B$$

Can see: whether the pattern of fits, agreements, dependences and collisions favours uncoordinated multiple sources over single fabrication. **Cannot see:** truth. Independent records of a shared legend also satisfy H_multi. Only dating constrains that, and dating is not in the file. **The fragile part:** α is a judgement about how correlated the items are, not an estimate from data. It is reported at three values (1.0, 0.5, 0.33) precisely because it cannot be estimated.

4.10 The rejection-histogram falsifier (F4)

$$fire\ if\ (\ n\_A2 + n\_A5\ )\ /\ n\_rejected > 0.40$$

Observed: 69/102 = **67.6%**. Fired.

Why this is the most interesting statistic in the repository. It does not measure the texts at all. It measures **the coder**. The threshold was set by predicting in advance which two gates a motivated thumb would rest on.

> **Naming a number beforehand is the only mechanism that can catch a bias you cannot feel while committing it.**

Each rule with the specific failure it exists to stop, and what breaks if it is relaxed. Compiled from `DESIGN-RATIONALE.md` and the myth-side pre-registrations.

#	RULE	THE FAILURE IT PREVENTS	IF RELAXED
1	Pre-register and commit before scoring	A rule written afterward gets shaped by what was found, **and the shaping is invisible from the inside**	Every number becomes uninterpretable
2	The sample is part of the registration, not just the instrument	A text found after seeing the results, added as a bug fix, flips a failed prediction to a pass	Correction 1's Diodorus page
3	Never repair the instrument to make a text pass	Adding dismemberment vocabulary to `crimes` to catch Pentheus	The instrument stops measuring and starts confirming
4	Tag new concepts so their effect stays separable	Cannot tell a bug fix from a theory change	`score.py --no-new` becomes impossible
5	Explicit surface forms, no stemming	`tear-` catches "tears", `torn-` catches "tornado"	Matches stop being auditable
6	Declare sensitivity analyses in advance	Choosing after seeing which way it helps	S1 and S2 would never have been run
7	Disclose contamination before scoring	Knowing John's division formula while writing the lexicon	The disclosure cannot be added later without looking defensive
8	A falsifier that fires is reported at the top	Burying F4 in a footnote	A falsifier that never fires is decoration
9	Show what survives, and why **structurally**	Asserting survival	Only a directional argument counts
10	The hostile compilers pick the docket	Choosing the easy items	Result becomes a selection, not a ceiling
11	Layer 2 inverts: the apologetics site never supplies the docket	The apologists pick the docket	Layer 2 is **void, not weakened**
12	A rescue counts only if something outside the coder's head licenses it	Both invented harmonies and lazy dissolutions	Plausibility measures the coder, not the text
13	Record the **first** failing gate, in fixed order	Coder preference disguised as a histogram	Coding stops being reproducible
14	Record findings that cut **for** the other side	A record where every entry is true and the set is biased	The spread is the evidence
15	Fix the ceiling of what a study can show, in advance	Overselling a good result later, including to yourself	Layer 2 gets read as proof of a miracle
16	Apply your own discount to yourself first	Discounting the other side's unauditable premise but not your own	The study is worthless
17	Corrections go against the story	Judas 4 → 3 was rhetorically worse and was made anyway	The scale stops outranking the instinct
18	Tier your own citations, V and K, never merged	51 rows of uniform confidence, of which 32 are unchecked	The file lies about its own strength
19	Name the escape hatch before you reach it	"So these are false positives, therefore the theory survives"	The escape gets taken and feels like reasoning
20	No re-running with a wider net until something passes	Round 3's window at 1000/250	A new operationalization is a new pre-registration

> **The one rule under all of them.** Never tune the instrument, and never tune the sample, to pass. Everything else is machinery for making that rule enforceable **by someone who does not trust the person who wrote it, including when that person is me.**

Every occasion on which this project caught itself. Ordered by how instructive it is, not by date.

1. The bug fix that would have flipped a failure into a pass. `CORRECTION-1.md`. A fetcher bug truncated one text. Fixing it also surfaced a corpus page that had been missed. Adding that page moves MYTHOGRAPHY 13.95 → 13.67 against TRAGEDY 13.88, turning a **failed** pre-registered prediction into a **passing** one. Every step was defensible as a bug fix. **The file's own words: it would have been entirely defensible, and none of it felt like cheating.** The page is on disk and excluded.

2. The instrument one line from being repaired to pass. `RESULTS.md`'s own next-action list said to add dismemberment vocabulary to `crimes` so it would catch Pentheus. It is a category error: `crimes` is what the victim is accused of, the sparagmos is what the mob does. What caught it was a **definitional** distinction, not a statistical one. The vocabulary went into `violence` instead.

3. The contamination disclosure that came true. Round 4 §7 disclosed in advance that `no_fault` and `division` were at risk of being motivated, because John's division formula was already known. S1 then showed the `core` half carries **none** of the Gospel effect and the disclosed half carries **all** of it. The disclosure was not a formality. It was correct.

4. A pre-declared falsifier that fired on the coder. F4 at 67.6% against a 40% ceiling. The threshold was set by predicting in advance which gates a motivated thumb would rest on. It rested there.

5. A null test that tested nothing, admitted as a design error. Round 3's Null B. A minimum favours uniformly-spread sets over clumpy ones, so frequency-matched random draws beat any coherent lexicon by construction. Recorded as "my error, in the pre-registration."

6. A null test that tested the wrong half. Round 4's D4 matched unigrams only, and the entire effect lives in the phrases. The half producing the whole result has never been null-tested. Recorded as a defect in the test design, and deferred to round 5 with its own commit rather than rerun.

7. A correction that overstated its own scope. The first pass on Correction 1 claimed five truncated files. Wrong: the search was case-insensitive and matched the ordinary word "continued." Recorded, on the ground that a correction which overstates its scope is not a correction.

8. A diagnostic output read past twice. The `CONTINUED` pagination marker was in the very first diagnostic output and was dropped into a list of ignored singletons. It was found by accident, on the way to something else.

9. A validation metric that could not fail. `NEXT.md` on speaker attribution: coverage sits at ~99% whether or not attribution is correct, because every word between two tags is assigned to *someone*. A missed speaker's lines are silently credited to the previous speaker. **A broken parse looks perfect.** An early version dropped `DEÏANIRA` and put her speeches on the Nurse, who came out at 61.5% of the play instead of 12.3%. Coverage was 99% for both runs.

10. A Greek claim corrected by writing it as a script. PAT1. Writing the article test as code caught two errors in the claim it was written to support: a postpositive δέ broke the test, and Mark 14:66 is *mia tōn paidiskōn*, not anarthrous. Both left recorded in the code comments.

11. A case grammar error caught by fetching the primary. C0443. Acts 22:7 had been recorded as accusative; it is genitive. Found on the first five K→V promotions.

12. Three claims corrected by reading a source instead of a blog. Eusebius *Ad Marinum*. He does **not** rank his two solutions, both are attributed to other people, and the second view explicitly refuses to choose. Three claims made from a blog summary were wrong. **The tier system is worthless if a K row gets quoted as V.**

13. An error in favour of the other side, published. The registry had Africanus backwards: he makes Matthew **natural** and Luke **legal**, the reverse of the Machen defence, so two rows were bundling incompatible rescues with an "or."

14. A correction made against the rhetorically stronger position. Judas 4 → 3. A lone hard contradiction is a cleaner story than a cost 3. The correction was made anyway, and the Bayes value for cost 3 was chosen on principle **before** running it, so the amendment came out Bayes-neutral.

15. A gate abused by all three parties, including us. A2, "not the same respect." Augustine's staff solution is gate A2 with no licence. The skeptic lists abuse it. So did we, at 40.5%. **All three caught by the same rule.**

16. An escape hatch named in advance so it could not be taken. Round 4 D6. "These counterexamples are false positives, therefore the universal negative survives" is refused in the file, because an instrument that produces non-counterexamples **cannot test a universal negative in either direction.** Recorded as *untested*, not as *survived*.

17. A rubric applied more strictly than written, and applied to everything. The conversion instrument's L bit as written let silence score. The stricter reading, requiring positive regard, was applied to **every** text and not only to the one that needed it. Quixote scores under it. Plath does not. Recorded because it changed a score.

18. A power problem foreseeable in advance and not foreseen. Round 3. `peak_min` is a minimum over four counts with median 2. `RESULTS-3.md`: "this was foreseeable before running it and was not foreseen."

Everything above is what the project says about itself. This part is what an outside reader finds when they check it. Two of these are not recorded anywhere in the repository.

8.1 What the git timestamp actually proves

The whole discipline rests on one sentence, repeated in every pre-registration: *committed before any item was scored*, and *the git history is the evidence*.

So here is the git history, in full.

ROUND	PRE-REGISTRATION	RESULTS	GAP
R1	21 Jul 01:37	21 Jul 01:39	**2 minutes**
R2	21 Jul 03:00	21 Jul 03:14	**14 minutes**
R3	22 Jul 00:55	22 Jul 03:22	2 h 27 m
NT count	24 Jul 22:53	24 Jul 23:09	**16 minutes**
Layer 2	25 Jul 00:56	25 Jul 00:58	**2 minutes**
R4	11 Aug 02:07	11 Aug 02:17	**10 minutes**
Conversion	12 Aug 16:50	12 Aug 17:12	22 minutes

Median gap: 14 minutes. Every pre-registration and its results were committed inside a single continuous working session.

> *A hostile reader's objection, stated at its strongest: you had the numbers on screen and committed the prediction file first. Git proves the order in which two files were written to disk. It does not prove that no scoring happened before the first one was written.*

This objection is real and it is not answered anywhere in the repository.

The defence that does survive, and it is stronger than the timestamps

Check **which files** land in each commit, not just when.

COMMIT	CONTAINS	DOES THE SCORING CODE EXIST YET?
`fc094bb` PREREG 3	`PREREGISTRATION-3.md` only	**No**
`4188308` RESULTS 3	`RESULTS-3.md`, **`nulltest.py`**, `results_nulltest.csv`	arrives here
`9310c30` PREREG 4	`PREREGISTRATION-4.md`, `dissent.py`, `tiers.py`, `build_corpus3.py`	**No scorer**
`0c765d0` RESULTS 4	`RESULTS-4.md`, **`score_dissent.py`**, `results_dissent.csv`	arrives here
`d671056` PREREG conv	`conversion/PREREGISTRATION.md` only	**No**
`cf4bfa0` RESULTS conv	`RESULTS.md`, **`analyze.py`**, **the entire corpus**	arrives here

> For rounds 3, 4 and the conversion study, git proves something **stronger than the order of two prose files**: the program that produced the number **did not exist** when the prediction was frozen. `nulltest.py` was not on disk when `PREREGISTRATION-3.md` was committed. Neither was `score_dissent.py` when round 4's predictions were frozen. The conversion corpus had not even been downloaded.

That is the argument to make. Not the timestamp. The file manifest.

And here is where it does not hold

COMMIT	CONTAINS
`00a9556` **PREREG 2**	`PREREGISTRATION-2.md`, `lexicon.py`, **`score.py`**, **`results_corpus_v1.csv`**, **`results_corpus_v2.csv`**

Round 2's pre-registration commit contains **the scorer and two files of scored results**. Those are `corpus/` scores, not `corpus2/`, which is consistent with the smoke test the pre-registration itself discloses. But the machinery was fully operational and committed at the moment the predictions were frozen, and the gap to results is 14 minutes.

Round 1 is weaker still. Two minutes, and the scorer already existed from the previous commit.

> ***The sharp point.*** *Round 2 produced p = 0.0045, which is the* ***only surviving positive result in the entire repository****. It is also the round with the* ***weakest file-level protection****. And per* `CORRECTION-1.md` *the script that computed it was never saved, so* ***it could not be reproduced.***
>
> *The one number the project had left was the one number that was both least protected by the procedure and impossible to check.*

> **Resolved in part, 14 August 2026. See §1.18.** The script was reconstructed under a spec committed before the scorer existed, across all 48 defensible configurations. The original procedure is **identified** by two committed statistics that are not the p-value (the null mean and sd), and the committed result **reproduces** under it at p = 0.0005 to 0.0070.
>
> **What is not resolved.** The pre-declared verdict rule returns **UNSTABLE**, because half the grid sits above 0.05. The committed effect size **+6.10 still does not reproduce exactly** from its own CSV. And the file-manifest weakness above is unchanged: round 2's own pre-registration commit still contains the scorer and two files of scored results. **A reconstruction three weeks later does not retroactively improve the protection round 2 was run under.**

8.2 A rule that was broken two weeks after it was written

`PREREGISTRATION-2.md` §1, 21 July: *"Minimum length 800 words. Excludes* `hyginus_fabulae_2`, `_3`, `_4`*. 71 texts scored."*

`CORRECTION-1.md`, 22 July: *"the sample is part of the registration."*

`tiers.py`, 11 August: **contains no length rule at all.** Recomputed from `corpus3/`:

TEXT	WORDS	TIER IN ROUND 4	STATUS IN ROUND 2
`hyginus_fabulae_2`	613	MYTH	**excluded**
`hyginus_fabulae_3`	689	MYTH	**excluded**
`hyginus_fabulae_4`	777	MYTH	**excluded**

Round 4 silently re-admitted exactly the three texts round 2 excluded on a stated rule. The rule change is not declared in `PREREGISTRATION-4.md`.

The effect, recomputed from `results_dissent.csv`:

	AS RUN, N = 71	800-RULE APPLIED, N = 68
MYTH median rate	0.547	0.561
myth texts above lowest Gospel (D6)	**7**	**6**

No verdict changes. D6 fails either way. D1's ordinal still holds. And the direction is mixed rather than helpful: two of the three re-admitted texts score 0.000, which lowers the myth median and **helps** D2 and D3, while the third scores 1.631 and **adds a counterexample against** D6.

> **This is exactly Correction 1's finding, recurring.** There the excluded page would have flipped a verdict and was caught. Here the change flipped nothing, and was **not** caught. The discipline is not "did it matter." The discipline is "was it declared." A rule that only binds when someone is watching is not a rule.

`RESULTS-4.md` §4 notices the **symptom** (hyginus at 613 words is "a rate artifact on a short text and nothing else") without noticing that a pre-registered exclusion rule had lapsed.

8.2b The patristic stratum is not in git, and it says it is

Found 14 August, incidentally, while checking the repository was clean after committing round 5.

`interlock/ntcount/PATRISTIC-HARD.md` §0, on how its ranking of the Hard Twelve was protected:

> "The ranking below was written and saved to this file **before any source was fetched for this round.** That ordering is the only thing stopping the list from being quietly reshaped by which fathers turned out to have good answers. **If this file is in git, the commit order is checkable by someone who does not trust me.**"

It is not in git. Neither is `PATRISTIC.md`. Neither is the `patristic/` directory holding the primary sources that make a tier-V patristic quotation recheckable.

```
$ git ls-files --error-unmatch interlock/ntcount/PATRISTIC-HARD.md
error: pathspec did not match any file(s) known to git
```

And five tracked files carry uncommitted modifications from the same session:

FILE	UNCOMMITTED CHANGE
`interlock/ntcount/audits.py`	+193 lines
`interlock/greek/check.py`	+80 lines, **including the PAT1 check**
`interlock/ntcount/REGISTRY.md`	58 lines changed
`interlock/greek/VERIFIED-greek-checks.txt`	+54 lines
`interlock/README.md`	17 lines, the reading order that **points at the two untracked files**

The timestamps place all of it on 25 July between 18:32 and 18:50. The last commit that day was 01:30. So an entire evening's work, 1,284 lines of new analysis plus 332 lines of changes, was written and never committed, and has sat uncommitted for twenty days.

Why this is worse than the two findings above it

> ***The one thing that file names as its own verification does not exist.*** *§8.1 found that the git gaps are minutes and argued the file manifest is the real evidence. Here there is no manifest and no timestamp at all. The claim that the D1 to D7 ranking preceded the source-fetching is exactly the kind of claim this project refuses to accept on anyone's word,* ***and it is currently resting on my word.***
>
> *It is also the section that grades other people's honesty. The Hard Twelve file is where Augustine is marked BAD for an unlicensed sense-distinction and Eusebius is marked BAD for reading the number of events off the number of sources.* ***A survey that audits defenders for unverifiable warrants is itself, at this moment, an unverifiable warrant.***

What must not be done about it

Do not commit it now. A commit today stamps 25 July work with 14 August, which in a project whose entire discipline is commit ordering would be the single most misleading action available. The honest options are to commit it with the discrepancy stated in the message, or to leave it and record that the patristic stratum carries **no ordering evidence** and that its §0 claim is unverifiable.

Either way `PATRISTIC-HARD.md` §0 needs amending, because as written it points at a guarantee that is not there.

8.3 Multiplicity, and why it is less bad than it looks

Roughly 30 named predictions and falsifiers across five pre-registrations, with **no multiple-comparison correction anywhere.** At $\alpha = 0.05$, 30 predictions produce about 1.5 false passes by chance.

Why this is not fatal here. Multiplicity inflates false **passes**, and this project's results are overwhelmingly failures. Multiplicity cannot manufacture 14 failures.

Where it does bite: round 2, which is where the surviving result lives. Seven predictions in one round. Bonferroni at 7 gives $0.0045 \times 7 =$ **0.0315**, still under 0.05. So the one number that matters survives the correction, narrowly. **This calculation is not in the repository and should be.**

8.4 Four more objections a hostile reader will raise

1. All the permutation tests are one-sided. Legitimate under pre-registration, since the direction was committed in advance. But it doubles the nominal α relative to a two-sided test, and nowhere is that stated.

2. The four Gospels are not four independent observations. Round 4's permutation test treats GOSPEL as $n = 4$ exchangeable draws. Matthew and Luke use Mark. Under any standard view of Synoptic dependence the **effective** n is smaller than 4, so $p = 0.00034$ overstates the evidence. The pre-registration discusses translation matching and register at length and **never raises literary dependence.** This is a real gap and it is not recorded anywhere.

3. The most load-bearing control in the project has n = 1. P5, translation robustness at 17.2%, ran on **exactly one paired play**. `RESULTS-2.md` says it "should be repeated on more paired translations." It has not been. Every cross-text comparison in three instruments rests on one Agamemnon.

4. The conversion study's file is called `RESULTS — run 1`. Its primary test (P2) and its group comparison (P3) **never ran**, because Set B was never fetched. The file states this in §0, plainly. But the title and the commit message ("the conversion instrument fails P1 and P4") describe a study that produced two failures, when what actually happened is that **two of four predictions were never tested at all.**

8.4b Six more, found by reading the files against each other

5. One gate is doing two different jobs, and nothing separates them. Set out in Part III with worked items. A2 rejections split into two visibly different kinds: ones citing a lexeme or a verse in the same passage (*barē* against *phortion*; John 4:2 correcting John 4:1; *hōs enomizeto*), and ones citing a convention (rhetorical quantifiers, inclusive reckoning, Semitic idiom). The first kind carries its licence. The second kind is a judgement about register. **They are coded identically.** F4 detected the aggregate as a number without naming the mechanism, and splitting A2 into A2-lexical and A2-conventional would have made the histogram interpretable at no cost.

6. The gap typology produced almost nothing and this was never noticed. Recomputed from `run/coding_sheet.csv`: the six-type G1 to G6 scan yielded **7 tagged items of 39**, and the flat-collision scan, added as a patch after Pilot 1, supplied the other 32. The theoretically motivated half of the interlock instrument did four fifths of nothing.

7. The interlock ratio has a numerator of five. Across the whole run there are **5 FITs**, and the headline statistic is FIT over FIT plus collisions. `RESULTS.md` already discounts the fit ratio for selection. It does not say that the number is also this small. Five items, coder-chosen, from the famous list.

8. "Strong structured pilot, not a confirmatory result" appears on every headline and is never operationalised. The phrase carries a real commitment, but no file states what would upgrade it, beyond "a future κ from any independent coder." No sampling plan, no target, no candidate. It functions as a standing disclaimer rather than a plan, and a disclaimer that never expires stops being read.

9. Two studies report a fired falsifier as the headline and one does not report a study that never ran. F4 and G1 are at the top of their files, in bold, exactly as the rules require. Pre-registration C's non-execution appears nowhere: not in `NEXT-R.md`, not in the interlock `README.md` reading order, not in any handoff. **The discipline covers failures and does not cover omissions**, and an omission is the easier one to hide because nothing was produced to look at.

10. The Bayes lever's α is doing more work than any other single number in the project, and it cannot be estimated. B ≈ 72.7 is quoted at default and α = 0.5. The α exponent discounts for the items not being independent of one another, and it is a judgement, not an estimate from data. `bayes.py` reports three values precisely because it cannot be pinned. So the honest form of the result is not a number, it is a range across an assumption nobody can check, and the range is wide: the same evidence at α = 1.0 and α = 0.33 differs by orders of magnitude. **Quoting 72.7 alone, as I did in Part III, drops the thing that qualifies it.**

8.5 What an outside reader would find *strong*

Stated for balance, since Part VII is a list of errors and a record of only errors would be as biased as a record of only successes.

- **Falsifiers that actually fired**, twice, and were reported at the top of their files rather than in footnotes.
- **Sensitivity analyses registered in advance that destroyed the round's own headline.** S1 and S2 are the strongest evidence in the repository that the procedure is real, because a project trying to look good would not have written them.
- **Contamination disclosed before scoring, and then confirmed by the data.**
- **A discount applied to itself first.** The run-1 fit ratio was discounted for coder selection before the same discount was applied to Habermas.
- **Escape hatches named in advance so they could not be taken.**
- **The file manifests**, per 8.1, which are better evidence than the timestamps the files themselves point to.
- **Real blinding tooling**, per Part III: opaque witness codes, a sealed key file, and gaps committed per witness before the key is opened. Then the limit of that tooling stated in the same section.
- **A gate that fired zero times, printed anyway** (A7), on the ground that a gate which never fires is evidence the others were not free passes.
- **A core flag removed when the frozen definition did not support it**, which made the study's own headline easier to satisfy, and was recorded as a discrepancy rather than absorbed.
- **The single most famous harmonising move refused.** "There were two temple cleansings" was available, and the item was admitted at cost 2 instead, with the abuse risk of that gate named in the same paragraph.
- **Roughly 960 junk hits deleted from a lexicon whose author wanted it to score high.** Bare `stone` at 211, bare `cast` and `drive` at 749. Every removal lowers the numbers.
- **Weak joints marked in public**, such as C0423 "drifts toward 3 if Mark's absolute phrasing is pressed."
- **A rule that bit the project's own best public argument.** W9: the Philip/Bethsaida coincidence cannot be used in public, because the project's own instrument flags it as selection-inflated.
- **Pre-registration C**, which is the best-designed file in the repository even though it produced nothing. Part IV.

> ***Read this box before the numbers.*** *Everything in this part was computed on* ***already-scored data****, after the results were known. Under the project's own rules that makes all of it* ***exploratory and post-hoc****. None of it can be cited as a test, none of it rescues a failed prediction, and none of it overturns one. Each is a* ***hypothesis for a future round with its own commit****. Analysis 9.1 in particular produced exactly the shape of temptation the whole repository exists to refuse, and that is said in 9.1 rather than in a footnote.*

Three of these four were named in the repository as things that needed doing and were never done. Seed `20260814` throughout.

9.1 The density confound, finally measured

`RESULTS-2.md` called this "the main confound and it is unresolved" and its Next list, item 1, said: "score against a per-text baseline of narrative event-words." Never built. Built here.

Method. A neutral event lexicon with nothing Girardian in it: *went*, *came*, *said*, *spoke*, *took*, *gave*, *saw*, *made*, *sent*, *brought*, *answered*, *asked*, *called*, *cried*, *heard*, *rose*, *stood*, *sat*, *put*, *set*, *found*, *left*, *turned*, *looked*, *fell*, *ran*, *walked*, *entered*, *arrived*, *returned*. Rate per thousand words, 73 texts over 800 words.

GENRE	N	NEUTRAL EVENT RATE	GIRARD RATE	RATIO
MYTHOGRAPHY	9	**28.65**	13.67	0.477
TRAGEDY	12	21.45	13.88	**0.647**
OVID	11	21.73	9.85	0.453
EPIC	36	22.68	9.48	0.418
CONTROL	5	20.77	7.77	0.374

Finding 1. `RESULTS-2.md`'s diagnosis is confirmed for the first time. Mythography really is the densest narrative prose in the corpus, and by a wide margin: 28.65 event-words per thousand against 21 to 23 for everything else. "Apollodorus and Diodorus are bare plot summaries, every word is event" was an impression. It is now a number.

Finding 2, and it goes against the confound story. Across all 73 texts the correlation between event rate and Girard total is **negative**: $r = -0.243$, permutation $p = 0.040$.

> So **density does not generally inflate Girard scores. It deflates them.** Epic has a high event rate and a low Girard score. The claim "the instrument partly measures how much plot per word a text carries" is **not supported as a corpus-wide effect.** What is true is narrower and stranger: mythography is an outlier on density *and* on score, and no general density mechanism explains that pairing.
>
> I also checked lexical variety as an alternative proxy. Type-token ratio on a fixed 4,000-token sample correlates with the Girard total at only $r = +0.198$, $p = 0.097$, and mythography has the **lowest** TTR in the corpus (0.269) while scoring near the top. **Two plausible density proxies, and neither explains the mythography result.** The confound named in July is real as a description and unproven as a mechanism.

Finding 3, and here is the trap. Normalise the Girard rate by the neutral event rate and the genre ordering changes:

```
event-rate ranking : MYTHOGRAPHY > EPIC > OVID > TRAGEDY > CONTROL
Girard    ranking : TRAGEDY > MYTHOGRAPHY > OVID > EPIC > CONTROL
Girard / event    : TRAGEDY 0.65 > MYTHOGRAPHY 0.48 > OVID 0.45 > EPIC 0.42
> CONTROL 0.37
```

The normalised ordering is **TRAGEDY > MYTHOGRAPHY ≈ OVID ≈ EPIC > CONTROL**, with tragedy clear of mythography by 0.17 where the raw rate had them tied at 0.04. **That is P1's predicted ordering, restored.**

> ***And that is exactly why it cannot be reported as a result.***
>
> *I have just built, after seeing P1 fail, the normalisation that turns a failed pre-registered prediction into a passing one. This is `CORRECTION-1.md` again, at the level of the **metric** instead of the sample. There the excluded page moved MYTHOGRAPHY 13.95 → 13.67 and flipped P1; here a denominator does the same job more elegantly. Both are defensible. Both were motivated by knowing the answer.*
>
> ***P1 remains a partial fail.*** *The event-normalised metric is a round-5 hypothesis, and it is only worth anything if its predictions are committed before it is scored on texts that have not been scored on it.*

9.2 Removing "king" from the marks bucket

`NEXT.md` item 3: "Fix `marks`. 44% 'king' is not a victim mark, it is a genre signal." Never done. The worry has been standing since 21 July.

Method. Drop *king, kings, queen, queens, kingly, queenly* from the `marks` bucket, rescore, and check whether any verdict moves.

GENRE	N	TOTAL AS-IS	DE-KINGED	MARKS AS-IS	MARKS DE-KINGED
MYTHOGRAPHY	9	13.67	11.42	3.33	1.08
TRAGEDY	12	13.88	11.48	4.37	1.98
OVID	11	9.85	8.50	2.22	0.86
EPIC	36	9.48	7.93	2.53	0.98
CONTROL	5	7.77	5.48	2.73	0.44

The bucket loses between 55% and 84% of its mass. Nothing moves.

	AS-IS	DE-KINGED
P2/P3, lowest named HIGH vs highest named LOW	10.07 vs 8.39, **holds**	8.50 vs 6.98, **holds**
Oedipus the King rank of 73	2	**3**
Bacchae rank	36	**35**

> **This is good news for the instrument and it retires a standing worry.** The "king" objection was the most obvious thing wrong with the lexicon and it turns out to be **cosmetic**: it inflates every text roughly proportionally rather than sorting them. Even *Oedipus the King*, the worst case by construction, drops one place.
>
> One detail runs the other way and is worth having. **The controls lose more proportionally than the tragedies** (marks 2.73 → 0.44, an 84% drop, against 4.37 → 1.98, 55%). So "king" was inflating the LOW texts harder than the HIGH ones, and removing it slightly **widens** the gap the prediction needed.

Still post-hoc, still not a pass. But `NEXT.md` item 3 can be marked as **probably not worth doing**, which is itself a saving.

9.3 How much of round 4 rests on one phrase family

`RESULTS-4.md` records that sixteen of John's twenty-six hits are the "some said / others said" formula, and asks whether that is Girard's broken unanimity or a translator's reporting convention. It does not test what happens if you remove it.

Method. Strip the crowd-speech concepts one at a time and rerun D1 and D2. Label permutation, 50,000 draws.

LEXICON	GOSPEL	TRAGEDY	MYTH	GOSPEL > MYTH, P
full, as run in round 4	1.387	0.667	0.547	**0.00042**
minus `some_others`	0.972	0.656	0.449	0.0135
minus `some_others` + `not_all`	0.874	0.611	0.438	0.0242
minus all three crowd-speech concepts	0.655	0.524	0.346	**0.0484**

The ordinal MYTH < TRAGEDY < GOSPEL survives every removal. D2 degrades smoothly, from p = 0.0004 to p = 0.048, and never inverts.

> **So round 4 is not a single-idiom artifact.** Removing the one formula the results file worried about costs a factor of thirty in p and still leaves it nominally significant. The effect is a **family** effect, spread across division, some-said, not-all and spare-him.
>
> **Which is arguably worse, not better.** A single formula is a quirk of one translator's rendering of one construction. A whole **crowd-speech register** is a systematic property of how the ASV narrates divided crowds, and it is **more** plausible as a translation convention, not less. `NEXT.md` warned that register was the largest threat to this test. This says the threat is broader than the results file located it.
>
> And it does not touch S1 at all. Every one of these concepts is `unan`. The `core` half still carries none of the Gospel effect.

9.4 The three re-admitted texts

Reported in full at Part VIII §8.2. Recomputed here for completeness: removing the three texts that pre-registration 2 excluded on the 800-word rule and round 4 silently re-admitted moves the MYTH median from 0.547 to 0.561 and the D6 counterexample count from 7 to 6. **No verdict changes.**

9.5 What these four analyses are worth

#	ANALYSIS	STATUS
9.1	density confound	**the most valuable and the most dangerous.** Confirms the description, refutes the general mechanism, and produces a metric that would flip P1. Round 5 hypothesis, nothing more
9.2	de-kinged marks	closes a standing worry as cosmetic. A saving, not a finding
9.3	crowd-speech family	refines round 4 against itself. The register threat is broader than located
9.4	re-admitted texts	confirms the rule lapse changed no verdict

None of them is a result. Three of them were asked for by the repository and should have been done there, before anyone knew what the answers were.

PART X — WHAT IS STILL BROKEN

In the order the repository itself ranks them.

#	ITEM	WHY IT MATTERS	STATUS
~~1~~	~~**The round-2 null script is not in the repo**~~	**CLOSED 14 Aug, round 5.** Reconstructed, original procedure identified, committed result reproduces under it. §1.18	**done**
~~1~~	~~**A genuinely phrase-aware null has never been built**~~	**CLOSED 14 Aug, round 6.** Built, and the phrase half passed at 0 of 2000. The translation control passed too. §1.19	**done**
~~1~~	~~**Localization**~~	**TESTED 14 Aug, rounds 7 and 8.** Round 7's null was wrong; round 8 corrected it and U1 passed at p = 0.0015. §1.20	**done**
~~1~~	~~**Everything lives in `unan` and nothing in `core`**~~	**DIAGNOSED 14 Aug.** `core` is three sub-families that cancel: the empty-accusation family runs 3.68× toward the Gospels, the advocacy family 17× away. §1.21	**explained**
1	**A corpus that has never been scored**	The `core` split cannot be tested on anything on disk: one spare text and six wrong-tier apocrypha. Round 9 needs new texts with the fetch list frozen first	**the new item 1**
2	**No denominator for the NT count**	"Two hard contradictions out of *what*" decides the entire verdict	specified, not built
3	**No control for the NT count**	The apocrypha are on disk. F3 is untested, so there is no evidence the instrument discriminates at all	not run
4	**The density confound**	Mythography ties tragedy because bare plot summary is the densest prose in the corpus	unresolved since round 2
5	**`marks` is 44% the word "king"**	It detects royalty, not victim marks	unfixed
6	**The phrase half of the dissent lexicon has never been null-tested**	It carries the entire round-4 effect	round 5, uncommitted
7	**Localization**	Two rounds died for lack of it. Victim voice share is the only localized measure available, and the diagnostic for it is already done	queued, unbuilt
8	**Set B of the conversion study**	P2 is the only test in that instrument that was ever going to mean anything	fetch was killed
9	**Dostoevsky and Proust**	Where Girard's strongest conversion claims live	not retrieved
10	**32 of 51 audits are tier K**	Their Greek and manuscript claims are leads, not facts	four could be settled in an hour
11	**Layer 2 Axis 2**	A cost 4 there invalidates the eliminative structure regardless of Axis 1	not scored
12	**No corpus integrity check**	Correction 1's truncated text was found by accident. A truncated text is invisible to every metric currently computed	not built
13	**More paired translations**	The whole project rests on one Agamemnon	not done
~~14~~	~~**The restraint control**~~	**CLOSED 14 Aug.** Run against its own pre-registration. H3 and H5 fired, H1/H2/H4 did not. Part IV	**done**
15	**Splitting gate A2** into lexical and conventional	The gate F4 fired on is doing two different jobs and nothing separates them	not recognised in the repo
16	**The Eusebian canon tables**	A stake-free enumeration of the cross-checkable set, made in the fourth century. A sampling frame for the missing denominator	needs its own pre-registration
17	***To Stephanus* 1-16**	Already on disk, bears on four docket items, the highest-value unread source in the project	unread

If the round-2 null script is rewritten and disagrees with the committed number, the rule already written says: the disagreement is the finding and gets reported, not reconciled.

PART XI — QUESTIONS YOU ANSWER

The answers are not elsewhere in this document. Write them in the margin.

On the git finding (7.1).

1. Is the file-manifest defence enough, or does round 2 need to be re-run from scratch under a procedure that leaves better evidence? What would that procedure look like?

2. What would a pre-registration procedure look like that a hostile reader could check **without** trusting the ordering of commits at all?

On the lapsed length rule (7.2).

3. The rule was declared, written down, given its own correction file, and then silently dropped two weeks later by the same project. What mechanism would have caught it? Name one that is not "be more careful."

4. Does the fact that the lapse changed no verdict make it better, worse, or neither? Argue it both ways before you answer.

On the instruments.

5. Rounds 3 and 4 died of the same cause: whole-text rates discard position. Is localization actually solvable with lexical methods, or is that cause fatal to the whole approach?

6. Round 4 §4 refuses the escape hatch and records Girard's universal negative as **untested**. Is "untested" the honest verdict, or is an instrument that cannot distinguish an Orphic epithet from a dissenting voice evidence that the *construct* is not lexically measurable at all?

7. The conversion study's §5.2 finding says the three pairs describe what happens to a **character** while "novelistic" names where the **author** stands. Is that a defect in the instrument, or a correction to Girard?

On the whole thing.

8. Which of the three instruments would you defend in front of someone who already thinks the enterprise is motivated? Which would you abandon?

9. Fourteen failures, nine passes, two fired falsifiers, one surviving p-value that cannot be reproduced. Is that a failed project or a working one? The answer is not in this document.

10. What is the one sentence you could say on camera that you would still defend in five years?

On the study that was never run (Part IV).

11. Pre-registration C has a control, a denominator, five numbered falsifiers, six confounds declared against itself, and an ordering rule that codes the unfavourable corpus first. It is the most disciplined file in the repository and it produced nothing. Does writing it count for anything at all, or is a pre-registration with nothing behind it worth exactly zero?

12. H2 fires if the **control** fails to behave as predicted, which would invalidate the instrument no matter how the canon scored. Name one other instrument you have built, in any project, that could detect its own yardstick being broken.

13. Confound C2 says that if restraint was itself a criterion of canonization, finding restraint in the canon measures the selection rule and not the composition process. Is that confound survivable, or does it mean the study should not be run at all?

On the analyses in Part IX.

14. 9.1 produces a metric that turns P1 from a fail into a pass, and 9.1 refuses to claim it. Write the pre-registration that would let a future round claim it. What corpus, what predictions, committed when?

15. 9.3 shows round 4's effect is a crowd-speech **family**, not one formula, and argues that is worse rather than better. Do you agree? A translator's systematic habit against a translator's single quirk. Which is more damaging to the claim?

16. 9.2 says the "king" problem is cosmetic and can be closed. That retires a worry that has been outstanding since 21 July on the strength of one post-hoc computation. Is

that legitimate, given every other rule in this project?

On the shape of the whole enterprise.

17. The interlock ratio has a numerator of five. The dissent test's whole effect lives in phrases never null-tested. The four-bucket detector cannot separate Gospel from myth by design. Is there a version of this project where the instruments are strong enough to matter, or is the honest output always going to be a claim about somebody's list rather than about the texts?

---

18. Every failure in this repository was reported. No omission was. Pre- registration C is missing from every handoff file. What would a rule look like that catches a study that was never run, the way F4 catches a coder who leaned?

PART XII — PROGRESS LOG

One line per pass. Date, what you read, what went fuzzy, what you would say differently.

DATE	SECTION	WHAT WENT FUZZY

Reproduction commands

```
cd ~/Myths
.venv/bin/python score.py corpus2              # rounds 1-2, whole-text scorer
.venv/bin/python score.py corpus2 --v1         # original lexicon
.venv/bin/python score.py corpus2 --no-new     # v2 matching, v1 concepts only
.venv/bin/python window.py corpus2             # peak_min locator
.venv/bin/python nulltest.py                   # round 3, both nulls
.venv/bin/python score_dissent.py              # round 4, every prediction
cd conversion && ../.venv/bin/python analyze.py
cd interlock/rescue && python3 bayes.py
cd interlock/greek && python3 check.py
git log --reverse --date=format:'%m-%d %H:%M' --pretty=format:'%ad  %h  %s'
```

`numpy` needs `.venv/bin/python`; Arch's system Python is externally managed. `corpus2/` and `corpus3/` are gitignored and rebuilt by `get_corpus2.py` and `build_corpus3.py`.